

Privacy Lens: An Evidence-Bounded Android Privacy Review Framework with Governed Legal-Evidence Packs (Revision 16)

Zhiheng Zhang

Student ID: 54560484

Supervisor: Richard Sinnott

August 2026

Abstract

Mobile privacy-review tools face an evidence-to-decision gap: a permission observation can indicate activity, but it cannot by itself establish purpose, necessity, lawful basis, acquisition completeness, or legal compliance. When an interface suppresses these missing facts, a technically correct calculation can still produce an unreliable conclusion. This thesis addresses that gap through Privacy Lens, a local-first Android prototype that converts bounded permission-audit summaries into traceable prompts for human review.

Following a design-science method, the architecture treats provenance, temporal scope, missing context, and uncertainty as first-class evidence rather than treating a count as a legal verdict. A fail-closed engine rejects malformed records and unsafe temporal aggregation, isolates simulator data, and records the rule-pack version applied to each finding. The European Union General Data Protection Regulation provides the evaluated legal objective; jurisdiction-specific mappings remain outside the Android, repository, and interface layers through a replaceable pack contract.

The implemented product organises the review journey around Overview, Findings, and Settings. Findings remain on the device, Android backup and remote updates are disabled, the evaluated package has no network or sensitive runtime permission, and users can clear local evidence directly. Revision 15 added source-status records, calendar-valid governance dates, basis-specific Article 6 evidence, transparency, minimisation, retention-justification, and DPIA-screening questions. Revision 16 exposes the governed source register in the interface, introduces non-colour status structure, stabilises empty-to-populated and disclosure interactions, and applies a coherent editorial visual system without weakening the legal caveats.

Evaluation separates rule conformance, hostile admission, temporal behaviour, pack identity, legal-evidence completeness, release packaging, manifest facts, physical installation, and rendered interface evidence. In a fixed-seed 1,000-case corpus, the engine produced 800 true positives and 200 true negatives without disagreement against the specified threshold oracle. A further 1,800 independently calculated pack cases, basis-specific and DPIA evidence probes, governance mutations, strengthened accessibility contracts, and release-privacy contracts passed. The 1.4.0 ARM64 APK and AAB built successfully; the APK completed five final cold-start rounds on one OPPO PERM00 without a package-specific crash-buffer entry. A four-screen final visual matrix covered overview, populated findings, pack selection, and source-status presentation.

These results establish an initial store-submission engineering candidate within the recorded scope. They do not establish legal compliance, unrestricted Android visibility, population accuracy, accessibility conformance, long-run reliability, or store approval. The contribution is an evidence-bounded architecture and traceability method that separate statutory motivation, engineering heuristics, observed evidence, missing legal facts, and prohibited conclusions.

Contents

1	Introduction	6
1.1	Problem Statement: The Evidence-to-Decision Gap	6
1.2	Why Mobile Permission Evidence Matters	6
1.3	From Technical Signal to Review Finding	7
1.4	Consent, Attention, and Interface Burden	8
1.5	Accountability Rather Than Automated Adjudication	8
1.6	Regional Extensibility as a Governance Problem	9
1.7	Research Method	9
1.8	Research Questions	11
1.9	Research Objectives and Contributions	11
1.10	Scope and Claim Boundary	12
1.11	Thesis Structure	12
2	Literature Review and Foundations	12
2.1	Review Protocol and Source Quality	13
2.2	Regulatory Principles and Technical Signals	13
2.2.1	From Legal Principle to Review Question	14
2.2.2	Lawful-Basis Evidence Is Basis Specific	14
2.2.3	Data Protection by Design and Interface Conduct	15
2.3	Android Visibility and Permission Semantics	15
2.3.1	A Taxonomy of Mobile Privacy Evidence	15
2.3.2	Historical Operations and the Ordinary-App Boundary	16
2.4	Human Factors and Calibrated Trust	17
2.4.1	Cognitive Load and Progressive Disclosure	17
2.4.2	Trust Calibration and Explanation Quality	18
2.4.3	Accessibility as an Evidence Requirement	18
2.5	Replaceable Rule Systems	18
2.5.1	Pack Lifecycle and Review Governance	19
2.5.2	Rule-Based and Learned Decision Models	19
2.5.3	Automated Validation and Oracle Independence	20
2.6	Comparative Positioning Against Prior Android Studies	20
2.7	Research Gap: From Observation to Calibrated Review	21
2.8	Conceptual Model: Evidence, Rule, Claim, and Action	22
2.9	Store Delivery as a Trust Boundary	23
2.10	Synthesis	23
3	Requirements and System Design	23
3.1	Design Objectives and Requirements	23

3.2	Design Alternatives and Threat Model	24
3.2.1	Selection Rationale	25
3.3	Architecture and Decision Lifecycle	26
3.3.1	End-to-End Decision Lifecycle	27
3.3.2	Failure-State Design	28
3.3.3	Data-Minimisation Decisions	29
3.3.4	Requirement Priorities and Release Gates	30
3.4	Stakeholders and Functional Requirements	31
3.4.1	Stakeholders and Usage Scenarios	31
3.4.2	Functional Decomposition	32
3.4.3	Non-Functional Requirements	32
3.5	Evidence and Regulation-Pack Design	33
3.5.1	Evidence Schema and Provenance	33
3.5.2	Admission and Fail-Closed Validation	34
3.5.3	Temporal Isolation Model	35
3.5.4	Rule-Pack Contract in Detail	35
3.5.5	State, Persistence, and Migration	36
3.5.6	Security and Abuse Cases	36
3.6	Traceability and Evaluation Design	37
3.6.1	Traceability from Requirement to Evidence	37
3.6.2	Evaluation Strategy	37
3.6.3	Design Invariants	37
3.6.4	Decision Records for Future Extension	38
4	Implementation	38
4.1	Implementation Overview and Module Boundaries	39
4.2	Rule and Evidence Pipeline	40
4.2.1	Regulation-Pack Implementation	40
4.2.2	Pack Governance Lifecycle	41
4.2.3	Runtime Validation Pipeline	42
4.2.4	Signal Computation	43
4.2.5	Finding Construction and Communication	44
4.2.6	Repository Behaviour	45
4.3	Android Bridge and Scheduling	45
4.4	Product Interface and Accessibility	45
4.4.1	Interface Implementation	46
4.4.2	Brand and Accessibility Implementation	46
4.5	Release Configuration and Privacy Controls	47
4.6	Implementation-to-Requirement Traceability	48

5	Evaluation and Results	49
5.1	Evaluation Overview and Headline Results	49
5.1.1	Source and Rule Verification	49
5.1.2	Legal-Evidence Completeness Verification	50
5.1.3	Release Build and Manifest Verification	50
5.1.4	Physical-Device Evaluation	51
5.1.5	Usability and Trust Assessment	51
5.2	Protocol and Reproducibility	52
5.2.1	Evaluation Questions and Acceptance Logic	52
5.2.2	Claim-to-Evidence Acceptance Matrix	53
5.2.3	Reproducibility Controls	54
5.3	Engine and Regulation-Pack Results	54
5.3.1	Deterministic Corpus Interpretation	54
5.3.2	Adversarial Input Results	55
5.3.3	Temporal and Replay Results	55
5.3.4	Accountability and Presentation Results	56
5.3.5	Regulation-Pack Results	56
5.4	Android Artefact and Device Results	57
5.4.1	Build Reproduction	57
5.4.2	Manifest and Installed-Package Reconciliation	57
5.4.3	Physical Interaction Protocol	58
5.4.4	Visual Review Method	58
5.5	Interpretation and Validity	59
5.5.1	Result Boundary	59
5.5.2	Evidence Triangulation	59
5.5.3	Oracle Independence and Correlated Defect Risk	61
5.5.4	Metric Interpretation and Error Analysis	61
5.5.5	Store-Candidate Evidence Review	62
5.5.6	Evaluation Repeatability and Change Control	63
5.6	Evaluation Synthesis	64
6	Discussion and Limitations	65
6.1	Interpretation of Findings	65
6.1.1	What the Prototype Establishes	65
6.1.2	Construct and Internal Validity	65
6.1.3	External and Ecological Validity	65
6.1.4	Legal and Regional Validity	66
6.1.5	Usability and Accessibility	66
6.1.6	Interpreting the Engineering Result	66

6.2	Technical Validity Boundaries	67
6.2.1	Acquisition Authority and Observation Quality	68
6.2.2	Threshold Meaning, Calibration, and Decision Error	69
6.2.3	Temporal Evidence and Replay Semantics	70
6.2.4	Regulation-Pack Governance Beyond Software Modularity	71
6.3	Human, Ethical, and Operational Trust	72
6.3.1	Trust, Comprehension, and Willingness to Use	72
6.3.2	Privacy and Security of the Reviewing Application	73
6.3.3	Research Ethics and Responsible Evaluation	74
6.3.4	Store and Operational Readiness	75
6.4	Cross-Disciplinary Assurance Case	75
6.5	Prioritised Next Work	76
7	Conclusion	77
7.1	Answers to the Research Questions	77
7.2	Contributions and Research Significance	79
7.3	Final Product Position	80
7.4	Implications for Privacy-Engineering Practice	80
7.5	Future Work	81
7.6	Closing Statement	81

1 Introduction

This chapter defines the problem that Privacy Lens is designed to solve and the claims that the thesis is prepared to defend. The central concern is not whether Android exposes a permission fact, but whether that fact can be transformed into a useful review prompt without acquiring legal meaning that the evidence does not contain. The chapter therefore moves from the mobile privacy context to the evidence-to-decision problem, the research questions, and the contribution boundary.

1.1 Problem Statement: The Evidence-to-Decision Gap

Mobile applications ask users to make privacy decisions while the evidence needed to evaluate those decisions remains fragmented. A permission declaration states capability, not necessarily use; a runtime observation can indicate activity, but not purpose, lawful basis, necessity, or proportionality. The Regulation requires legal and organisational analysis in addition to technical safeguards, including the principles in Article 5 and data protection by design and by default in Article 25 [1, 2].

Privacy Lens addresses a narrower engineering problem: how can an ordinary Android application present available permission-audit evidence without transforming incomplete telemetry into a false legal conclusion? The application collects no remote evidence, has no account or advertising layer, and stores findings locally. It exposes whether a finding came from device-visible observations or a synthetic demonstration, displays missing-evidence limitations, and prepares questions for human review.

The work began as a GDPR-oriented prototype. The evaluated release retains the GDPR as its legal objective but removes jurisdiction-specific interpretation from the core decision mechanism. Legal mappings are loaded through replaceable regulation packs. This separation makes regional extension possible while preventing the interface, repository, and Android bridge from silently inheriting one jurisdiction’s terminology.

1.2 Why Mobile Permission Evidence Matters

Mobile privacy is difficult to reason about because a handheld device combines intimate data, continuous availability, and many independent software actors. A user may encounter an application as a single icon, yet that application can depend on operating-system services, embedded libraries, cloud endpoints, analytics components, advertising identifiers, and background tasks. The apparent simplicity of a permission prompt therefore hides a larger processing system. The prompt may describe access to a sensor or data category, but it does not normally explain every purpose, recipient, retention period, inference, or downstream combination. Contextual integrity theory is useful here because it treats privacy as the appropriateness of information flows within a social context rather than as secrecy alone [3]. A technically permitted flow can still be surprising or inappropriate when its purpose, recipient, or persistence differs from the context understood by the user.

The mismatch matters across application domains because camera, microphone, location,

contacts, and activity signals can reveal sensitive behaviour even when an operating-system permission is technically valid. Privacy Lens therefore does not infer purpose or legality from a permission name. It preserves the observed signal, its source, its time window, and the questions that remain unanswered. This domain-neutral position keeps the architecture applicable to different application categories without claiming knowledge that the evidence contract does not contain.

Android’s permission model provides necessary protection, but permission state and processing purpose are different constructs. A granted permission indicates that an application may invoke a protected capability under specified platform conditions. It does not state whether the capability was invoked, why an invocation occurred, whether derived data were retained, or whether the processing satisfied a legal obligation. Conversely, a denied or absent permission does not prove the absence of all related data because an application may receive information through another service, user input, previously collected records, or a less sensitive signal from which attributes can be inferred. The prototype therefore treats permission observations as one evidence family within a broader review, not as a complete representation of processing.

Prior Android studies show why permission state must be interpreted with behaviour and context. Wijesekera and colleagues instrumented Android for a 36-person field study and found substantial mismatch between protected-resource access and participant expectations [4]. Cao and colleagues studied 1,719 participants across ten countries and regions; expectation and application-provided explanation were associated with permission decisions [5]. More recently, Prange and colleagues followed 132 Android users and observed that revocations often occurred in clusters and particularly affected rarely used applications or permissions judged non-essential [6]. These studies supply behavioural evidence that Privacy Lens itself does not possess. The present contribution is complementary: it does not predict what a user will accept, but preserves enough provenance and uncertainty for an eventual reviewer or participant to judge a technical prompt in context.

1.3 From Technical Signal to Review Finding

The central research problem can be described as an evidence-to-decision gap. At one end of the gap are platform facts: a package identifier, permission category, access count, observation window, acquisition source, and timestamp. At the other end are questions that matter to a person or regulator: Was processing expected? Was it necessary for the stated purpose? Was a valid lawful basis available? Were special-category conditions satisfied? Were recipients and retention disclosed? Was the data subject able to exercise meaningful control? No deterministic transformation can recover answers that were never present in the input.

Two common design errors follow from ignoring this gap. The first is semantic inflation, in which a technical event is labelled as a legal violation. This gives a small signal the authority of a much broader assessment. The second is evidential silence, in which an empty or unavailable audit is displayed as a reassuring green state. On stock Android, ordinary applications often lack authority to inspect unrestricted cross-application history. Treating an empty result as absence of activity would therefore reward the system precisely when its visibility is weakest.

Privacy Lens introduces an intermediate output: a review finding. A finding can say that the available count crossed a research threshold, that evidence is unavailable, or that no technical concern was identified within the observed window. It can also state which contextual facts are missing. This format preserves utility without claiming certainty. The distinction is comparable to the difference between a screening instrument and a diagnosis: screening can identify cases that deserve attention, but its result must be interpreted in context and by an appropriately qualified actor.

The design also recognises that explanation is not an optional layer added after computation. Explanation begins with data modelling. If the repository stores only a severity colour and a message, the interface cannot later reconstruct the source, rule version, temporal scope, or missing context. Revision 15 retained provenance fields and added basis-specific legal-evidence references and source-status records to the decision context. Revision 16 makes that structure visible through source-authority cards, text markers, status rails, and provenance chips rather than allowing polished colour alone to imply authority. The interface is a view of the evidence object rather than a separate narrative that can drift away from it.

1.4 Consent, Attention, and Interface Burden

Privacy interfaces operate under severe attention constraints. McDonald and Cranor quantified the impractical time cost of reading the privacy policies encountered by an ordinary user [7]. Obar and Oeldorf-Hirsch demonstrated how readily people accept online terms without substantive engagement [8]. These findings should not be interpreted as evidence that users are careless. They show that notice-and-choice systems frequently demand more time and legal comprehension than a routine interaction can support.

A warning application can reproduce this failure if it presents a wall of legal text, asks the user to interpret raw telemetry, or interrupts every time a count changes. More information is not automatically more usable transparency. Information must be staged according to the decision a user can reasonably make. Privacy Lens therefore begins with three questions: what requires attention, what evidence is missing, and what can be done next. Rule references and technical details remain available, but they follow the status and action rather than preceding them.

This approach also affects notification strategy. Repeated alerts for an unchanged state can teach users to ignore the product. A review system should distinguish new evidence, a material escalation, and a stable finding. Although the present candidate does not claim a longitudinal notification study, its repository retains enough state to support transition-based notification in future work. The architectural principle is that attention is a limited resource and should be consumed only when evidence meaningfully changes.

1.5 Accountability Rather Than Automated Adjudication

The GDPR’s accountability principle requires a controller not only to comply with the Regulation but also to be able to demonstrate compliance [1]. A consumer-facing application cannot fulfil the controller’s entire accountability obligation, because it does not possess the

controller’s records of processing, contracts, purposes, retention decisions, risk assessments, or organisational measures. It can nevertheless support accountability by making a narrow technical observation reproducible and by showing which questions remain unanswered.

This positioning changes the desired failure mode. A conventional classifier may optimise for a single label and treat rejected input as an implementation inconvenience. An accountability tool should prefer an explicit hold or evidence gap to a falsely reassuring result. Unsafe types, unrecognised sources, impossible windows, and missing context therefore fail closed. The system does not invent a value to keep the pipeline moving. This behaviour is especially important at a native bridge, where JavaScript assumptions can be violated by malformed, stale, or unexpected platform data.

Accountability also requires version knowledge. A finding that says only GDPR is not reproducible if rules or interpretations later change. Privacy Lens records a pack identifier and version, an official source, and a particular rule reference. Historical findings are not silently reinterpreted when the active pack changes. This creates a minimal audit trail: a reviewer can identify which data were available, which pack was applied, and why a particular message appeared.

1.6 Regional Extensibility as a Governance Problem

The request for regional switching creates both an architectural opportunity and a governance risk. It is straightforward to define a common software interface for multiple packs. It is much harder to ensure that each pack represents its jurisdiction accurately. Legal systems differ in scope, terminology, actors, exemptions, enforcement, and the relationship between national and sectoral rules. A design that assumes every jurisdiction can be translated into GDPR article equivalents would merely hide legal coupling behind a new file format.

Revision 15 therefore separates software extensibility from legal validity and separates binding law, final guidance, and consultation material within the source register. The registry proves that the application can discover, select, persist, and apply different pack objects. It does not certify the content of every possible pack. The EU GDPR pack is the evaluated legal objective. The Research Baseline pack is deliberately non-legal and demonstrates the mechanism without suggesting that another jurisdiction has been reviewed. A future regional pack must identify its authority, effective date, scope, official source, reviewer, localisation status, fixtures, migration policy, and limitations.

This distinction is central to the thesis contribution. Decoupling is valuable not because it makes law interchangeable, but because it makes jurisdiction-dependent assumptions visible and reviewable. When a legal mapping is isolated, it can be versioned, challenged, replaced, and tested without changing the Android bridge or product shell. The architecture therefore supports expansion while making the cost of responsible expansion explicit.

1.7 Research Method

The work follows a design-science research method. Design science treats the construction and evaluation of an artefact as a route to knowledge about a problem and its possible solution [9].

Peffer and colleagues organise this work into problem identification, definition of solution objectives, design and development, demonstration, evaluation, and communication [10]. Privacy Lens follows that sequence iteratively rather than presenting implementation as a single linear build.

First, legal and platform constraints are translated into claim boundaries. Second, these boundaries become data and interface requirements. Third, the requirements are implemented across TypeScript, React Native, Kotlin, storage, and Android build configuration. Fourth, each layer is demonstrated and evaluated with evidence appropriate to that layer. Finally, contradictions between intended claims and observed behaviour are used to revise both code and thesis. Table 1 makes the relationship between research method, artefact, and evidence explicit.

Table 1: Design-science stages and thesis evidence

Stage	Action in this study	Inspectable evidence
Problem identification	Define the evidence-to-decision and authority gaps	Research questions, claim boundary, prior Android and legal literature
Objectives	Require safe admission, calibrated communication, pack isolation, and reproducibility	Requirements F1–F5, U1–U2, P1, and D1
Design and development	Construct the engine, pack contract, repository, bridge, and interface	Versioned source modules and architecture traceability
Demonstration	Exercise controlled and device-visible review journeys	Synthetic workflow, release APK, installed-device screenshots, and UI trees
Evaluation	Test each claim with a layer-appropriate oracle	Boundary suites, artefact hashes, manifest reconciliation, and physical interaction record
Communication	State supported claims and unresolved validity threats	Thesis, review reports, versioned PDFs, and explicit limitations

The method combines constructive and empirical elements. The constructive element is the artefact: the rule-pack contract, compliance engine, repository, Android bridge, and interface. The empirical element consists of deterministic test vectors, adversarial boundary cases, build outputs, merged-manifest inspection, installed-package inspection, ADB interaction, UI hierarchies, screenshots, and filtered runtime logs. The method does not include a controlled participant experiment or independent legal assessment; those omissions are recorded as limitations rather than replaced with speculative conclusions.

An important methodological rule is that evidence is not pooled across incompatible layers. A thousand logical cases do not become a thousand device trials. A successful installation does not validate every WorkManager schedule. A well-rendered warning does not prove comprehension. A rule reference does not prove infringement. Maintaining these separations allows the final evaluation to be positive about what worked without obscuring what remains

unknown.

1.8 Research Questions

The thesis answers four questions that connect evidence admission, communication, extensibility, and evaluation.

RQ1: How can Android permission-audit signals be transformed into deterministic, explainable review findings while unsafe or incomplete inputs fail closed? This question concerns the runtime evidence contract, decision semantics, temporal isolation, and replay behaviour. It is evaluated through exact-boundary, adversarial, and stateful logical tests.

RQ2: How can provenance, temporal scope, and missing legal context be communicated so that a technical warning is not mistaken for a legal verdict? This question concerns finding construction, dashboard projection, status language, source disclosure, and the distinction between synthetic evidence and device-visible observations. It is evaluated through presentation fixtures and rendered workflows, with participant comprehension retained as future work.

RQ3: How can jurisdiction-specific rules be isolated behind a stable regulation-pack contract without weakening traceability? This question concerns explicit dependency injection, registry admission, immutable pack metadata, historical identity, and visible caveats. It is evaluated by applying the same admitted evidence under two registered packs and comparing the resulting findings.

RQ4: What claims are supported by deterministic tests, release builds, and bounded physical-device evaluation, and which claims remain unverified? This question concerns the relationship between software evidence and research conclusions. It is answered through a claim-to-evidence matrix that keeps logical, package, device, legal, and deployment claims separate.

1.9 Research Objectives and Contributions

The research has four operational objectives. It must make each decision reproducible, communicate uncertainty without neutralising usefulness, isolate regional interpretation from generic product infrastructure, and produce an Android artefact that can be evaluated beyond source code. These objectives are realised through permission, source, time-window, pack, rationale, and missing-context fields; three non-verdict status classes; a registry of immutable pack definitions; and release APK and Android App Bundle outputs targeting SDK 36.

The **architectural contribution** is an evidence-bounded pipeline in which admission, temporal preparation, rule evaluation, persistence, projection, and presentation preserve the same claim boundary. The **methodological contribution** is a layered evaluation record that distinguishes agreement with a specified oracle from acquisition coverage, device behaviour, legal validity, and user comprehension. The **practical contribution** is a local-first Android candidate whose Overview, Findings, and Settings journeys expose source, selected pack, missing evidence, and data controls without requesting sensitive runtime permission.

The contribution is therefore not an automated GDPR judge. It is an accountability-support prototype that makes a bounded technical inference reconstructable and leaves legal conclusions to an appropriately authorised human reviewer. This positioning is narrower than the strongest claim a compliance application might advertise, but it is the strongest claim supported by the implemented evidence path.

1.10 Scope and Claim Boundary

The evaluated data path covers permission summary records available to the application or generated by its controlled simulator. Stock Android restricts ordinary applications from reading unrestricted activity belonging to other applications. Absence of an observation is consequently an evidence gap, not proof that an access did not occur. The current rule thresholds are research baselines; they are not statutory limits. The included EU GDPR pack is a traceable mapping aid, not legal advice.

Evaluation covers source compilation, deterministic logical tests, adversarial boundary tests, release assembly, manifest inspection, installation, launch, interaction, and short-run observation on one OPPO PERM00 device. It does not establish multi-OEM compatibility, 24-hour scheduler reliability, battery performance, population accuracy, accessibility conformance, independent legal validity, or application-store approval.

1.11 Thesis Structure

Chapter 2 establishes the legal, Android, HCI, and software-architecture foundations. Chapter 3 derives requirements and the evidence-bounded architecture. Chapter 4 documents the implemented application and regulation-pack mechanism. Chapter 5 reports reproducible evaluation evidence. Chapter 6 analyses validity, usability, deployment, and legal limitations. Chapter 7 answers the research questions and defines the next validation programme.

In summary, the thesis begins from a semantic constraint: technical evidence can support review only when its acquisition scope and missing context remain visible. The next chapter positions this constraint within regulatory principles, Android platform semantics, human factors, and prior approaches to rule-based privacy support.

2 Literature Review and Foundations

This chapter examines four bodies of work that constrain the design: data-protection principles, Android observation mechanisms, human-centred warning design, and replaceable rule systems. The purpose is not to claim that Privacy Lens supersedes specialised legal, static-analysis, or enforcement tools. It is to identify the unresolved integration problem between a bounded technical observation and the conclusion a user is invited to draw from it.

2.1 Review Protocol and Source Quality

The chapter uses a structured scoping review rather than claiming a systematic review or meta-analysis. Sources were selected to answer four design questions: what a permission observation can establish, which legal facts it omits, how people interpret permission and warning interfaces, and how a replaceable decision mechanism can remain auditable. The search was updated through 10 August 2026 and included official EU and Android material, peer-reviewed security and usable-privacy studies, standards, and foundational information-systems research. Seminal older work was retained when it defined a construct still used in the design, such as contextual integrity, cognitive load, trust calibration, usability, or design science.

Source authority was ranked by the proposition being supported. Statutory meaning is grounded in the GDPR text and final EDPB guidance, not vendor summaries. Platform capability is grounded in Android documentation and bounded by installed-device evidence. Behavioural claims use original field or user studies and preserve their sample sizes and contexts. Engineering-method claims use original design-science publications. Standards and NIST privacy-engineering material are used as design guidance, not as proof of legal compliance [11, 12]. Consultation materials are retained only when their non-final status is explicit; they do not displace binding law or final guidance. Commercial forecasts, uncited practitioner claims, and sources whose title or bibliographic record could not be verified were excluded from the Revision 15 argument.

This protocol has limits. It does not report database-by-database recall, duplicate screening, or a PRISMA flow, and therefore must not be read as exhaustive coverage of mobile privacy research. Its purpose is narrower: make each source class and its evidential role explicit so that a reviewer can distinguish legal authority, platform specification, empirical user evidence, and design rationale.

2.2 Regulatory Principles and Technical Signals

The GDPR regulates processing rather than permission counts. Article 5 establishes principles including lawfulness, transparency, purpose limitation, data minimisation, accuracy, storage limitation, integrity, confidentiality, and accountability. Articles 6 and 9 concern lawful bases and special-category data; Article 25 requires data protection by design and by default [1]. EDPB guidance describes Article 25 as a continuing obligation implemented through appropriate technical and organisational measures [2]. A technical observation may support an accountability inquiry, but cannot by itself demonstrate whether these legal conditions are satisfied.

This distinction determines the framework’s vocabulary. A threshold exceedance is a *review signal*. Missing purpose, lawful-basis, necessity, data-subject, controller, retention, and sharing evidence is shown explicitly. The UI never derives a finding labelled “GDPR violation” or “compliant”. Regulation-pack references provide a reason to investigate and a link to an official source; they do not replace legal interpretation.

2.2.1 From Legal Principle to Review Question

Operationalisation is necessary in any regulatory technology, but it must preserve the semantic distance between a principle and a sensor. Article 5(1)(a) concerns lawfulness, fairness, and transparency. A permission count can contribute to transparency by revealing a pattern that was previously invisible, yet it cannot establish lawfulness or fairness. Article 5(1)(b) concerns purpose limitation. A count contains no purpose unless purpose is supplied as contextual evidence. Article 5(1)(c) concerns data minimisation. Frequency can be relevant to necessity, but a high frequency may be required for continuous navigation while a single collection may be excessive for an unrelated purpose. Article 5(2) concerns accountability. An audit record can support demonstration, but it is only one component of the records, decisions, safeguards, and governance required from a controller [1].

Articles 6 and 9 further show why a permission signal cannot become a verdict. Article 6 provides several possible lawful bases; consent is not automatically the only basis. Article 9 establishes a prohibition and exceptions for special categories, but not every Android dangerous permission maps directly to special-category processing. Location, contacts, microphone, and camera data are context-dependent. Audio might contain health information, ordinary conversation, environmental noise, or no retained personal data. A technically observed invocation does not identify which of those states occurred.

The pack therefore maps a signal to a sequence of questions. For a location-related pattern, the reviewer may ask what purpose was declared, whether continuous precision was necessary, whether a less intrusive source was available, whether the application was in the foreground, whether location was retained or transmitted, and whether the user could reasonably anticipate the flow. For a microphone-related pattern, the reviewer may ask whether recording was user-initiated, how activation was signalled, whether raw audio left the device, and how long any derivative was retained. These questions are not answers; their value is that they prevent the visible count from crowding out legally decisive facts.

2.2.2 Lawful-Basis Evidence Is Basis Specific

Recording the name of an Article 6 basis is not equivalent to demonstrating that its conditions are met. Consent requires evidence capable of supporting the controller’s burden under Article 7 and must remain freely given, specific, informed, and unambiguous; withdrawal must be possible without detriment [1, 13]. Contractual necessity asks whether the processing is objectively necessary for performance, not merely mentioned in terms. Legal obligation and public task require an appropriate Union or Member-State-law anchor. Vital interests require a necessity assessment. Legitimate interests require a purpose, necessity, and balancing analysis, while the EDPB Article 6(1)(f) material used in this revision remains labelled as consultation material rather than final guidance [14].

Transparency and retention are cross-cutting rather than basis substitutes. Articles 13 and 14 require the applicable information to be provided, while Article 5(1)(e) ties retention to necessity rather than to an arbitrary number of days. Revision 15 therefore asks for a notice reference, a minimisation assessment, and a retention justification in addition to the numeric

period. Where a screening concludes that Article 35 applies, a DPIA reference is required. The endorsed DPIA guidance remains authoritative background; the 2026 EDPB template is recorded separately as consultation material pending finalisation [15, 16].

2.2.3 Data Protection by Design and Interface Conduct

Article 25 directs controllers to implement appropriate measures and privacy-protective defaults, taking account of state of the art, cost, context, and risk. EDPB guidance emphasises that this is a continuing obligation rather than a certification achieved once [2]. For Privacy Lens, data protection by design affects both internal processing and the claims made to users. Local storage, disabled backup, no analytics, no advertising, no account, explicit deletion, and minimised permissions reduce the application’s own data footprint. Fail-closed validation and immutable pack metadata reduce the risk that incomplete evidence becomes a misleading result.

Interface conduct is also relevant. A privacy tool could employ alarming colours, default users into sharing diagnostic data, hide evidence limitations, or imply that purchasing a service resolves risk. EDPB guidance on deceptive design patterns provides a broader reminder that formal information can still be presented in a manipulative way [17]. Privacy Lens uses restrained severity language, places limitations near the result, and keeps the demonstration optional. The implementation is not claimed to have been audited against every deceptive-design criterion, but the guidance informs the design objective of calibrated rather than maximised reliance.

2.3 Android Visibility and Permission Semantics

Android permissions are part of a layered security model. Manifest declarations, runtime grants, AppOps state, operating-system services, and OEM behaviour do not expose one universal event stream to an ordinary third-party application. Android’s permission guidance emphasises minimising permission requests and using platform-supported workflows [18]. Therefore, a credible review tool must state its acquisition authority and must not equate an empty result with absence of processing.

Static manifest analysis and dynamic monitoring answer different questions. Static tools can identify declared capabilities and embedded components; dynamic or taint-based systems can observe selected flows but may require instrumentation, system modification, a VPN, or elevated authority. PiOS illustrates the value and complexity of static data-flow analysis [19]. Privacy Lens deliberately occupies a lighter position: it evaluates bounded audit summaries and communicates their coverage rather than claiming universal enforcement or surveillance.

2.3.1 A Taxonomy of Mobile Privacy Evidence

Existing approaches can be organised by the claim their evidence supports. Manifest analysis answers what capabilities an application declares and which components are packaged. It is reproducible and can be performed without executing the application, but it cannot establish

that a declared capability was exercised. Static data-flow analysis attempts to connect sources and sinks in program code. It can reveal possible flows beyond manifest declarations, although reflection, native code, dynamic loading, and environment-specific paths complicate coverage. PiOS and related systems demonstrate the analytical value of tracing potential information flows while also illustrating that such analysis is specialised rather than a routine end-user feature [19].

Runtime instrumentation answers which monitored operations occurred during a particular execution. Its result depends on exercised paths, instrumentation placement, operating-system authority, and the definition of an event. Network interception answers which traffic crossed an observed interface and may expose destinations or payload structure. It cannot see purely local processing and may be limited by encryption, certificate pinning, protocol choice, and traffic that bypasses the observation path. Platform privacy dashboards can offer authoritative information for selected permissions, but their presentation and export capabilities are determined by the operating system.

Enforcement systems answer yet another question: whether a protected operation was blocked or modified under a policy. Enforcement is stronger than observation for the operation under control, but it can require device-owner status, system integration, rooting, a custom operating system, or application rewriting. These deployment choices change the threat model. A VPN monitor, for example, gains visibility by becoming an intermediary and must itself be trusted. A rooted monitor can inspect more state but weakens assumptions that apply to ordinary consumer devices.

Privacy Lens does not attempt to dominate this taxonomy. It accepts a structured summary from a bounded source and focuses on interpretation, traceability, and communication. This makes it complementary to manifest scanners, platform dashboards, enterprise management systems, and forensic monitors. A future authorised deployment could supply richer evidence through the same decision contract, but the application should continue to display the authority and coverage of that source.

2.3.2 Historical Operations and the Ordinary-App Boundary

The existence of an Android API is not evidence that any application can use it for any package. Access control can depend on signature permissions, usage-access grants, AppOps policy, device-owner status, process identity, Android release, and OEM modification. Android’s AppOps documentation states, for example, that callers without the privileged `WATCH_APPOPS` permission are limited to their own UID for relevant active-state queries and watchers [20]. Historical-operation methods may therefore be public in an SDK while useful cross-package results remain restricted in ordinary deployment. The physical-device evidence collected for this thesis exposed application-owned modes and scheduling state; it did not demonstrate unrestricted enumeration of other packages.

This boundary affects evaluation design. A complete system has an acquisition layer and a decision layer. The acquisition layer determines what the operating system or an authorised integration provides. The decision layer validates, aggregates, maps, stores, and presents that

evidence. Synthetic injection can test the second layer exhaustively without proving the first. Conversely, an acquired record can demonstrate that a bridge works without proving the correctness of every decision boundary. The thesis reports these layers separately.

2.4 Human Factors and Calibrated Trust

Privacy notices impose substantial reading and comprehension costs [7, 8]. Adding another dense dashboard can reproduce the problem it intends to solve. Trustworthy warning design should therefore support quick orientation while preserving access to rationale and provenance. ISO 9241-11 frames usability through effectiveness, efficiency, and satisfaction in a specified context of use [21]. WCAG 2.2 and Android accessibility guidance further motivate non-colour status cues, meaningful labels, and sufficiently large touch targets [22, 23].

These sources support four design choices. First, the overview leads with the current evidence boundary. Second, colour is paired with iconography and text. Third, details are progressively disclosed through finding cards. Fourth, destructive actions such as clearing local findings are separated from routine navigation. These are conformance-oriented design decisions, not proof of usability; participant studies remain necessary.

2.4.1 Cognitive Load and Progressive Disclosure

Cognitive Load Theory distinguishes load inherent to a task from load introduced by its presentation [24]. Privacy review has substantial intrinsic complexity because it combines technical events, purposes, law, expectations, and uncertainty. A product cannot remove that complexity, but it can avoid adding unnecessary navigation, decoration, or jargon. The earlier prototype mixed health dashboards, research projects, experiment controls, scoring views, and privacy findings. That structure made the user’s immediate task unclear. The three-destination Revision 14 architecture reduces extraneous choice and uses a stable pattern for each finding.

Progressive disclosure is not equivalent to hiding caveats. The most important caveat—the limited reach of evidence—appears on the overview before the user initiates a review. Detail cards then reveal the observed count, source, pack, reference, rationale, and missing context. This ordering allows rapid orientation while preserving auditability. A legal citation without an explanation would be too terse; a full reproduction of statutory text on every card would be too heavy. The interface therefore supplies a short reference and offers the official source separately.

Dual-process accounts distinguish fast, heuristic judgement from slower deliberation [25]. Severity colour can produce a rapid response, which is useful for prioritisation but dangerous if colour appears to communicate legal certainty. Textual labels, source chips, explicit caveats, and neutral action language are intended to interrupt an overly simple red-equals-illegal inference. This intention requires empirical testing. The thesis does not infer comprehension from the presence of these elements.

2.4.2 Trust Calibration and Explanation Quality

Trust is calibrated when reliance corresponds to actual capability. A system that understates uncertainty can induce over-reliance; a system that only displays caveats may become unusable and be ignored. Lee and See describe the importance of appropriate reliance on automation [26]. In Privacy Lens, calibrated trust is pursued through claim symmetry: a positive-looking result and a warning must both show their evidence basis. A no-technical-concern state means that no rule fired within the admitted evidence, not that processing was lawful. A needs-review state means that a rule fired, not that a violation occurred.

Explanation quality has at least four dimensions. Fidelity asks whether the explanation corresponds to the implemented rule. Completeness asks whether the important inputs and limits are represented. Comprehensibility asks whether the intended audience can understand it. Actionability asks whether the explanation supports a reasonable next step. The deterministic rule and stored pack metadata support fidelity. Missing-context questions support completeness. The card hierarchy aims at comprehensibility. Suggested review questions support actionability. Only the first two dimensions are strongly evaluated in the present work; comprehension and actionability need participant evidence.

2.4.3 Accessibility as an Evidence Requirement

Accessibility affects whether transparency is available in practice. WCAG 2.2 requires that information is not conveyed through colour alone and includes criteria concerning contrast and target size [22]. Android guidance recommends meaningful content descriptions and interaction targets that are large enough to operate reliably [23]. The implementation pairs icons and words with colour and supplies accessibility labels for primary controls.

These measures are design evidence, not conformance certification. True accessibility evaluation should include screen-reader order, announcements after asynchronous review, large-font reflow, switch access, keyboard focus, high-contrast settings, reduced motion, colour-vision deficiencies, and comprehension with users who rely on assistive technology. The distinction mirrors the wider thesis: source inspection can establish that a label exists, whereas a user study is needed to establish that the resulting interaction is effective.

2.5 Replaceable Rule Systems

Hard-coding statutory names into a decision engine creates maintenance and validity risks. A legal text may change, another jurisdiction may use different concepts, and a translated label may not preserve the same meaning. The appropriate software pattern is a stable engine contract with versioned, independently reviewable rule data. A pack should declare its jurisdiction, version, official source, caveat, supported signals, thresholds, references, and explanatory text. The engine should reject unknown packs and preserve the selected pack identifier in every finding.

Replaceability does not mean equivalence. Two packs cannot be treated as interchangeable simply because they implement the same TypeScript interface. Each requires independent

legal review, localisation, validation fixtures, and governance over updates. The included “Research Baseline” pack demonstrates switching without representing a second jurisdictional legal interpretation.

2.5.1 Pack Lifecycle and Review Governance

A credible pack lifecycle begins before implementation. The maintainer identifies the intended jurisdiction, material and territorial scope, regulated actors, effective date, official consolidated source, language versions, and the technical signals actually available. A qualified reviewer then maps signals to questions, not verdicts, and records assumptions and exclusions. Engineering converts the reviewed mapping into immutable data and fixtures. Localisation is reviewed separately because legally important distinctions can be lost even when the code remains valid.

Release governance should treat a pack as a versioned dependency. A substantive legal amendment, authoritative interpretation, corrected translation, changed threshold, or altered missing-context question may require a new pack version. Findings retain the version used at evaluation time. Updating the active pack should not rewrite historical records. If a migration is needed, it should create a new evaluation linked to the old one rather than silently changing the original rationale.

Pack selection also raises a scope question. Device location alone is not a reliable selector because applicable law can depend on establishment, targeting, residence, processing context, contracts, and sector. Revision 14 therefore requires explicit user selection from registered packs and presents the jurisdiction label and caveat. Automatic selection could be added only with a defensible scope model and an override.

2.5.2 Rule-Based and Learned Decision Models

Rule-based systems are attractive for a legal-support interface because a reviewer can reconstruct the condition that fired. A fixed threshold also has obvious limitations: it can be bypassed by splitting activity across windows, may not represent purpose, and may treat different application categories alike. Learned anomaly detection can adapt to patterns, but it introduces training-data bias, distribution shift, opacity, and uncertainty. A hybrid system could combine transparent guardrails with contextual or personalised models, provided that model provenance and uncertainty remain visible.

The present engine intentionally uses explicit research thresholds. The contribution lies in safe admission, temporal isolation, versioning, explanation, and boundary disclosure rather than in claiming optimal thresholds. Population calibration would require representative applications, ground-truth definitions, stratification by purpose and category, independent annotation, and validation on data not used to choose parameters. Without that evidence, a mathematically sophisticated adaptive threshold could be less trustworthy than a simple rule whose limits are stated.

2.5.3 Automated Validation and Oracle Independence

Randomised tests can explore more input combinations than a short hand-written suite, but their evidential value depends on the generator and oracle. Uniform random counts may seldom exercise exact thresholds. A generator that only creates well-typed values cannot test the runtime boundary. An oracle copied from the production helper may reproduce the same defect. The evaluation therefore combines fixed seeds, boundary values, malformed types, permission mutation, temporal constructions, and an explicit expected-result calculation.

Fuzzing and property-oriented testing are especially useful for invariants. Counts should remain finite and non-negative. Results should not depend on the order of unrelated package histories. Replaying an identical window should not increase totals. Overlapping aggregate windows should not be added. A simulator source should never become a native-observation source. Unknown packs should not load. These properties express safety behaviour more clearly than a single accuracy number.

Statistical language must also match the sample. A deterministic 1,000-case run provides strong regression coverage for its generator, but it is not a random sample of real applications or alleged infringements. Confidence intervals applied to such a run would describe the generated distribution, not the mobile ecosystem. Revision 14 reports the confusion matrix descriptively and reserves population claims for future independent datasets.

2.6 Comparative Positioning Against Prior Android Studies

Table 2 compares the unit of evidence and unanswered question in the closest empirical and platform work. The comparison is not a league table: the cited studies provide stronger behavioural evidence than Privacy Lens, while Privacy Lens contributes an artefact-level integration and traceability design that those studies were not intended to evaluate.

Table 2: Comparative positioning of Android privacy evidence and Privacy Lens

Work	Unit and method	Strongest supported contribution	Question outside its evaluated scope
Wijesekera et al. [4]	contextual-integrity field study; 36 participants; 27 million data points	demonstrates mismatch between contextual expectations and observed permission access	versioned legal prompts, missing-fact gates, and end-to-end finding provenance
Cao et al. [5]	survey of 1,719 participants across 10 countries/regions	quantifies cross-context variation in Android permission expectations and behaviour	local artefact admission, replay control, pack identity, and release traceability
Prange et al. [6]	field deployment with 132 participants	evaluates awareness and control around privacy-relevant smartphone features	legal-inference boundaries and replaceable jurisdiction mappings
Android AppOps [20]	official platform API and authority model	specifies observable operations and access restrictions, including own-UID limits without privileged authority	human interpretation, lawful-purpose facts, and calibrated user-facing claims
Privacy Lens	design-science artefact, deterministic tests, package evidence, and one-device walkthrough	preserves source, time, pack version, missing facts, and prohibited inference across a review pipeline	population behaviour, participant comprehension, legal approval, and broad device validity

The comparison prevents two symmetrical overclaims. Privacy Lens cannot inherit participant validity from studies whose interfaces, samples, and contexts differ. Conversely, the absence of a participant study does not erase its separately testable architecture contribution. The novel claim is therefore the explicit composition of evidence admission, replaceable rule identity, missing-context preservation, bounded presentation, and layer-matched evaluation—not discovery that permissions are contextual or that users need understandable controls.

2.7 Research Gap: From Observation to Calibrated Review

The reviewed approaches are individually strong at different layers. Manifest and static-analysis tools can identify declared or potential behaviour. Runtime and network monitors can observe selected executions. Platform dashboards can expose system-mediated histories. Enterprise or modified-platform systems can enforce policy. Consent and accessibility research can improve the presentation of user choices. None of these strengths, however, removes the need to state what an observation means, how far it reaches, and what it cannot establish.

The first unresolved issue is an *authority gap*. A screen can display a precise count while hiding whether the value came from a system service, an imported file, an instrumented experiment, or a simulator. The same number has different evidential weight under each source. Platform security further means that an ordinary application may observe only a subset of relevant operations. Prior technical work demonstrates the value of richer analysis, but also shows that coverage depends on authority, instrumentation, and execution conditions [18, 19]. A review interface therefore needs source and capability disclosure as part of the result, not solely in developer documentation.

The second issue is a *semantic gap*. Regulatory principles concern processing in context, whereas permission telemetry describes a platform operation. Frequency may help prioritise review, but it does not reveal purpose, lawful basis, necessity, recipient, retention, or applicable exception. A tool that replaces these missing facts with a legal label is not simply incomplete; it changes the proposition being evaluated. The literature on contextual integrity and data protection by design supports a context-sensitive interpretation, yet does not provide an automatic mapping from Android counts to legal outcomes [1–3].

The third issue is an *interaction gap*. Dense disclosure can be formally complete while remaining unusable. Cognitive-load and dual-process accounts explain why users may rely on colour, default position, or a short label even when detailed text is available [24, 25]. Accessibility guidance further shows that information is not meaningfully available if it depends on colour, small targets, or unlabelled controls [22, 23]. A privacy-review interface must therefore support fast orientation without allowing the fast path to imply more certainty than the detail supports.

The fourth issue is a *maintenance gap*. Legal references, product copy, thresholds, and test fixtures often evolve in different files and at different speeds. Hard-coded jurisdiction names can survive after a rule changes; historical findings can be reinterpreted by a new active configuration; translated labels can lose distinctions from the reviewed source. A versioned pack boundary cannot guarantee legal correctness, but it creates a reviewable unit with identity, source, caveat, fixtures, and lifecycle.

The final issue is an *evaluation gap*. Software projects frequently aggregate heterogeneous successes into one readiness claim. Unit tests, a generated AAB, a device screenshot, and a legal citation are all useful, but they answer different questions. Conformance with a threshold oracle is not acquisition recall. Installation is not long-run reliability. A polished warning is not comprehension. A citation is not independent legal approval. An evidence-bounded evaluation must preserve these denominators and scopes.

Privacy Lens addresses the intersection of these five gaps. Its contribution is compositional rather than substitutive: it does not replace platform telemetry, legal review, or participant research. It provides a contract through which a bounded observation can be admitted, evaluated, stored, and presented without losing source, time, rule identity, or missing context. The research problem is therefore not merely to detect a high count. It is to preserve epistemic discipline across the complete path from observation to user-facing claim.

2.8 Conceptual Model: Evidence, Rule, Claim, and Action

The literature supports a four-part conceptual model. *Evidence* describes what was observed, by whom, over which interval, and with what contextual fields. *Rule* describes the deterministic condition and regulation-pack version applied to admitted evidence. *Claim* describes the bounded status that the interface may present. *Action* describes the next human step, such as obtaining purpose information, checking a notice, repeating collection under an authorised source, or clearing local records.

Figure 1 makes the permitted direction and the missing-context gate visible. A legal verdict is not a fifth stage: it sits outside the automated path because the admitted record lacks the complete facts and authority required for adjudication.




Figure 1: Evidence-to-action conceptual boundary

The model is directional. Evidence constrains which rule can run; the rule constrains the claim; and the claim constrains the action that can be justified. The interface must not reverse this logic by beginning with a desired compliance label and searching for a supporting count. Nor should it allow an action such as reporting an application to imply that the underlying legal question has already been resolved.

Uncertainty enters at each transition. Evidence may be incomplete, a rule may be a research baseline, a claim may omit legally decisive context, and an action may require

professional judgement. Rather than compressing these uncertainties into one confidence score, Privacy Lens retains categorical source, explicit missing fields, rule metadata, and textual boundaries. This choice follows the reviewed literature’s emphasis on contextual interpretation and calibrated reliance while remaining implementable in a deterministic prototype.

The model also explains the division of evaluation. Admission and replay tests concern evidence. Threshold and pack fixtures concern rules. Projection and screenshot checks concern claims. Participant and professional-review studies would concern actions and reliance. Keeping these layers separate allows the thesis to answer its engineering questions without presenting unevaluated human or legal outcomes as completed results.

2.9 Store Delivery as a Trust Boundary

Application-store readiness includes more than a successful APK. Google Play uses Android App Bundles for new application delivery, requires current target API levels, and requires developers to complete a Data safety declaration [27–29]. Signing identity, a public privacy-policy URL, support contact, screenshots, listing text, and console review remain owner-controlled steps. A research build can be an initial submission candidate while still being unpublishable until these external obligations are completed.

2.10 Synthesis

The literature yields a precise design requirement: a privacy-review system must preserve the boundary between an observation and the authority of the claim made from it. Android research defines the acquisition constraint; regulatory sources define the missing legal context; HCI and accessibility sources define the communication constraint; and rule-system research defines the maintenance constraint. The following chapter converts this model into explicit requirements, design alternatives, and evidence contracts.

3 Requirements and System Design

This chapter translates the research gap into a design that can be implemented and evaluated. It first defines the product position and release requirements, then compares rejected alternatives, traces the decision lifecycle, specifies the evidence and rule-pack contracts, and maps each design invariant to an acceptance source. This order follows the central dependency of the system: a user-facing claim is valid only when the evidence, rule, and state transitions that produced it remain reconstructable.

3.1 Design Objectives and Requirements

Privacy Lens is designed as a local review instrument. Its output is a structured prompt for investigation, not a legal decision and not a security-enforcement action. The system accepts bounded permission summaries, validates their shape and temporal relationships, applies the

active regulation pack, stores the resulting findings locally, and presents the result with its provenance and limitations.

Table 3 defines the requirements used for implementation and evaluation. The table is intentionally concise; later subsections explain why each requirement is necessary and which failure it prevents.

Table 3: Revision 16 requirements and acceptance evidence

ID	Requirement	Evidence
F1	Evaluate supported permission summaries deterministically.	Seeded and boundary suites
F2	Preserve source, time window, pack, rule, and missing context.	Type and repository tests
F3	Reject unsafe identifiers, numbers, sources, timestamps, windows, and contexts.	Adversarial tests
F4	Switch only among registered, versioned regulation packs.	Pack registry tests and Settings UI
F5	Separate synthetic demonstrations from device-visible evidence.	Source labels and simulator tests
F6	Reject packs whose governance lifecycle, source, reviewer scope, or release scope is inconsistent.	Registry-load and mutation tests
F7	Require evidence specific to the selected Article 6 basis before returning no technical concern.	Basis-matrix tests
F8	Preserve binding-law, final-guidance, and consultation-material source status without conflation.	Source-governance tests
U1	Provide overview, findings, and settings through stable navigation.	Physical screenshots and UI trees
U2	Communicate severity without colour alone and disclose evidence limits.	Source inspection and physical rendering
U3	Present source-authority status without visually conflating consultation material with law or final guidance.	Settings source register and UI contracts
U4	Preserve heading context when the ledger changes state or evidence detail expands.	Empty/populated and disclosure device paths
P1	Keep findings local, disable Android backup and remote updates, and remove application network permissions.	Privacy tests and merged manifest
D1	Build a versioned release APK and AAB with a stable package identity.	Release build and device install

3.2 Design Alternatives and Threat Model

The model includes malformed bridge values, forged simulator provenance, unsafe numbers, replayed windows, overlapping summaries, unbounded history, misleading empty results, and user over-reliance on authoritative language. It does not defend against a compromised

operating system, a maliciously modified application binary, falsified system telemetry, or a legally incorrect pack supplied by a trusted maintainer. Those risks require platform integrity, signing governance, and independent legal review.

3.2.1 Selection Rationale

Several alternative product positions were considered. The first was a binary compliance scanner that would display compliant or non-compliant for every application. This was rejected because permission frequency cannot establish purpose, necessity, lawful basis, controller obligations, or applicable jurisdiction. A binary label would hide missing evidence and encourage users to treat a technical heuristic as a legal decision. The adopted model separates potential conflict, ordinary review, evidence gap, and no-technical-concern states; the last remains explicitly limited. The legacy “likely non-compliant” value is accepted only for migration and is projected as a potential conflict rather than generated by new evaluation.

The second alternative was a single risk score. A score appears compact and supports sorting, but it combines unlike uncertainties. A high aggregate count, missing purpose, synthetic source, and consent contradiction do not have a common natural unit. Weighting them would introduce assumptions that are difficult to justify and easy to overlook. Privacy Lens instead stores contributing signals and renders a reasoned state. A future prioritisation score could be added for professional triage only if every contribution remains inspectable and calibration evidence exists.

The third alternative was a learned classifier trained on application labels or permission histories. Machine learning could model interactions not captured by fixed thresholds, yet the project lacked an independently labelled, representative, legally meaningful data set. Training on store categories, malware labels, or developer declarations would answer different questions. A deterministic rule engine was selected because it makes boundary behaviour, missing inputs, and change history reviewable. This choice does not claim that rules are universally superior; it matches the evidence available for the thesis.

The fourth alternative was to embed GDPR language in interface components. This is initially convenient because each screen can display the relevant article directly. It creates duplication, complicates review, and makes regional extension risky: changing jurisdiction would require editing presentation and perhaps persistence logic throughout the application. The pack contract centralises legally motivated content and forces the engine to record which pack generated a finding. Generic screens display pack-provided labels through validated fields.

The fifth alternative was to choose a region automatically from device locale or location. Locale does not reliably determine applicable law, organisational establishment, data-subject status, or processing context. Automatic selection could therefore present the wrong framework with unwarranted confidence. The adopted design requires explicit user selection from registered packs and shows the selected jurisdiction and caveat. It does not collect location to decide which privacy law to discuss.

The sixth alternative was a cloud account with central evidence storage. Centralisation could simplify synchronisation, remote dashboards, and model updates, but it would also

create a new repository of sensitive behavioural metadata and require authentication, server security, retention, international transfer analysis, and deletion workflows. The current product is local first, uses no account, and offers local deletion. This choice reduces deployment features but aligns the application’s conduct with its privacy objective.

The seventh alternative was to ingest unrestricted free-form pack files. This would maximise extensibility but expose users to misleading sources, malformed thresholds, pack confusion, and potentially executable content. The present registry admits bundled immutable definitions. Future signed declarative distribution may be appropriate, but only after schema, publisher, rollback, and review controls exist.

The eighth alternative was to treat absence of observations as zero activity. This would produce a cleaner dashboard, but it would conflate a real zero with unavailable acquisition. The selected model requires a capability-aware source and reports an evidence gap when coverage is missing. This is less reassuring and more accurate.

The ninth alternative concerned notifications. Emitting an alert for every evaluation would demonstrate activity but create repetition and fatigue. Repository logic instead considers whether a finding is new, has changed state, or has escalated. Silent storage remains possible for unchanged or low-priority outcomes. The design can later support controlled comparison of notification policies without changing decision semantics.

The final alternative concerned storage of raw observations. Keeping a complete event history would improve forensic detail but increase sensitivity and resource cost. The candidate stores bounded findings and aggregate history needed for current temporal reasoning. Event-level acquisition is deferred until an authority model, minimisation policy, and retention design can be evaluated. This is a deliberate privacy-performance trade-off rather than an accidental omission.

3.3 Architecture and Decision Lifecycle

The selected architecture separates acquisition, admission, temporal preparation, evaluation, persistence, projection, and presentation. Rather than allowing each screen to interpret raw Android data independently, every stage receives a bounded input and produces a state whose meaning is explicit. The following lifecycle, failure, minimisation, and release-gate decisions define how this separation behaves under both normal and adverse conditions.

Figure 2 shows the single decision path and the two governance inputs that constrain it. The regulation pack does not receive UI state, and the interface does not bypass the finding record to reinterpret raw evidence. This separation makes source identity and pack identity available at every downstream review point.




Figure 2: Evidence-bounded system architecture and decision path

3.3.1 End-to-End Decision Lifecycle

The lifecycle begins when the user initiates review, a scheduled worker requests work, or the controlled simulator generates an example. The acquisition component returns candidate records together with a source identity. It must not return a native zero merely because a platform capability is unavailable. Capability failure and an observed zero are different states and lead to different user decisions.

Each candidate crosses the admission boundary as unknown data. The validator checks the package identifier, permission key, count, window, timestamps, source, and optional context without relying on compile-time types. Rejection is structured. The record does not enter temporal history and cannot contribute to a reassuring status. The caller can present the failure as an evidence problem and retain enough diagnostic detail to correct the producer without exposing raw sensitive content unnecessarily.

An admitted record is placed in a key space defined by package and permission. Exact-window replay replaces the prior aggregate. Independent completed windows may be retained for cross-window analysis; overlapping aggregates are not summed. Simulator data remains excluded from accumulated cross-window claims. Retention and a per-key cap bound state. These operations prepare evidence before jurisdictional rules are applied.

The engine receives the selected immutable pack explicitly. It locates the supported rule, calculates daily excess, evaluates burst information where timestamps are available, examines compatible history, and inspects any admitted processing context. The result includes the observed and threshold values, signal types, status, communication priority, rationale, legal or research reference, missing-context questions, and pack identity. The calculation never mutates the pack.

The repository compares the result with the stored finding for the same logical subject. It distinguishes an unchanged state, a new concern, resolution or reversal, and risk escalation. This comparison governs notification eligibility. It also writes a bounded slice under the versioned repository key. A finding remains attached to its original pack even if Settings later

selects another pack.

Projection turns repository records into interface-facing models. It calculates overview totals by state, applies Findings filters, formats source and temporal labels, and exposes actions. Projection should not infer new legal semantics. If a required field is absent, it retains an unknown or evidence-gap presentation. This makes projection testable independently of React components.

Presentation then orders information according to user need. The overview states reach and limitation before the main action. Findings allow scanning by state and inspection of reasons. Settings places pack choice, caveat, simulator, privacy information, and deletion controls in separate groups. Accessibility labels name the action and state rather than depending on icon or colour. Destructive action requires deliberate confirmation where implemented.

The lifecycle ends with user action or explicit non-action. A user may seek a privacy notice, ask a developer for purpose and retention information, repeat review with authorised evidence, or decide that the technical signal does not require immediate attention. Privacy Lens does not contact the reviewed application, change its permissions, or submit a complaint automatically. Keeping the final decision outside the engine preserves human accountability.

Deletion is also part of the lifecycle. The user can clear local findings, and uninstall or platform data clearing removes the application data within Android's boundary because backup is disabled. Export or future synchronisation would create additional copies and must define deletion separately. The current lifecycle intentionally ends at local storage.

3.3.2 Failure-State Design

Failure states were designed as product states rather than exception messages. Acquisition unavailable means the current deployment cannot obtain suitable records; it is not displayed as a completed scan. Invalid evidence means a producer supplied data that fails the contract; it cannot be silently repaired into a normal case. Unknown pack means evaluation is blocked until a registered definition is selected; it does not fall back to the EU pack merely because that pack is bundled.

Storage failure also requires conservative behaviour. If persisted data cannot be parsed or migrated safely, the interface should not mix it with current findings. Recoverable records may be retained after field-level validation, while unrecoverable content should be quarantined or omitted with a visible explanation. Clearing data remains available so the user can recover control. Production logging should report the failure class without printing evidence payloads.

Work scheduling failure is distinct from review failure. WorkManager can defer periodic work because of platform policy, Doze, reboot, or vendor restrictions. The interface should present the last completed review time and source capability rather than imply continuous monitoring. A registered request demonstrates scheduling intent, not on-time execution. Future long-run tests should measure delay and recovery before stronger wording is used.

Partial evaluation is another possibility. Some records may be valid while others are rejected. Treating the batch as wholly clean would hide rejection; treating it as wholly failed would discard useful evidence. A suitable summary reports accepted findings together with

the count and reasons for excluded records. The current bounded flows establish structured rejection at record level; richer batch disclosure remains an extension point.

Pack-content failure has two forms. Schema failure is machine detectable and should block admission. Substantive legal error may only be identified by expert review or later correction. Pack provenance, versions, official sources, and historical identity allow the maintainer to issue a successor without pretending earlier findings never existed. The interface can mark a deprecated pack and recommend re-evaluation while preserving the original record.

Network failure is intentionally limited because core review and bundled packs operate locally. Privacy-policy and official-source links may be unavailable, but the application should retain the relevant title and caveat. If remote pack delivery or synchronisation is later added, offline operation and last-known-good behaviour become acceptance requirements rather than optional resilience features.

Visual failure includes clipped localisation, unreadable contrast, focus loss, hidden caveats, and controls obscured by system insets or keyboards. These failures may not appear in logic tests. Rendered screenshot matrices, hierarchy inspection, screen-reader use, font scaling, and device variation are therefore part of the planned evidence. The one-device acceptance check establishes only the candidate baseline.

Finally, misunderstanding is treated as a failure even when the software behaves as coded. If users interpret no technical concern as GDPR compliance, or mistake simulator records for device evidence, the product has failed its communication objective. Participant comprehension and reliance tests are needed to detect that class. The design reduces the risk through terminology and provenance, but does not declare it eliminated.

3.3.3 Data-Minimisation Decisions

The application needs enough evidence to explain a review signal, yet every retained field has privacy cost. Package identifier, permission family, count, time window, source, pack identity, and decision reason form the current minimum for reconstructing an aggregate finding. Optional context is admitted only when supplied for the review purpose. The application does not require account identity, contact lists, precise location content, microphone content, or advertising identifiers.

Counts are stored instead of raw sensor values. A location access finding does not require coordinates, and a microphone access finding does not require audio. This distinction is central: the system reviews evidence that an operation occurred, not the personal data produced by the operation. A future connector should preserve that separation unless a separately approved research question requires content.

Time resolution is also minimised. Aggregate start and end values support window validity, replay, and bounded temporal analysis. Timestamp arrays are accepted for controlled burst evidence, but they should not be collected by default merely because the schema can represent them. The source capability and purpose determine whether additional resolution is justified.

Local-first storage avoids creating a server-side behavioural database. The trade-off is that the user or controlled deployment is responsible for device access and loss, and professional

collaboration features are limited. The design disables Android backup so evidence is not restored unexpectedly to another device through ordinary platform backup. This choice should be re-evaluated if enterprise-managed backup becomes a stated requirement.

Bounded history limits denial-of-service and retention exposure. The cap is not a substitute for purpose-based deletion, but it prevents indefinite growth under current aggregate use. Repository slices and clear controls make the lifecycle understandable. Future event-level storage would require stronger indexing, encryption assessment, retention schedules, and selective deletion.

Notifications contain only the minimum text needed to indicate a changed review state. Detailed evidence remains inside the application. Lock-screen visibility can still expose the existence of a concern, so future production settings should consider notification privacy and user preference. The current evaluation does not claim that every device lock-screen configuration was tested.

External links are used for official materials and privacy documentation instead of embedding broad web content. The application does not send the finding to those sites. Users should be informed that opening a link transfers them to another controller and may reveal ordinary browser metadata. Remote analytics or link tracking would contradict the minimisation posture unless separately justified.

Data minimisation finally applies to evaluation artefacts. Screenshots, UI dumps, log excerpts, and device metadata should avoid unrelated notifications, account names, and installed-application lists. Synthetic demonstrations are preferred for public evidence. Physical records used for thesis acceptance remain scoped to the test application and bounded device session.

3.3.4 Requirement Priorities and Release Gates

Requirements were prioritised according to the harm caused by failure. Evidence admission, provenance, non-verdict language, source separation, replay safety, pack identity, and deletion are release blockers because failure could create misleading privacy conclusions or remove user control. Core navigation, readable hierarchy, package identity, release builds, and absence of sensitive runtime permissions are also blockers for the initial store candidate.

Important but non-blocking research extensions include multi-OEM acquisition measurement, long-duration WorkManager observation, full screen-reader matrices, participant studies, production signing, and independent legal review. They are non-blocking only for the bounded engineering candidate. They become blockers for stronger claims such as public production readiness, accessibility conformance, or validated GDPR guidance.

Optional features include cloud synchronisation, accounts, remote pack updates, application blocking, and learned ranking. These features do not repair current evidence gaps automatically and would introduce substantial security, privacy, or governance work. They should be considered only after the core assurance chain is stable.

The release gate follows dependency order. A source-level failure blocks package evaluation because the binary would not represent an accepted program. A package or manifest failure

blocks device acceptance. A device crash or broken primary journey blocks visual promotion. A missing legal or usability study does not turn an engineering check red; it keeps the associated substantive claim explicitly unverified.

This priority model avoids two extremes. It does not delay all prototype learning until every production activity is complete, and it does not allow a successful demonstration to erase unresolved authority or legal questions. Each gate authorises a named next activity. The Revision 15 source candidate authorises release rebuilding and controlled evaluation; it does not inherit predecessor device acceptance automatically.

3.4 Stakeholders and Functional Requirements

The same evidence contract serves different users, but their decisions and information needs are not identical. Requirements are therefore derived from explicit scenarios before they are divided into functional and non-functional categories.

3.4.1 Stakeholders and Usage Scenarios

The primary user is a person who wants to understand available privacy evidence without becoming an Android or GDPR specialist. This user needs a short account of what was observed, why it deserves attention, what is not known, and what action is reasonable. A second stakeholder is a developer or product owner using the application during internal review. This stakeholder needs reproducible rule identifiers, pack versions, timestamps, and evidence-source labels. A third stakeholder is a privacy or legal reviewer who may use the finding as an entry point into a broader assessment. This stakeholder needs the official source, missing-context questions, and an explicit statement that the application has not reached a legal conclusion.

The first usage scenario is an ordinary review with no prior data. The overview explains the scope before offering a review action. If the bridge returns no admissible observations, the interface reports an evidence gap rather than a clean bill of health. The second scenario is a populated review. Findings are grouped by status and can be inspected individually. The third scenario is learning through synthetic data. The user runs a controlled demonstration, which creates a varied but deterministic set of records marked as synthetic. The fourth scenario is changing legal perspective. The user selects another registered pack, sees its jurisdiction and caveat, and performs a new review without relabelling historical findings. The fifth scenario is withdrawal from the tool itself: the user clears local findings or uninstalls the application.

These scenarios exclude unsupported uses. Privacy Lens is not intended to block another application, intercept network traffic, prove that a sensor callback stopped, obtain evidence by evading platform restrictions, or submit a legal complaint automatically. It does not select applicable law from device location. Defining exclusions is important because a polished interface could otherwise suggest a broader product than the evidence pipeline implements.

3.4.2 Functional Decomposition

The review workflow is decomposed into acquisition, admission, temporal preparation, evaluation, persistence, projection, and presentation. Acquisition produces a candidate summary and a capability statement. Admission checks whether the record can be trusted as a well-formed input. Temporal preparation determines whether history can be combined without double counting. Evaluation applies one immutable pack. Persistence writes findings and migration metadata. Projection converts stored findings into dashboard counts and sections. Presentation renders this projection with actions and caveats.

This decomposition allows each failure to be represented at the correct layer. If the bridge is unavailable, the application should not create a zero-count native record. If a timestamp is unsafe, the engine should not try to repair it silently. If two aggregate windows overlap, the repository should not sum them. If a pack identifier is unknown, the registry should not fall back to whichever pack appears first. If storage migration fails, the interface should avoid presenting stale records as current evidence. The design favours explicit absence over invented continuity.

The simulator follows the same downstream path after acquisition. It is not a separate UI mock that bypasses validation. This is important because a demonstration that constructs presentation objects directly would not test the engine, repository, or projection. Synthetic source identity is introduced at acquisition and persists through every later stage.

3.4.3 Non-Functional Requirements

Security requirements include fail-closed input handling, conservative identifier formats, bounded history, source enumeration, immutable pack definitions, disabled backup, and no sensitive runtime permissions. Privacy requirements include local processing, no remote evidence service, no analytics or advertising SDK, and a clear deletion control. Reliability requirements include deterministic results for identical admitted inputs, replay replacement, pack-specific persistence, and migration from the previous repository schema.

Usability requirements include a clear initial action, a stable three-destination navigation model, non-colour status semantics, concise summaries, optional detail, meaningful accessibility labels, and separation of destructive controls. Maintainability requirements include one authoritative compliance path, typed interfaces, registry-based packs, reusable theme constants, tests for new packs, and removal of duplicate legacy services. Deployability requirements include a stable application identifier, explicit version, current target SDK, release APK and AAB generation, upload-key configuration, privacy-policy content, and documented owner actions.

Performance is treated as a constraint rather than a validated claim. The decision engine operates on bounded in-memory arrays and simple rule lookup. Storage is local and capped per package-permission key. WorkManager is selected for deferrable periodic work rather than exact scheduling. No quantitative battery, long-duration memory, or startup objective is claimed in Revision 16 because the available evidence does not cover those dimensions broadly enough.

3.5 Evidence and Regulation-Pack Design

This section specifies the two contracts at the centre of the framework. The evidence contract defines what the engine may know; the regulation-pack contract defines which reviewed rule data may be applied. Admission, temporal isolation, persistence, and abuse controls preserve the boundary between them.

Table 4 prevents the pack’s legal references from being mistaken for derivation of its numerical thresholds. The cited provisions motivate review questions. They do not contain daily or per-minute permission limits, and the application cannot infer the legally decisive facts listed in the fourth column.

Table 4: EU GDPR pack: legal motivation and prohibited inference

Legal anchor	Observed input	Permitted engineering prompt	Facts still required	Prohibited conclusion
Art. 5(1)(b), purpose limitation	Contacts count and time window	Ask whether access serves an explicit and compatible purpose	Purpose specification, compatibility analysis, recipients, user expectation	Contacts access breached purpose limitation
Art. 5(1)(c), data minimisation	Location, microphone, or contacts frequency	Ask whether the observed pattern was necessary and proportionate to the stated function	Function, alternatives, precision, duration, retention, data actually produced	A threshold exceedance was unlawful or excessive
Art. 6 and, where applicable, Art. 9	Declared lawful-basis and special-category context fields	Expose a missing basis or Article 9 condition for human review	Controller assessment, validity of basis, data category, exceptions, withdrawal and objection facts	Processing lacked a valid legal basis
Arts. 7, 13 and 14	Consent record and notice references	Ask for consent evidence and applicable transparency record	Validity, timing, information supplied, withdrawal, and data provenance	Consent or transparency was legally valid
Art. 35	DPIA screening flag and assessment reference	Block evidence-complete status when a positive screening lacks the DPIA record	Risk likelihood, scale, systematic monitoring, vulnerable people, safeguards, consultation	A DPIA was or was not legally required
Art. 5(2) and Art. 25	Source, pack version, rationale, missing fields, local controls	Preserve a reproducible prompt and evidence boundary	Wider technical and organisational measures, risk, state of the art, records, governance	The controller demonstrated compliance or the App is certified

The legal traceability matrix is deliberately asymmetric. Technical evidence may trigger a question, but a legal proposition is not allowed to flow backwards into the evidence record as if it had been observed. This follows the GDPR and final EDPB guidance while retaining the privacy-engineering distinction between a data action, its context, and the problems it may create [1, 2, 11].

3.5.1 Evidence Schema and Provenance

The evidence schema is designed around reconstructability. A permission summary includes a package identifier, supported permission key, non-negative safe-integer count, start and end

of the aggregate window, source enumeration, and optional processing context. The source is one of the controlled values supported by the application, such as native bridge, imported evidence, or simulator. A free-form source string is rejected because it could allow synthetic data to impersonate device evidence.

Context fields represent what a legal review would need but a permission count cannot provide. Revision 15 includes purpose, controller identity, lawful basis, basis-specific supporting reference, transparency notice, data-minimisation assessment, retention period and justification, Article 9 condition, and DPIA screening outcome. Their absence does not make the technical observation invalid. It changes the finding by adding missing-context questions and preventing an evidence-complete state. This distinction separates data-quality failure from legal incompleteness; a reference proves only that a record was supplied, not that its legal reasoning is correct.

A finding contains the admitted evidence plus decision metadata: status, risk, rule identifier, threshold, rationale, regulation-pack identifier and version, official source, evaluation time, and missing-context fields. The schema is intentionally verbose. Storage efficiency is less important than the ability to explain why a result appeared and which assumptions applied.

Provenance also constrains metrics. Simulator precision and recall are computed only for records whose expected labels are part of the controlled generator. Imported or native records do not receive an accuracy label because their legal and behavioural ground truth is unknown. Aggregating these sources into one accuracy metric would create a number with no coherent denominator.

3.5.2 Admission and Fail-Closed Validation

Validation occurs at runtime even though the application is written in TypeScript. Static types protect code paths compiled from known sources, but values crossing a native bridge, storage migration, import boundary, or untyped caller can violate those types. The validator therefore treats each field as unknown until checked.

Counts and timestamps must be finite JavaScript safe integers. This excludes negative values, fractional event counts, infinities, not-a-number values, and integers whose arithmetic cannot be represented exactly. The time window must have a strictly earlier start and a later end. An evaluation timestamp must be consistent with the window. Retention values must be positive and bounded. Package identifiers and permission keys must match conservative patterns rather than merely being non-empty strings.

Context validation follows the same rule. String fields must actually be strings before trimming; booleans must not be inferred from truthy text or numbers. The engine does not coerce a value such as the text `false` into a Boolean because this can reverse meaning. Unknown source and permission values produce an invalid or review-hold state instead of falling through to a default threshold.

Fail-closed does not mean the application crashes. It means the system returns a safe, explicit result that prevents incomplete evidence from being described as no concern. The UI can then explain which field was rejected or why review cannot proceed. This approach

supports both security and diagnosability.

3.5.3 Temporal Isolation Model

Aggregate windows create a double-counting problem. Suppose two records report counts for intervals that overlap. Without event identifiers, the system cannot know whether events in the shared interval appear in both counts. Summing them could manufacture a cross-window threshold. Revision 14 therefore aggregates only completed, non-overlapping windows for the same package and permission.

For windows $w_i = [s_i, e_i)$ and $w_j = [s_j, e_j)$, compatibility is defined as

$$e_i \leq s_j \quad \vee \quad e_j \leq s_i. \quad (1)$$

Adjacency is allowed because half-open intervals that meet at one boundary do not overlap. Exact replay is detected by equal package, permission, start, and end. The newer replay replaces the stored value rather than increasing the history size or total. This makes processing idempotent for a repeated aggregate export.

Retention uses a monotonic watermark derived from the latest admitted window end for the key. Basing expiry on a newly arrived record's timestamp alone would allow a late or forged old record to move the horizon backwards. Records older than the bounded horizon relative to the watermark are removed. Each package-permission history is also capped at a finite maximum. The cap is a denial-of-service safeguard, not a claim about the maximum legitimate event rate.

Cross-window detection is consequently conservative. It may fail to combine useful but overlapping aggregates, yet it avoids asserting a total that the evidence does not support. Event-level identifiers or disjoint platform buckets would permit a richer model. The current design states the limitation rather than guessing how much overlap occurred.

3.5.4 Rule-Pack Contract in Detail

The common pack contract separates descriptive metadata from executable rule data. Descriptive metadata includes identifier, display name, jurisdiction label, version, official source, effective or review date, and caveat. Rule data binds a supported permission signal to a research baseline, threshold logic, severity, rationale, reference label, and missing-context questions.

The registry is the only route from a persisted identifier to a pack object. This prevents arbitrary downloaded or user-entered text from being executed as a rule. Pack objects are immutable so an evaluation cannot observe a threshold changing midway through a run. The engine receives the pack as an explicit dependency, which makes test fixtures clear and avoids hidden global state.

Pack identity is copied into findings rather than represented only by a current settings value. If a user evaluates under the EU pack, selects the Research Baseline, and later opens an older finding, the finding still displays the EU pack and version. This historical stability is

essential for auditability.

The pack contract deliberately does not require every jurisdiction to supply an equivalent article number. A rule can use a provision, principle, official guidance reference, or non-legal research identifier. The UI renders a general reference label and source link. This avoids encoding the assumption that all legal systems share the GDPR’s structure.

3.5.5 State, Persistence, and Migration

The privacy context coordinates bridge capability, active pack, loading state, findings, demonstration status, and repository actions. State transitions are explicit: idle, reviewing, completed, and failed or unavailable. The interface disables duplicate actions while a review is running and provides a human-readable result when acquisition returns no evidence.

AsyncStorage is used for app-private persistence. The repository key is versioned. On load, older records are normalised into the current schema where safe. Fields that cannot be reconstructed are marked missing rather than invented. A migration should preserve source identity and original timestamps. If a record cannot satisfy current admission rules, it should be excluded from the active projection and surfaced as a migration issue where appropriate.

Clearing findings removes review records while preserving the application itself. Pack selection may be retained separately because it is a preference, but the UI states what the clear action affects. Android backup is disabled so uninstalling or clearing application data has the expected local-deletion meaning within the platform boundary.

3.5.6 Security and Abuse Cases

The first abuse case is source forgery: a caller marks synthetic input as native. Enumeration and bridge-controlled construction reduce this risk inside the application, although a modified binary is outside the threat model. The second is numeric abuse: extremely large or non-finite counts cause overflow or misleading totals. Safe-integer validation rejects them. The third is replay amplification. Exact-window replacement provides idempotence. The fourth is overlap amplification, addressed by temporal isolation.

The fifth case is history exhaustion. An attacker or faulty source supplies many distinct windows. Per-key caps and retention limit memory growth. The sixth is identifier pollution, where many malformed package or permission strings create unbounded partitions. Conservative formats and supported permission enumeration reject these values. The seventh is pack confusion, where an unknown identifier is presented under a trusted legal name. Registry-only loading and copied pack metadata prevent silent fallback.

The final abuse case is epistemic rather than computational: a user or stakeholder treats a finding as a legal determination. Interface caveats, neutral statuses, official-source links, and missing-context questions reduce this risk. They cannot eliminate it. Training, documentation, and qualified review remain necessary controls.

3.6 Traceability and Evaluation Design

The architecture is complete only when its requirements can be tested at the appropriate layer. This final section links requirements to evidence, identifies stable invariants, and records how future extensions should renew acceptance.

3.6.1 Traceability from Requirement to Evidence

Every requirement is paired with a planned evidence source. Deterministic evaluation is covered by seeded and exact-boundary fixtures. Provenance preservation is covered by type, repository, and presentation tests. Unsafe inputs are covered by adversarial runtime cases. Pack switching is covered by registry and identity assertions plus Settings rendering. Synthetic separation is covered by simulator-source tests and populated-device screenshots.

Navigation and visual hierarchy are evidenced by UI trees and physical screenshots, while accessibility claims remain limited to source and rendered cues. Local processing and permission minimisation are evidenced by dependency review, merged manifest, and installed-package inspection. Delivery is evidenced by successful release tasks, artefact hashes, package metadata, and installation. The matrix does not list legal compliance or broad user adoption because no test in the current protocol can establish them.

3.6.2 Evaluation Strategy

Evaluation is deliberately layered. Source-level checks establish type and lint conformance. Deterministic tests establish consistency with the published rule specification and adversarial rejection behaviour. Android builds establish package integration. Manifest and package inspection establish requested-permission facts. ADB interaction and screenshots establish bounded physical rendering and launch behaviour. Each result is reported at its own layer; no result is extrapolated into legal compliance or population effectiveness.

3.6.3 Design Invariants

The architecture is governed by invariants that should remain true even when screens, acquisition connectors, or legal packs change. First, no accepted finding exists without an admitted evidence source and bounded temporal scope. Second, no legal or research rule exists without a pack identifier and version. Third, unknown, malformed, or unavailable evidence cannot be translated into a reassuring status. Fourth, simulator provenance cannot be upgraded by persistence or presentation. Fifth, historical findings retain the rule identity used at evaluation time. Sixth, local deletion remains available without requiring an account or remote approval.

These invariants are more stable than component names. A future implementation could replace AsyncStorage with Room, React Native with another presentation framework, or bundled packs with signed declarative packages while preserving the same assurance properties. Conversely, a refactor that keeps every current class name but drops source identity or silent-fallback protection would violate the design. Stating invariants therefore makes architectural review less dependent on the incidental code structure of Revision 16.

The invariants also guide test selection. Admission tests challenge the first and third invariants. Pack-switch and persistence tests challenge the second and fifth. Simulator and cross-window tests challenge the fourth. Manifest, backup, and clear-data checks contribute to the sixth. This mapping helps maintainers decide which evidence to renew after a change. It also reveals gaps: authentication of an external producer, for example, is not yet an invariant guaranteed by the prototype and must not be implied by the controlled source label.

An invariant can fail through wording as well as computation. If a finding correctly stores evidence gap but the screen headlines the review as safe, the third invariant has failed at presentation. If an export omits the pack version, the fifth has failed outside the repository. End-to-end checks are therefore necessary even when unit tests are comprehensive.

Future release reviews should record each invariant as satisfied, failed, or not evidenced, with direct artefact references. A failed invariant blocks promotion. A not-evidenced item may permit a narrower research build, but its associated claim must remain absent. This practice keeps the product’s trust model explicit as functionality grows.

3.6.4 Decision Records for Future Extension

Future contributors should document material changes as concise decision records. A record should state the problem, evidence, alternatives, selected option, affected invariants, privacy consequences, pack consequences, migration, rollback, and verification required. This is especially important for changes that appear convenient but alter authority, such as adding remote telemetry, inferring jurisdiction automatically, importing an unreviewed rules file, or expanding Android privileges.

A decision record does not replace code review or legal review. It makes assumptions visible before they become distributed across implementation and copy. It also provides a durable explanation when later maintainers ask why the product refuses a seemingly simple shortcut. In a thesis prototype intended for store progression, this continuity matters because deployment work will occur after the original evaluation snapshot.

The first required record for a new acquisition connector should describe deployment authority and observed coverage. The first record for a new regulation pack should identify responsible reviewers and effective sources. The first record for remote services should include a data-flow and retention analysis. The first record for a new user-facing status should specify evidence semantics and comprehension tests. These prompts keep expansion aligned with the central evidence-bounded principle.

4 Implementation

This chapter explains how the Chapter 3 contracts were realised in the release candidate. Following the argument-first pattern established in the reference thesis, each implementation section begins with the design problem it resolves and then identifies the code path, state transition, or package fact that supplies the evidence. The chapter does not repeat the removed health-dashboard, experiment, or scoring features; those components are unrelated to the

current research questions and no longer form part of the product architecture.

The implementation has four cooperating domains. TypeScript supplies evidence types, validation, regulation packs, deterministic decision logic, and repository behaviour. Kotlin supplies Android capability reporting and deferrable workers. React Native renders the three-destination review journey. Gradle and the Android manifest define package identity, SDK targets, signing inputs, backup behaviour, and final permission declarations. The sections below follow this dependency order from rule data to delivery artefact.

4.1 Implementation Overview and Module Boundaries

The application uses Expo and React Native for the interface, TypeScript for the decision and state layers, AsyncStorage for app-private persistence, and Kotlin for the Android bridge and workers. Expo Router supplies route composition, while React Navigation primitives support the custom bottom navigation. AndroidX WorkManager represents deferrable background work. The release build uses the Android Gradle Plugin, Hermes JavaScript engine, and standard Android packaging tasks.

The stack was selected to separate concerns rather than to maximise the number of frameworks. TypeScript provides shared types for evidence, findings, and packs. Kotlin is limited to capabilities that require Android APIs or native scheduling. The interface does not import Kotlin details; it calls a typed service that reports capability and results. The compliance engine does not import React Native, storage, or Android modules. This makes deterministic tests runnable under Node after TypeScript compilation.

At the source-tree level, regulation definitions reside under the regulation directory, generic decision logic under compliance, application state under context, the bridge adapter under services, and reusable presentation components under a privacy-specific component directory. Theme constants are centralised. This organisation replaces earlier duplicate compliance services and prevents a screen from becoming the authoritative home of a rule.

Table 5 records concrete implementation witnesses for the architecture in Figure 2. The final column is important: a module name alone is weak evidence unless the reviewer can also identify the failure that the separation prevents.

Table 5: Design invariants and implementation witnesses

Invariant	Implementation witness	Failure prevented
Unknown input is not trusted	<code>parseAudit</code> and <code>evaluateSafe</code> in <code>RulePackComplianceEngine.ts</code>	Malformed bridge or imported data becoming a reassuring normal finding
Source identity survives evaluation	Audit source field, finding evidence, repository migration, and source chip	Simulator or imported evidence being mistaken for device observation
Regional logic is pack-owned	<code>RegulationPack</code> , closed registry, EU pack, and Research Baseline pack	UI, storage, or Android code silently embedding one jurisdiction’s terminology
Historical findings keep decision identity	Pack identifier, regulation label, rule reference, rationale, and immutable stored finding	A later pack selection silently reinterpreting an earlier result
The reviewer does not create a new remote data flow	App-private <code>AsyncStorage</code> , no account, analytics, advertising, or evidence upload, plus disabled backup	A privacy-review tool expanding the exposure it is intended to help examine
Release evidence is reproducible but bounded	Versioned package, environment-based signing inputs, APK/AAB hashes, merged manifest, and installed-package record	Treating a local build success as production signing, console approval, or broad device validity

4.2 Rule and Evidence Pipeline

The core implementation problem is to apply a jurisdiction-specific rule without allowing jurisdiction-specific assumptions to enter validation, persistence, or presentation code. The pipeline resolves this problem through an explicit pack dependency, an unknown-to-admitted runtime boundary, deterministic signal computation, a provenance-rich finding, and a repository that notifies only on meaningful change.

4.2.1 Regulation-Pack Implementation

The EU pack is an immutable TypeScript object identified as `EU_GDPR`. It declares the European Union and EEA as its jurisdiction label, cites Regulation (EU) 2016/679 as its version label, and links to the consolidated EUR-Lex source. The pack includes a legal caveat that the automated warning supports accountability review and does not determine infringement.

Three permission families are implemented. Location has a research baseline of 24, microphone a baseline of 8, and contacts a baseline of 4. Each baseline uses a multiplier of 1.5, producing prompt thresholds of 36, 12, and 6 respectively. These values are prototype parameters. They are not extracted from GDPR text, not population estimates, and not represented as legal limits. Their purpose is to exercise an explainable evaluation path and provide deterministic boundaries for testing.

The EU pack supplies a rationale and reference for every permission family. Location directs review toward necessity, proportionality, and documented lawful basis under Articles 5 and 6. Microphone review includes the possibility that special-category material may require Article 9 analysis. Contacts review emphasises purpose limitation, minimisation, and lawful basis. The reference is intentionally a prompt rather than an assertion that the cited provision applies to every record.

Missing-evidence logic is also pack-owned. The EU pack asks for specified purpose, Article 6 lawful basis, controller identity, retention period, and, when special-category processing is marked, an Article 9 condition. If processing relies on consent, consent has been withdrawn, and a positive access count remains, the pack returns its most urgent review status. If required context is missing, it returns insufficient evidence. Otherwise a threshold signal produces review required, while the absence of a signal produces no technical concern for the admitted check.

The Research Baseline pack uses higher prototype baselines of 36, 12, and 6 before applying the same multiplier. It uses region-neutral principles and does not require an Article 6 or Article 9 field. Its caveat explicitly states that it is not law or legal advice. The difference demonstrates that the engine can consume pack-owned rules and missing-evidence logic without embedding GDPR branches in generic code.

4.2.2 Pack Governance Lifecycle

Technical modularity is necessary for regional expansion but insufficient for trustworthy pack publication. Table 6 defines the governance states that should surround the code-level registry. A pack may enter the application only after its identity and technical fixtures are valid; it may be represented as legally reviewed only when a qualified reviewer and dated decision record exist. Supersession never rewrites findings created under an earlier pack.

Table 6: Regulation-pack governance lifecycle and release authority

State	Required decision record	Permitted product action
Draft	jurisdiction, scope, official sources, effective date, language, assumptions, exclusions, author	development only; cannot appear as an approved legal pack
Technical candidate	schema validation, immutable identifier/version, fixtures, caveat, source URL, migration and rollback plan	registered testing and controlled demonstration with candidate disclosure
Legally reviewed	reviewer identity and qualification, review date, provisions examined, contested mappings, expiry or review trigger	user selection with the reviewed scope and residual caveat; never an automated compliance verdict
Approved release	owner approval, signed artefact hash, localisation record, compatibility range, release notes	production distribution for the approved version and locale set
Superseded or revoked	successor or revocation reason, effective time, affected versions, correction and user-notice plan	historical findings remain readable under original identity; new evaluation uses an authorised successor or is blocked

The current EU GDPR pack is a *technical candidate*: its official anchors, caveat, deterministic fixtures, and version are inspectable, but no independent qualified legal approval record exists. The Research Baseline is a non-legal demonstrator and cannot progress through the legal-review state. Revision 15 implements the non-cryptographic part of this lifecycle as a mandatory version-2 pack contract: authored and reviewed dates, effective date where applicable, review authority and scope, release scope, reviewed locales, change triggers, and lifecycle state are validated when the registry loads. Calendar dates are parsed rather than

checked by shape alone, and the review date cannot precede authorship. An engineering reviewer cannot label a pack legally reviewed or approved, candidate packs cannot claim production scope, and unknown identifiers fail closed. Cryptographic pack signing, reviewer credential verification, expiry enforcement, and remote revocation remain unimplemented.

Each legal pack also carries a source register. A source records its title, HTTPS location, issuing authority, status, and check date. The status vocabulary distinguishes binding law, final guidance, and consultation material. The registry requires at least one binding-law record and requires the primary URL to appear in the register; duplicate, malformed, or impossible-date records fail admission. These controls make source status inspectable but do not authenticate the named authority or prove that a maintainer classified a document correctly. The EU pack consequently remains a technical candidate even though its primary Regulation and final EDPB guidance are official.

4.2.3 Runtime Validation Pipeline

The safe evaluation entry point accepts an unknown value rather than trusting its compile-time annotation. The first check requires a non-array object. The package identifier must be a string between one and 255 characters using a conservative character set. Permission type must belong to the owned set of location, microphone, and contacts. Source, if present, must be simulator, native bridge, or imported. This sequence prevents later operations from assuming properties that have not been established.

The access count must be a non-negative safe integer. The window start and end must also be safe integers, the start must precede the end, duration cannot exceed one day, and the end cannot be implausibly far in the future. When event timestamps are supplied, their array length must equal the access count and every timestamp must lie within the audit window. These conditions ensure that peak-rate calculations do not operate on inconsistent evidence.

Processing context is validated field by field. Purpose, controller identity, Article 9 condition, transparency notice, minimisation assessment, retention justification, and basis-specific evidence references are optional strings at admission. Consent-withdrawn, special-category, user-initiated, and DPIA-required values are optional Booleans. Lawful basis must belong to the Article 6 enumeration used by the prototype. Retention days must be a non-negative safe integer. Values are rejected rather than coerced because coercion can change meaning at a trust boundary. After syntactic admission, the EU pack reports absent cross-cutting and basis-specific items as legal-evidence gaps rather than rejecting the technical observation.

Every rejection returns an explicit code and message with a rejection timestamp. The throwing evaluation method is retained for callers that require an exception, but the application can use the safe result to display a controlled evidence gap. This dual interface allows strict internal use without forcing the UI to recover from an untyped exception.

4.2.4 Signal Computation

After admission, the engine loads the rule for the active pack and computes the prompt threshold

$$T_p = \max(1, \lceil B_p M_p \rceil), \quad (2)$$

where B_p is the pack baseline and M_p is the deviation multiplier. Daily excess is $\max(0, x - T_p)$, where x is the admitted count. A daily-total signal is emitted only when excess is positive, so equality with the threshold remains below the signal boundary.

When event timestamps are present, the engine sorts them and evaluates a sliding sixty-second interval. Two indices move monotonically through the sorted array, so peak calculation is linear after sorting. Permission-specific burst limits are 12 for location, 6 for microphone, and 4 for contacts. A burst signal is emitted when the observed peak is strictly greater than the configured limit.

The history key combines package and permission. Previous entries outside the one-day watermark are discarded. An entry with the same start and end is treated as replay and replaced. Only entries ending no later than the current start contribute to the completed rolling total. Overlapping entries stay in bounded history but do not contribute to that total. A cross-window signal is disabled for simulator evidence and requires at least one compatible completed window.

Risk derives from the greatest normalised excess across daily, burst, and rolling dimensions. Ratios below the first threshold remain low; larger ratios map through medium, high, and critical bands. Risk controls presentation and notification priority, while the pack separately controls the compliance-review status. Keeping these concepts distinct prevents a large technical deviation from bypassing missing legal context.

Algorithm 1 formalises the single-record path. It makes three ordering constraints explicit: admission precedes arithmetic, temporal compatibility precedes aggregation, and the pack decides review status only after missing context has been enumerated. Presentation is deliberately absent from the algorithm because it consumes the stored finding rather than recomputing a verdict.

Algorithm 1 Evidence-bounded evaluation of one candidate record

Require: candidate record r , registered pack p , bounded history H

Ensure: rejection with a reason, or a provenance-preserving finding f

```
1:  $a \leftarrow \text{ADMIT}(r)$   $\triangleright$  type, range, source, time, and context checks
2: if  $a$  is rejected then
3:   return structured rejection
4: end if
5:  $q \leftarrow \text{RULEFOR}(p, a.\text{permission}); T \leftarrow \max(1, \lceil q.\text{baseline} \times q.\text{multiplier} \rceil)$ 
6:  $d \leftarrow \max(0, a.\text{count} - T)$ 
7:  $b \leftarrow \text{PEAKINSIXTYSECONDS}(a.\text{timestamps})$ 
8:  $C \leftarrow \{h \in H : h.\text{key} = a.\text{key} \wedge h.\text{end} \leq a.\text{start} \wedge \text{COMPATIBLE}(h, a)\}$ 
9:  $w \leftarrow \sum_{h \in C} h.\text{count} + a.\text{count}$ 
10:  $s \leftarrow \text{SIGNALS}(d, b, w, q, a.\text{source})$ 
11:  $m \leftarrow \text{MISSINGCONTEXT}(p, a.\text{context})$ 
12:  $c \leftarrow \text{PACKSTATUS}(p, s, m, a.\text{context})$ 
13:  $f \leftarrow \text{CONSTRUCTFINDING}(a, p, q, s, m, c)$ 
14: return  $f$  with source, interval, rule, pack version, rationale, caveat, and missing facts
```

Let n be the number of event timestamps, h the bounded history length for a package-permission key, and k the fixed number of context fields. Admission is $O(n + k)$, timestamp sorting is $O(n \log n)$, the sliding peak pass is $O(n)$, and compatible-history selection is $O(h)$. Rule lookup and pack-status classification are constant time for the current fixed registries. Total evaluation is therefore $O(n \log n + h)$ time and $O(n + h)$ working space. These are implementation-complexity bounds, not performance measurements; the thesis has not established latency, energy, or memory behaviour across devices.

4.2.5 Finding Construction and Communication

The finding identifier combines package and permission so the repository can update a stable subject. Evidence fields contain daily count, peak calls per minute, rolling count, and source. Decision fields contain threshold, excess count, active signals, risk, detection time, regulation identifier and name, legal or research reference, and rationale. Compliance fields contain status, applicable principles, missing evidence, and caveat.

Communication is derived after the finding exists. The title identifies the permission review. The summary translates status and signal identifiers into human-readable language. The recommended action first requests missing evidence; an urgent result suggests pausing processing where appropriate and obtaining human review; other results recommend examining purpose, necessity, and proportionality. Notification priority is urgent only for the strongest status or critical risk, silent for no technical concern, and standard otherwise.

This construction ensures that UI wording is based on stored decision state. A card is not permitted to infer a legal label merely from colour or excess count. It receives the pack name, source, caveat, and missing context that were produced with the finding.

4.2.6 Repository Behaviour

The repository uses a versioned AsyncStorage key. Listing parses the stored array, replacement writes at most the bounded presentation set, clearing removes the key, and saving updates an existing package-permission finding or appends a new one. The latest application context also normalises records created by earlier schemas so that pack identity and evidence fields remain available where they can be reconstructed safely.

Notification state is based on change. A new finding can notify; a reversal of active state can notify; an increase in risk rank can notify. A byte-for-byte identical record is not treated as changed. This policy prepares the product for attention-aware notification without claiming that a longitudinal notification campaign has been conducted.

The decision engine’s internal temporal history and the persistent finding repository serve different purposes. Engine history supports signal computation within a running context. Repository findings support user review across sessions. Conflating them would either store unnecessary event aggregates indefinitely or make visible findings disappear whenever the process restarts.

4.3 Android Bridge and Scheduling

The Kotlin module reports whether the native privacy inspector is available and exposes bounded audit and demonstration operations to React Native. WorkManager workers are registered for periodic or one-time tasks according to Android’s scheduling model. WorkManager is deferrable: the operating system may delay work according to battery, idle state, quotas, and vendor policy. The implementation therefore does not describe a registered interval as proof of exact execution time.

The audit worker reads only evidence available under the application’s deployment authority. On the tested ColorOS device, that authority did not expose unrestricted third-party history. The bridge returns capability and available summaries rather than fabricating other-package events. The simulator worker follows a controlled path and marks every generated record as simulator source.

Native package declarations were moved from the anonymous template namespace to the stable Privacy Lens namespace. Activity, application, bridge, package, and worker source files use the same identifier as Gradle and the manifest. This reconciliation matters for release upgrades, signing continuity, deep links, provider authorities, and store identity.

4.4 Product Interface and Accessibility

The interface translates stored findings into a review journey without introducing new compliance semantics. Rather than exposing every technical field simultaneously, it uses stable destinations and progressive detail while keeping source, pack, and evidence reach close to the status they qualify.

4.4.1 Interface Implementation

The Overview screen has four layers. A brand header establishes identity and provides a direct Settings entry. A high-contrast editorial hero states “See the evidence. Keep the conclusion human”, identifies the active pack and research-candidate status, and presents the primary review action. A framed review snapshot distinguishes needs-review, evidence-gap, and no-concern counts with icon, text, and a structural top accent. An evidence-reach card uses a persistent rail and explicit capability badge to explain whether the native bridge is available and why an empty result is not proof of no access.

The Findings screen is designed for both empty and populated states. The empty state explains what will appear and directs the user back to a review or demonstration. The populated state begins with a dark ledger summary and renders reusable finding cards. Each card combines a status rail, icon, text marker, permission, package, source chip, regulation-pack chip, detected signals, rationale, and missing context. Source and regulation are placed near the result because they alter how the result should be interpreted. The empty-to-populated transition remounts the list at its heading rather than retaining a stale offset. Evidence disclosure is a dedicated minimum-height control, and details are inserted after that control so Android focus retention does not scroll the card title out of view.

The Settings screen separates preference, explanation, provenance, and deletion. Pack selection is restricted to registry entries. Selecting a pack displays its jurisdiction, version, description, caveat, and governance state. A source-register section renders every registered source with authority, checked date, and a distinct binding-law, final-guidance, or consultation label. The visual hierarchy therefore exposes the legal-source distinction added in Revision 15 instead of burying it in code. Clearing findings retains a destructive visual treatment and explanatory copy.

The custom bottom navigation uses three equally weighted targets with icon, label, accessibility role, selected state, and sufficient touch area. It replaced the platform tab implementation after physical coordinate and hierarchy evidence revealed ambiguity in the previous navigation setup. Navigation semantics no longer depend on obsolete routes.

4.4.2 Brand and Accessibility Implementation

Revision 16 formalises the visual language as a restrained editorial system. Evergreen supplies the principal reading surface, brighter teal marks actions, mint identifies supportive and selected states, sand separates synthetic demonstration and caution, and warm off-white reduces background glare. Two border strengths, four radii, and bounded elevation tokens distinguish canvas, cards, selected controls, and floating navigation. Rounded panels group related information without relying on heavy separators. Typography uses strong size and weight contrast between proposition, section heading, count, status, metadata, and explanatory copy. Decorative orbits and accents remain non-semantic and cannot change the interpreted result.

The application icon combines a shield with a lens aperture. The shield suggests protection while the aperture suggests inspection; neither implies surveillance of a person. Adaptive

foreground and monochrome assets support launcher contexts. Splash imagery, application icon, colours, and product name are generated consistently across Android densities.

Accessibility labels are supplied for navigation and important actions. Icons and colour do not carry status alone: status text, a marker, and a structural rail remain available. The evidence disclosure is a separate 44-pixel-or-larger control rather than making an arbitrarily tall card one button. Bottom navigation exposes tab and selected-state semantics and adds a structural selected container. Source cards expose link labels that include source status and title. These implementation facts and UI-tree evidence support an accessibility-oriented design claim, while screen-reader traversal, large-font reflow, switch access, and assistive-technology user testing remain future acceptance work.

4.5 Release Configuration and Privacy Controls

The Revision 16 source sets minimum SDK 24, compile and target SDK 36, version code 5, and version name 1.4.0. The Expo configuration, JavaScript package metadata, settings footer, and Gradle declaration use the same version. Release tasks produced an ARM64 APK for controlled device installation and an AAB for Play delivery. The build reads upload-store path, passwords, and alias from environment variables. Because those values were absent, the accepted QA artefacts use the debug configuration and are not production-upload signed. Package inspection and physical installation were renewed for 1.4.0 rather than inherited from the legal-priority predecessor.

The source manifest explicitly removes inherited Internet, network-state, and legacy external-storage permissions using manifest-merger directives. Inspection of the 1.4.0 APK found only foreground-service, wake-lock, boot-completed, and an application-scoped dynamic-receiver protection permission contributed by the background-work stack. Android backup and remote Expo updates remain disabled. A release-privacy contract checks version agreement, manifest removals, backup and update settings, and the absence of application network-client APIs. These package facts do not prove that every future dependency or connector will preserve the property.

Package reduction was part of implementation. Dependencies used only by removed charts, health views, haptics, image helpers, secure storage, web browser wrappers, gestures, SVG rendering, and network-status features were removed when no longer needed. Generated build and cache directories remain ignored. This reduces attack surface and maintenance burden, although transitive dependencies are still reviewed through the lock file and final manifest.

Findings and the active-pack preference remain in app-private local storage. The release contains no account system, advertising SDK, behavioural analytics SDK, remote evidence database, or evidence-upload workflow. Settings exposes a direct clear-findings action, and Android backup is disabled so that ordinary uninstall or data clearing does not silently restore the review history. These choices do not make local data invulnerable, but they reduce the number of actors and data flows introduced by the reviewing application itself.

The repository also contains a bilingual privacy-policy draft, user guide, product audit, and store-readiness checklist. These documents distinguish binary behaviour from owner-controlled

deployment steps. A public release still requires a stable HTTPS policy URL, support contact, production upload key, console declarations, and internal-track verification. Keeping these requirements outside the code would be insufficient; keeping them only in the interface would be misleading. They are recorded as release evidence linked to the exact build.

Codebase reconciliation supports the same privacy boundary. Legacy health-dashboard, scoring, experiment, charting, duplicate compliance-service, and unused template components were removed because they neither answered a current research question nor belonged to the reviewed product journey. Their dependencies were removed with them. The resulting application has one decision path, one privacy context, one navigation model, and one current user guide, reducing the risk that an obsolete screen or service contradicts the evaluated architecture.

4.6 Implementation-to-Requirement Traceability

The implemented system can be traced directly to the Chapter 3 requirements. **F1**, deterministic evaluation, is realised by immutable pack rules, strict threshold comparisons, explicit signal calculation, and a stable finding identifier. The implementation does not consult wall-clock randomness or network state when deciding a result. Given the same admitted evidence, compatible history, and pack version, the engine produces the same finding semantics.

F2, provenance preservation, is realised by carrying package, permission, source, start, end, observed count, pack identifier, pack version, rule reference, rationale, and missing-context fields into the finding object. Repository replacement preserves these fields, and the Findings screen displays the source and pack near the status. Provenance is therefore not reconstructed from a current setting after the decision.

F3, fail-closed admission, is realised by accepting unknown runtime values and validating their identifiers, numbers, windows, timestamps, source, and optional context before evaluation. Every rejection produces a typed code and timestamp. A rejected record cannot enter temporal history or become a no-technical-concern finding. This mapping is central to the thesis because a permissive default would reverse the intended claim boundary.

F4, registered pack switching, is realised by a closed registry and explicit pack dependency. The engine does not infer a pack from device locale, arbitrary text, or a persisted object. Settings can select only registered identifiers, and historical findings retain the pack that produced them. The non-legal Research Baseline demonstrates switching while its caveat prevents it from being presented as a validated jurisdiction.

F5, simulator separation, is realised through the controlled source enumeration, simulator-owned record construction, exclusion from cross-window evidence, repository preservation, and visible synthetic labels. The simulator exercises the real downstream pipeline, but its values do not acquire device-observation authority when they reach the dashboard.

The usability requirements map to the product shell. **U1** is realised by Overview, Findings, and Settings, a three-target navigation component, purposeful empty states, and one dominant review action. **U2** is realised by textual status labels, icons, non-colour distinctions, evidence-reach copy, source and pack chips, and accessibility labels. These implementation facts support

inspectable design conformance; they do not substitute for participant or assistive-technology studies.

The privacy and delivery requirements map to the assembled Android package. **P1** is realised by app-private storage, disabled backup, direct deletion, no evidence server, and removal of inherited external-storage declarations. **D1** is realised by the stable package identifier, explicit version, current SDK configuration, environment-based upload-key inputs, APK and AAB tasks, branded resources, and installed-package inspection. The QA signing fallback keeps local reproduction possible while remaining visibly distinct from production signing.

This traceability demonstrates why the implementation chapter excludes unrelated legacy features. A component without a requirement, research question, or acceptance source would add maintenance and interpretation risk without strengthening the thesis. Chapter 5 now evaluates each retained path at the source, rule, package, device, and presentation layers.

5 Evaluation and Results

This chapter evaluates whether the implemented system preserves the evidence boundary defined in Chapters 2 and 3. The protocol separates logical conformance, hostile input handling, temporal state, pack identity, legal-evidence completeness, Android packaging, physical interaction, and visual communication. Results are reported before interpretation so that a positive engineering outcome is not silently promoted into a legal, population, or store-approval claim.

5.1 Evaluation Overview and Headline Results

The evaluation separates seven evidence layers: source conformance, deterministic rule behaviour, adversarial boundary behaviour, legal-evidence governance, Android package integration, physical-device behaviour, and rendered-interface communication. This separation prevents a passed unit test from being reported as device reliability, usability, or legal validity. Revision 15 supplied the strengthened legal contract; Revision 16 renews the 1.4.0 release artefacts and adds a four-screen device matrix, non-colour status structure, a visible source-status register, and targeted interaction-state checks.

5.1.1 Source and Rule Verification

The complete application passed `tsc -noEmit` and ESLint. The compliance sources were compiled before execution, preventing stale generated JavaScript from being treated as current evidence.

A fixed-seed logical campaign generated 1,000 cases. Eight hundred were expected review signals and 200 were below-threshold controls. The observed confusion matrix was

$$(TP, TN, FP, FN) = (800, 200, 0, 0), \quad (3)$$

which gives precision and recall of 1.0000 against the specified oracle. The oracle and engine implement the same published threshold model; this is a conformance and regression result, not an independent estimate of real-world infringement detection. Zero observed errors do not establish a zero population error rate.

Extended suites passed for unsafe types and integers, malformed context, unknown source, permission-key mutation, temporal overlap, adjacency, replay, retention, simulator isolation, dashboard presentation, missing-context disclosure, and regulation-pack identity. The tests confirm the enumerated boundaries. They cannot establish correctness for telemetry the application never receives or for legal interpretation outside the fixtures.

The retained deterministic campaign contains 1,800 cases across both packs and all three supported permission families. Its expected daily signal and threshold are calculated in a small reference function rather than by calling the production helper. Values at and immediately above each threshold confirm the strict comparison, while a fractional product confirms ceiling behaviour. Revision 15 extends governance mutations to impossible calendar dates, absence of a binding-law source, and explicit retention of consultation-material status. This improves oracle diversity and pack-contract coverage without turning generated cases into population evidence.

5.1.2 Legal-Evidence Completeness Verification

The legal tests now distinguish the six Article 6 bases instead of treating a selected enum value as complete evidence. Consent requires an Article 7 evidence reference; contract requires a necessity assessment; legal obligation and public task require an Article 6(3) mandate; vital interests requires a necessity assessment; and legitimate interests requires a three-part assessment. Cross-cutting probes require transparency, minimisation, retention justification, and a DPIA screening result. A positive DPIA screen without an assessment reference remains evidence-incomplete. Malformed new strings and Boolean fields fail with the same structured context error as earlier unsafe context.

These tests establish implementation of the declared question matrix, not legal sufficiency of the referenced records. The engine does not open, authenticate, or assess the legal reasoning in a reference. A withdrawn-consent observation is emitted as `POTENTIAL_CONFLICT` with urgent review; the previous `LIKELY_NON_COMPLIANT` value is migration-only. This wording change reduces adjudicative overreach without asserting that every post-withdrawal access is unlawful, since timing, another legal basis, exceptions, and processing facts may still matter.

5.1.3 Release Build and Manifest Verification

Gradle completed `assembleRelease` and `bundleRelease` with target SDK 36. Table 7 records the generated artefacts.

Table 7: Evaluated Android release artefacts

Artefact	Bytes	SHA-256
app-release.apk	59,407,715	FB87A9B32E6F1A5FBC83F9EA35F0051D381 524DBD0EE82DE3E7BCCA0CF58EBF7
app-release.aab	29,607,554	7EE5B3A1DFE4C6AFA42DAEB41BC1424A447 32EC4BD8637547D52751F13DB3061

The 1.4.0 merged manifest contained wake lock, boot rescheduling, foreground service, and an app-scoped dynamic-receiver protection permission. Internet, network-state, and external-storage permissions were absent. Inspection confirmed application ID `com.zhihengzhang.privacylens`, version code 5, version name 1.4.0, minimum SDK 24, and target SDK 36. The signer remained the Android debug QA certificate rather than an owner-controlled production key.

5.1.4 Physical-Device Evaluation

The 1.4.0 networkless release APK was installed by ADB on one OPPO PERM00 device. Five force-stop and cold-start rounds completed. Median and mean `TotalTime` were 1,941 ms and 1,998.2 ms; sampled total PSS ranged from 106,713 to 108,019 KB; and the crash buffer contained no package-specific entry. A final four-screen matrix covered the Overview, populated Findings, pack-selection, and source-register states without obvious clipping or overlap. These are short-run observations, not a startup service-level objective, leak proof, battery evidence, legal validation, or broad performance claim.

The three primary destinations rendered at the device’s captured application resolution. Coordinate-driven exploration initially appeared to show a navigation-target problem because the physical display and captured application coordinate spaces differed. A bounded coordinate sweep and UI-tree inspection identified the offset; subsequent navigation and custom bottom-bar interaction succeeded. This calibration episode is retained because it demonstrates why raw ADB coordinates must not be interpreted without the target’s current window geometry.

The controlled demonstration completed a 50-record synthetic workflow. The populated overview displayed four review items, zero evidence gaps, two no-concern items, and descriptive precision and recall of 100% for the simulator’s known labels. The Findings screen marked the source as synthetic and displayed the EU GDPR pack. These values validate the demonstration path only; they are not device-observation accuracy.

5.1.5 Usability and Trust Assessment

The product was assessed against three release-candidate questions.

Table 8: Expert product-screening assessment

Criterion	Supporting evidence	Boundary
Visually coherent	consistent brand, hierarchy, spacing, cards, iconography, and three-destination navigation on the tested device	one device; no participant preference study
Trustworthy	local-first disclosure, source and pack labels, explicit missing evidence, no legal-verdict language	not a legal audit or security certification
Interaction feasibility	clear first action, concise summaries, optional details, demonstration, and direct data clearing	expert walkthrough only; no participant or longitudinal adoption measure

The earlier simulated multidisciplinary screen assigned 95.1/100 under a release-candidate engineering rubric. The stricter publication-target rubric instead resets the predecessor App-thesis system to 70/100 and assessed Revision 15 at 79/100. Under the same rubric, Revision 16 is assessed at 84/100: legal validity and governance 24/30, software evidence 23/25, interface and accessibility 19/20, user-psychology evidence 10/15, and academic communication 8/10. The five-point movement from Revision 15 reflects the visible source register, non-colour status cues, renewed release and device evidence, and correction of two state-transition defects. It does not award points for external legal approval, participant outcomes, assistive-technology conformance, long-duration reliability, or journal acceptance.

5.2 Protocol and Reproducibility

The protocol begins from six evaluation questions and assigns a separate acceptance condition to each. Reproducibility controls then ensure that the tests exercise current generated code, fixed temporal values, and explicit boundaries rather than relying on stale output or chance coverage.

5.2.1 Evaluation Questions and Acceptance Logic

The evaluation was organised around six questions. First, does the source compile and satisfy static rules? Second, does the decision engine agree with its specified thresholds and statuses? Third, does it reject malformed or adversarial evidence safely? Fourth, do regulation packs remain independent and traceable? Fifth, can the Android release artefacts be assembled, inspected, installed, and launched? Sixth, does the rendered product communicate the expected hierarchy and evidence boundaries on the available physical device?

Each question has a distinct acceptance condition. Source acceptance requires a clean TypeScript application check and ESLint run. Rule acceptance requires exact expected confusion-matrix values for the fixed corpus and passing boundary assertions. Runtime-

boundary acceptance requires every malformed case to return the specified rejection code without an uncaught exception. Pack acceptance requires registry allowlisting, different thresholds where defined, retained pack identity, and a non-legal caveat for the Research Baseline.

Android build acceptance requires both release tasks to complete and the artefacts to exist with recorded hashes. Permission acceptance requires agreement between the merged manifest and installed package, with legacy external-storage permissions absent. Physical acceptance requires successful streamed installation, launcher start, stable primary navigation, captured final states, and no matching fatal error in the bounded log scan. Visual acceptance requires inspection of actual screenshots rather than source code alone.

The protocol deliberately does not set an acceptance condition for GDPR compliance, legal accuracy, population detection accuracy, long-run scheduler reliability, battery consumption, multi-device compatibility, or store approval. Those outcomes require evidence not collected in this campaign.

5.2.2 Claim-to-Evidence Acceptance Matrix

Table 9 is the evaluation’s decision record. A claim passes only if the named evidence exists and agrees at the relevant layer. The last column prevents a pass in one row from being borrowed by a stronger claim in another discipline.

Table 9: Claim-to-evidence acceptance matrix

Claim	Required and observed evidence	Decision	Explicit non-inference
Specified rule conformance	compiled current sources, 1,000 fixed-seed cases, exact threshold boundaries, finite metrics	pass	no population accuracy or lawful-processing classifier
Hostile admission and temporal safety	malformed types, unsafe numbers, source mutation, replay, overlap, and cross-window suites	pass	no proof against every adversarial input or external producer
Pack separation and traceability	registered EU and Research packs, differing rules, retained identifiers, sources, caveats, fixtures	pass	no legal approval of EU mappings or validity of future packs
Release artefact identity	APK/AAB existence and hashes, merged manifest, installed-package facts, stable package identifier	pass	no production signing, console acceptance, or broad compatibility
Bounded physical operation	streamed install, launch, navigation, UI tree, screenshots, and filtered fatal-log scan on OPPO PERM00	pass	no long-run reliability, scheduler guarantee, or multi-OEM result
Human and legal outcomes	participant protocol, assistive-technology campaign, independent legal decision record	not evaluated	no usability population claim, accessibility conformance, trust adoption, or GDPR compliance

The matrix also defines contradiction handling. If source tests pass but the assembled package exposes a forbidden permission, the package claim fails. If screenshots look coherent but a participant study later shows systematic misunderstanding, the human outcome fails without erasing the recorded visual implementation. If legal review rejects a mapping, the pack’s legal-review state fails even when its deterministic fixtures still pass. Evidence closer to

the disputed claim takes precedence; heterogeneous passes are never averaged into a universal readiness score.

5.2.3 Reproducibility Controls

The compliance sources are compiled immediately before the generated JavaScript runners execute. This control was introduced because a failed TypeScript command can otherwise leave older generated output in place, producing a misleading green test. The acceptance sequence checks the compiler exit code before invoking Node. Generated compliance output is ignored by Git so it cannot become an accidental source of truth.

The main logical corpus uses a fixed linear-congruential seed. A fixed seed makes a failure reproducible while still exploring varied generated values. The 1,000 rounds preserve an 80/20 design: four of every five cases are expected review signals and one is a control. A fresh engine is created for each main-corpus case so temporal state from one generated subject cannot leak into another.

Boundary cases are not left to random chance. For every supported permission, counts at threshold minus two, minus one, equal, plus one, and plus two are exercised. The assertion encodes the strict boundary: only values above the prompt threshold become active under the daily-total rule. Burst tests similarly check exactly at and exactly above each permission-specific limit.

Time is fixed in extended tests. Windows and timestamp arrays are created relative to a constant epoch value rather than the wall clock. This prevents a test from becoming invalid while it runs and makes future failures comparable. Production validation still compares implausible future times against the current clock, but fixtures choose values far enough in the past to avoid accidental rejection.

5.3 Engine and Regulation-Pack Results

The engine evaluation proceeds from nominal threshold conformance to hostile admission, temporal replay, accountability context, and pack independence. This order tests not only whether a rule fires, but whether unsafe evidence is prevented from reaching the rule and whether the resulting meaning survives into presentation.

5.3.1 Deterministic Corpus Interpretation

The confusion matrix compares the engine’s active-state decision with the generator’s specified expected label:

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}. \quad (4)$$

The resulting precision and recall of 1.0000 mean that the implementation agreed with the oracle for all generated cases. This supports a regression claim: later refactoring did not alter the expected threshold semantics for the corpus. It also supports a fixture-quality claim: the

corpus contains both positive and negative cases and the metric implementation returns the expected counts.

It does not show that a threshold crossing identifies unlawful processing. The expected label is generated from the same documented research model that the engine is meant to implement. Although the oracle is expressed independently of the evaluated result, both are derived from one specification. A defect in that specification could be reproduced consistently. There is no independently labelled sample of real applications or legal cases.

The tests also verify metric edge conditions. Empty samples, all-negative samples, and samples with missed positives must produce finite precision and recall rather than not-a-number or infinity. This matters because a dashboard can otherwise display an apparently precise but meaningless metric when a denominator is zero.

5.3.2 Adversarial Input Results

The initial security cases reject unsupported permission values, not-a-number counts, negative counts, reversed windows, and prototype-like permission keys. Extended cases add unsafe integers and positive infinity. Every such count returns the invalid-count code. Timestamp evidence is rejected when array length differs from count, when an element falls before the start, when it falls after the end, or when it is fractional.

Processing context is attacked through type confusion. Numeric purpose, array controller identity, object Article 9 condition, string consent-withdrawn, numeric special-category, and string user-initiated values all return invalid context. An unknown lawful basis and negative retention period are also rejected. These cases are significant because JavaScript truthiness or an early string operation could otherwise convert malformed evidence into a plausible status or throw outside the controlled error path.

Source validation rejects a forged source. This prevents a caller from using an arbitrary label that resembles an authoritative acquisition path. Package and permission history are isolated: a high-count record for one package cannot make another package active. Histories are also isolated by permission family.

The assertions demonstrate the implemented rejection surface, not complete security. They do not cover a compromised native module, a modified application binary, corrupted AsyncStorage beyond the migration parser, malicious operating system, or every possible Unicode identifier. Security conclusions remain bounded to the enumerated input contract.

5.3.3 Temporal and Replay Results

Burst tests create timestamp arrays exactly at and just above the location, microphone, and contacts limits. At-limit arrays produce no burst signal; one additional event produces the signal. A separate imported-evidence example concentrates 20 location events within one minute and is detected.

Cross-window tests use two below-threshold location summaries. The first remains inactive. When the second begins after the first has completed, their rolling total exceeds the threshold

and the engine emits a cross-window signal. The same construction is repeated with adjacent windows to establish that equality between a prior end and current start is admissible.

An overlapping construction uses one interval that intersects the next. The second finding's rolling count remains its own count and no cross-window signal appears. This is the expected conservative behaviour because aggregate evidence cannot identify duplicate events in the overlap. An exact replay is then evaluated twice before a compatible next interval. Rolling total confirms that the repeated window contributed once, not twice.

Simulator histories are excluded from cross-window signalling. The demonstration may still produce a daily or burst signal for a record, but it cannot present accumulated simulator activity as imported temporal evidence. This is a provenance safeguard, not a statistical assumption.

The temporal tests operate in process memory and do not demonstrate persistence of aggregate history across process death. They establish the logic of the current engine instance. A production design requiring durable cross-window evidence would need a versioned evidence-history repository and migration tests separate from the visible finding repository.

5.3.4 Accountability and Presentation Results

A low-count contacts record without processing context returns insufficient evidence rather than no concern. Its communication priority is standard because the absence of legal context warrants review. A low-count location record with purpose, contract basis, controller identity, one-day retention, and user-initiated context returns no technical concern and a silent priority.

A microphone record marked as consent-based, special-category, and accessed after consent withdrawal returns the urgent review status. The fixture does not prove a real GDPR infringement; it confirms that the pack's stated escalation branch is implemented. A special-category microphone record without an Article 9 condition lists that condition among missing evidence. An invented lawful basis and invalid retention value are rejected at admission.

Dashboard projection preserves three different states from the accountability fixtures. Summary totals show one action item, one evidence gap, and one no-concern result. Filters surface the expected finding. Every presentation state has both a textual label and marker, ensuring that source-level semantics do not depend on colour alone.

5.3.5 Regulation-Pack Results

The registry exposes at least two packs and recognises only the two expected identifiers. A forged identifier is rejected. The same location audit with a count of 40 is evaluated under both packs. The resulting findings retain different regulation identifiers and different thresholds because each pack supplies its own baseline.

This test is stronger than checking that a dropdown changes text. It passes the same admitted evidence through two independently instantiated engines and inspects the resulting finding objects. The Research Baseline result also contains a caveat stating that it is not law. The test does not validate either pack's legal interpretation; it validates decoupling, identity,

and disclosure.

The registry now validates every bundled pack before exposing it. The contract requires dated governance metadata, a review authority and scope, reviewed locales, change triggers, compatible lifecycle and release-scope combinations, an HTTPS source for legal frameworks, and complete finite rule parameters. The test suite mutates these fields to show that an engineering-only record cannot promote a pack to legal approval. Independent domain review and cryptographic distribution remain outside this technical validator.

5.4 Android Artefact and Device Results

The Android evaluation connects source acceptance to the assembled and installed product. Build reproduction establishes artefact creation; manifest reconciliation establishes package declarations; physical interaction establishes a bounded working journey; and visual review establishes that the principal claim boundaries remain visible on the tested display.

5.4.1 Build Reproduction

The release build initially encountered Windows path and dependency-layout problems. A hoisted local dependency layout allowed native CMake paths to remain within workable bounds. React Native’s new architecture was disabled for the release candidate because enabling it in this Windows checkout reproduced path-related native build failures. This is a build-environment decision, not evidence that the legacy architecture is preferable for future releases.

Gradle reported deprecated features that will require attention before Gradle 9 compatibility. It also reported an SDK XML version warning caused by differences between installed Android Studio and command-line components. Neither warning prevented the recorded release build, but both are retained as maintenance debt rather than suppressed.

The release task generated bundles for multiple Android ABIs. The APK was installed on the ARM64 physical device. The AAB was not uploaded to Play Console, processed by Play signing, or delivered through an internal track. Its success establishes local bundle generation only.

Hashes were recorded immediately after the final permission-manifest rebuild. Any source or build-tool change requires rebuilding and computing new hashes. The files are generated artefacts and are not treated as immutable simply because their names remain the same.

5.4.2 Manifest and Installed-Package Reconciliation

Source-manifest inspection alone was insufficient because dependencies contribute declarations through manifest merging. The initial merged manifest exposed legacy external-storage permissions inherited from transitive Expo components. The evaluated source added explicit merger removals for read and write external storage, rebuilt both artefacts, and inspected the merged result again.

The first candidate still inherited Internet and network-state declarations even though

application code contained no network client. Revision 14 changed these from documented-but-unused permissions into explicit manifest removals and rebuilt the artefacts. The final manifest contains only wake-lock, boot-completed, foreground-service, and an application-scoped dynamic-receiver protection permission. This package-level check supports the stronger, still bounded statement that the tested APK cannot initiate ordinary application networking because the permission is absent.

The package dump also confirmed application identity, version name, version code, minimum SDK, target SDK, primary ABI, install state, and signing version. The release used the QA fallback signing identity because production environment variables were absent. A successful package inspection is therefore not proof of final upload-key configuration.

5.4.3 Physical Interaction Protocol

The device was identified by ADB before installation. The final release APK was streamed with replacement enabled, the package was force-stopped, and the launcher intent was injected. Package metadata was read after installation so that a stale or differently named package could not be mistaken for the new build.

Interaction evidence included screenshots and UI hierarchy dumps for overview, populated findings, pack selection, and the source-status register. The Revision 16 walkthrough exposed two reproducible defects: retained scroll offset after the ledger changed from empty to populated, and context loss when whole-card disclosure moved focus. Keyed ledger state and a dedicated 44-pixel-minimum disclosure control corrected both before the final matrix was captured. The five-round final run retained an empty package crash buffer. These checks validate bounded rendered behaviour on one device; they do not validate the legal content of either pack.

An important measurement issue arose from coordinate spaces. The physical display reported 1080 by 2400 pixels, while the captured application window and hierarchy used 1080 by 2244. Initial taps derived from the full display therefore appeared to miss navigation targets. A bounded coordinate sweep and hierarchy inspection established the offset, after which navigation succeeded. The event is reported because excluding it would conceal a real threat to ADB-based UI-test validity.

Future automation should select accessibility nodes or stable resource identifiers instead of relying on raw coordinates. It should record display size, application bounds, system-bar insets, orientation, density, and font scale for every run. Screenshots should be associated with the installed package hash and build hash so that visual evidence cannot drift from the tested binary.

5.4.4 Visual Review Method

Visual assessment considered hierarchy, identity, first-action clarity, state distinction, provenance visibility, limitation placement, empty states, navigation consistency, destructive-action treatment, and overall coherence. Revision 14 recorded a numeric expert-proxy score only to operate the project's internal iteration gate. Because no validated participant scale was administered, the number is not an estimate of population preference or trust.

The overview succeeded by making the product proposition and review action dominant while keeping the evidence-reach warning in the first viewport. Summary cards use count, icon, label, rail, and colour so status is not colour-only. Findings place synthetic source, active pack, and a text status marker next to the decision. Settings expose source authority as binding law, final guidance, or consultation material. The floating navigation uses both shape and text treatment for its selected state. Large-font testing could not be executed because the OEM shell denied the required setting change, and TalkBack was not evaluated; the result is therefore not an accessibility-conformance claim.

The result satisfies a project-defined candidate gate rather than a universal usability standard. One evaluator and one physical device cannot establish willingness to adopt across a population. The strongest supported statement is that the rendered candidate is coherent, presents its boundaries, and is suitable for internal-track and participant evaluation.

5.5 Interpretation and Validity

The final stage asks how much authority the combined evidence supports. It triangulates independent artefacts, examines correlated-oracle risk, interprets metrics, applies the project-defined store-candidate gate, and defines which changes would invalidate the result.

5.5.1 Result Boundary

The evidence supports the claim that the specified code compiles, the rule engine conforms to its fixtures, the Android artefacts build, the final package installs and launches, the main interface renders coherently, and provenance remains visible in the tested workflows. It does not support claims about 24-hour execution, battery cost, memory leaks, unrestricted third-party observation, multiple OEMs, legal compliance, accessibility conformance, production signing, or store acceptance.

5.5.2 Evidence Triangulation

No single method in the campaign was designed to carry the complete acceptance claim. Source inspection is good at confirming constants, branches, labels, and data flow, but it cannot establish which code entered an APK or how a screen rendered. Automated tests execute specified behaviours repeatedly, but their assertions can share assumptions with the implementation. Package inspection observes the assembled artefact, but it cannot prove the meaning of every runtime path. Device interaction observes a concrete installation, but the sample is narrow. The evaluation therefore uses convergence among methods while preserving the scope of each.

Figure 3 presents this assurance structure. Solid arrows connect evidence produced in the study; dashed arrows mark validation that remains external or future. The shape prevents a large logical test count from visually overwhelming the smaller but qualitatively different device, human, and legal evidence classes.




Figure 3: Evaluation triangulation and claim boundary

The deterministic tests and source review converge on threshold semantics. Source review identifies the pack-provided baseline, multiplier, daily threshold, burst limit, and source restrictions. Boundary assertions then test values immediately around those decision points. The fixed corpus tests varied inputs at scale. Agreement across these methods reduces the chance that the reported rule differs accidentally from the implemented one. It still does not establish that the rule is a valid legal classifier.

Runtime-boundary tests and presentation tests converge on uncertainty handling. Malformed evidence is rejected at admission, low-context evidence becomes an evidence gap, and presentation projection retains that state. This end-to-end check matters because a correct engine could be undermined if the UI translated an unknown state into reassuring copy. The tests show that selected fixtures preserve meaning through the current projection; screenshot review confirms the principal labels are visible in the rendered flow.

Build logs, output existence, checksums, manifest inspection, installed-package metadata, and physical launch converge on artefact identity. Each closes a different substitution risk. A successful Gradle message without an output file would be insufficient. A local file without a checksum could be confused after rebuilding. Source manifest facts could differ after dependency merging. An installed package could be a stale version. Reading these signals together supports the conclusion that the evaluated handset ran the recorded release candidate.

Visual review and UI-hierarchy inspection also complement each other. A screenshot shows layout, colour, clipping, and hierarchy as a person sees them. A hierarchy dump exposes text, bounds, and accessibility-node structure that may not be obvious visually. Neither replaces screen-reader interaction, large-font testing, or participant comprehension. Their combined use supports a bounded expert assessment and identifies the next accessibility evidence required.

Triangulation is not vote counting. If one authoritative check contradicts another, the discrepancy must be resolved rather than outvoted. For example, an installed-package dump showing a sensitive permission would invalidate a source-only conclusion even if several documents said otherwise. The evaluation gives greater weight to evidence closest to the released artefact for package claims and to independent human or legal review for substantive claims.

5.5.3 Oracle Independence and Correlated Defect Risk

The main corpus labels are derived from the same published threshold model the implementation is intended to realise. This is appropriate for conformance testing: it asks whether the program implements the specification. It is insufficient for empirical validation because a mistaken specification can yield perfect agreement. The measured 1.0000 precision and recall are therefore implementation-oracle agreement values, not estimates of performance on real legal cases.

Several design choices reduce accidental coupling. The generator records expected polarity directly from its construction, a fresh engine limits cross-case state, boundary cases are written explicitly, and extended tests target properties beyond the main daily-total rule. Nevertheless, developer knowledge shaped both sides. The campaign did not use a separately developed executable oracle, blinded labels, or an external benchmark.

Mutation testing would provide one useful next check. Deliberately changing a comparison from greater-than to greater-than-or-equal, weakening source validation, removing replay replacement, or allowing overlapping history should cause specific tests to fail. Surviving mutations would identify assertions that execute code without distinguishing correct behaviour. Property-based generation could then explore many valid and invalid shapes while preserving stated invariants.

An independent reference implementation could further reduce correlated defects. It should be written from the regulation-pack contract by a reviewer who has not read the production engine, ideally in a different language or style. Both implementations could evaluate a versioned interchange corpus. Disagreement would trigger case analysis rather than majority voting. This process would validate formal semantics, though it would still not validate legal relevance.

For legal or ecological validation, the oracle must change. Controlled event traces can provide ground truth about whether acquisition observed an action. Independent privacy reviewers can label whether a scenario warrants investigation, using a documented protocol and recording disagreement. Neither source should be described as perfect legal truth. The evaluation question, sampling frame, and annotation uncertainty must be reported alongside performance metrics.

This distinction explains why the thesis retains the term deterministic corpus rather than benchmark. A benchmark normally suggests comparison against an external reference population. The current corpus is a rigorous regression artefact that protects rule semantics as the code changes. Its value is high, but different from an independently labelled real-world set.

5.5.4 Metric Interpretation and Error Analysis

Precision and recall were included because they expose false-positive and false-negative cells more clearly than accuracy in an imbalanced corpus. The designed 80/20 mixture would allow a trivial always-positive strategy to appear 80 per cent accurate. The full confusion matrix prevents that reading: all 200 controls were negative and all 800 constructed signals were positive under the specified oracle.

The exact totals also make the experiment reproducible. The campaign did not sample a changing external population; it evaluated a fixed generated set. Confidence intervals around the observed proportions would describe repeated sampling from an assumed generator, not uncertainty about the fixed run. More importantly, they would not address specification validity. For this reason the report presents exact agreement and claim boundaries instead of using a narrow interval to imply real-world certainty.

Error analysis remains meaningful even when the main matrix has no disagreements. Adversarial suites reveal categories the primary generator does not represent: unsafe numeric values, prototype-derived identifiers, malformed context, source forgery, temporal overlap, replay, and history isolation. These tests show that corpus size alone is not coverage. A small, carefully targeted case can provide more security evidence than hundreds of additional values drawn from the same safe range.

Future field studies should report more than aggregate precision and recall. Results should be stratified by permission, application category, source authority, foreground status, OEM, Android version, and window resolution. Reviewers should inspect false positives and negatives for common causes. Missing acquisition records should not be treated as model negatives. Cases excluded because evidence is insufficient should be counted and explained rather than removed from the denominator without notice.

Calibration should be evaluated separately from discrimination. A status hierarchy may rank more concerning cases correctly while its displayed wording still conveys excessive certainty. Participant studies can test whether confidence changes appropriately between authoritative, synthetic, and incomplete evidence. This human calibration outcome is central to the product's trust claim and cannot be inferred from the engine matrix.

Operational metrics also need denominators. A launch success rate requires a defined number of clean starts and devices. A crash-free percentage requires a meaningful observation period and collection method. A jank percentage depends on a representative interaction trace and enough rendered frames. The current evaluation reports bounded events directly rather than producing population-style percentages from a short sample.

5.5.5 Store-Candidate Evidence Review

The project-defined store-candidate gate was assessed across product identity, core journey, disclosure, packaging, privacy posture, accessibility foundations, and operational ownership. Product identity was supported by a consistent name, icon, visual language, and package identifier. Core journey was supported by first-launch overview, findings navigation, detailed rationale, pack selection, demonstration control, and local-data deletion. Disclosure was supported by visible source labels, pack identity, evidence limits, and non-verdict language.

Packaging evidence included release APK and AAB generation, target SDK 36, final checksums, manifest reconciliation, signature verification, and installation of the APK. Privacy posture included no network or sensitive runtime permission in the final package, disabled backup and remote updates, local-first storage, and a release contract that scans application source for network-client APIs. Accessibility foundations included textual state labels, non-

colour cues, explicit control labels and hints, selected/expanded state, a polite live region, target sizing, and a simple navigation structure.

Supply-chain review began with 30 npm findings: two critical, 17 high, ten moderate, and one low. Compatible lock-file updates plus explicit patched overrides for `brace-expansion` 1.1.18 and `postcss` 8.5.26 removed the critical findings and reduced the report to 20 findings, comprising ten high and ten moderate reports rooted in Expo/Metro build-time image parsing and Apple configuration tooling. The final Android artefact was rebuilt after the updates. This reduction is useful maintenance evidence, but it is not a zero-vulnerability claim; the remaining advisories require a separately validated Expo major transition or upstream resolution.

The gate was not interpreted as proof that every store prerequisite is complete. Production upload signing was absent. The privacy policy required owner-controlled HTTPS publication. Console declarations, listing assets, support details, content rating, internal-track delivery, and store review remained external. These gaps do not negate the engineering candidate; they determine why it must not be described as a publicly released production version.

The visual adjectives in the project gate were converted into inspectable expert-review criteria. Visual coherence meant consistent hierarchy, spacing, typography, colour restraint, branded identity, and absence of visible broken states in the inspected flows. Claim-boundary coherence meant that evidence provenance, pack governance, and limitations were prominent, destructive controls were explicit, and the product did not present a legal approval badge. Interaction feasibility meant that the evaluator could understand the purpose, navigate the product, switch and persist the review framework, and control data without a blocking defect. The resulting 95.1 score authorises the next engineering stage; none of these observations measures population preference, trust, adoption, or willingness to continue use.

These definitions create a repeatable expert gate, but they do not convert subjective judgement into population evidence. A single evaluator can miss cultural preferences, accessibility barriers, unfamiliar terminology, or long-term fatigue. The appropriate next step is internal-track and participant evaluation using the actual distributed candidate. The expert gate answers whether that investment is justified, not whether the market has accepted the product.

The candidate judgement also depends on the declared product category. A local research and accountability aid can be useful with bounded imported or simulated evidence. A product marketed as universal device surveillance would fail because the acquisition authority is not established. Store copy must preserve the former positioning. Changes to marketing are therefore capable of invalidating readiness even if the binary remains identical.

5.5.6 Evaluation Repeatability and Change Control

The evidence applies to a specific source revision and generated artefacts. Repeating only the final UI flow after changing engine code would not renew rule evidence. Rebuilding only the APK after changing privacy copy would not renew screenshot evidence. The project therefore treats acceptance as a dependency graph: a change invalidates every downstream artefact that incorporates or presents it.

Engine or pack changes require TypeScript compilation, logical suites, boundary suites, presentation projections, release rebuilds, hashes, and targeted device review. Android manifest or Gradle changes require merged-manifest inspection, package dump reconciliation, rebuilds, installation, and launch. UI changes require lint, package build, hierarchy and screenshot capture, accessibility checks, and renewed expert review. Policy or store-description changes require reconciliation with package behaviour and data flows.

Reproducible scripts reduce omission but do not remove the need to inspect their outputs. A command may exit successfully while testing stale generated code, selecting the wrong device, or reading a prior file. The campaign therefore compiles generated runners immediately before execution, checks package identity after installation, and associates evidence with checksums. Future automation should emit a machine-readable manifest linking commit, tool versions, commands, artefact hashes, device facts, and screenshots.

Environment details also matter. Windows path length and native dependency layout affected release compilation. The chosen workaround and architecture flag belong in the build record because a different environment may produce different outputs. Tool warnings about future Gradle compatibility remain maintenance observations even when the current build passes.

A release tag should be immutable, while packs and policies use explicit version identifiers. If a later investigation finds an error, the project should publish a successor and a correction note rather than rewriting the historical evidence. This thesis follows the same principle by creating Revision 16 chapter files instead of replacing Revision 15. Preservation makes it possible to see which claims changed and why.

Repeatability finally includes honest failure recording. Network, dependency, device-coordinate, and build-path failures can be part of the evidence because they expose operational assumptions. Removing them from the narrative would make future reproduction harder. A trustworthy evaluation explains both the final pass and the conditions needed to obtain it.

5.6 Evaluation Synthesis

The results answer the engineering and interface portions of the Revision 16 research problem. The engine accepts and rejects the enumerated evidence classes deterministically; replay and overlap rules prevent the tested forms of aggregate amplification; pack switching changes rule-owned output without changing generic infrastructure; and legal-evidence tests preserve source status, basis-specific gaps, and non-verdict language. Renewed Android evidence connects Revision 16 source to 1.4.0 release artefacts and a bounded rendered workflow, while Revision 15 evidence remains historical rather than overwritten.

The strongest supported conclusion is therefore a versioned chain rather than a single score. Revision 16 retains the strengthened legal-evidence contract, passes current source acceptance, builds as version 1.4.0, and completes a bounded physical-device and visual-evidence round. The chain still stops before legal adjudication, unrestricted acquisition, population performance, broad usability, accessibility conformance, long-run reliability, production signing, and store approval. Chapter 6 examines why those boundaries remain material and which future evidence

could move them.

6 Discussion and Limitations

This chapter interprets the results through the same bounded-claim discipline used in evaluation. It first states what the prototype establishes, then examines technical, legal, human, and operational validity. The discussion does not repeat the result tables; it explains why each remaining limitation changes the authority of the conclusion and which future evidence would be capable of resolving it.

6.1 Interpretation of Findings

6.1.1 What the Prototype Establishes

Privacy Lens demonstrates that a permission-review product can be useful without claiming complete visibility or automated legality. The core value is not the threshold alone; it is the evidence contract surrounding the threshold. Source, window, pack, rule reference, and missing context make the decision traceable. Fail-closed validation prevents malformed evidence from becoming a reassuring result. The interface then communicates the distinction between an observed signal, a synthetic example, and an evidence gap.

The replacement of a GDPR-specific engine with regulation packs also improves scientific clarity. It reveals which statements belong to generic software and which depend on a jurisdictional interpretation. This supports future extension while making clear that interface compatibility is not legal equivalence.

6.1.2 Construct and Internal Validity

The deterministic confusion matrix measures agreement with a defined threshold oracle. It does not measure whether that threshold identifies unlawful processing. Because test fixtures and implementation share the same conceptual specification, correlated defects remain possible. Boundary suites reduce but do not eliminate this risk. Independent implementation review, mutation testing, and event-level reference traces would strengthen internal validity.

The simulator is intentionally useful for reproducible demonstration, but it can inflate perceived confidence if its labels are mistaken for field ground truth. Revision 14 mitigates this by preserving synthetic provenance through the repository and interface. The remaining risk is user over-reliance on any precise metric; future builds should consider hiding precision and recall outside an explicitly labelled research panel.

6.1.3 External and Ecological Validity

Physical evaluation used one OPPO PERM00 device and a short acceptance window. The Revision 16 final run comprised five cold starts, with a median 1,941 ms launch, mean 1,998.2 ms, and sampled PSS range of 106,713–108,019 KB without a package-specific crash-buffer entry. OEM process management, Android versions, accessibility settings, screen sizes, font

scaling, languages, and long-run scheduling can change behaviour. These samples are an anomaly check, not evidence of a startup objective, 24-hour WorkManager reliability, battery impact, leak freedom, or multi-device stability.

The ADB coordinate calibration problem also illustrates a methodological limitation: automated interaction evidence depends on the relationship between physical display coordinates, application-window bounds, navigation bars, and screenshots. Future UI automation should anchor actions to accessibility nodes or resource identifiers and record window geometry with every run.

6.1.4 Legal and Regional Validity

The GDPR pack maps technical signals to questions motivated by official provisions [1]. Revision 15 makes more of the missing legal work explicit by asking for basis-specific support, transparency, minimisation, retention justification, Article 9 context, and a DPIA screening outcome. A supplied reference is evidence that a record exists in the input; it is not proof that the record is accurate, applicable, current, or legally persuasive. The engine cannot determine controller identity, processing purpose, lawful-basis validity, necessity, proportionality, consent validity, data-subject expectations, recipients, safeguards, or jurisdiction. Its output remains an accountability aid, and a qualified reviewer must examine the underlying processing and applicable law.

Adding another regional pack requires more than translation. The maintainer must identify scope, legal authority, effective date, regulated actors, exceptions, remedies, and concepts that have no GDPR analogue. Each pack needs independent legal sign-off, versioned fixtures, source preservation, localisation review, and an update policy. The current Research Baseline pack is not evidence that a second legal regime has been validated.

6.1.5 Usability and Accessibility

The redesigned interface applies recognised accessibility and usability principles [21–23]. Revision 16 adds text and shape alongside colour, minimum-size disclosure controls, a stable selected-navigation structure, and an explicit legal-source status register. The four-screen device matrix supports a positive expert assessment, and the walkthrough found and corrected two interaction-state defects before final capture. However, conformance-oriented inspection cannot establish comprehension, efficiency, satisfaction, trust calibration, or willingness to continue using the product. An attempted large-font change was rejected by the OEM shell permission model, and no screen-reader campaign, switch-access test, colour-vision simulation, successful large-font matrix, or participant study was completed. These remain required before making an accessibility-conformance or user-adoption claim.

6.1.6 Interpreting the Engineering Result

The evaluation is strongest when it is read as a chain of bounded claims rather than as a single pass or fail judgement. At the innermost layer, unit and integration tests establish

that a specified input produces the specified state transition. At the next layer, release assembly and manifest inspection establish that the tested TypeScript is incorporated into an installable Android artefact with the recorded package properties. Physical-device interaction then establishes that selected routes render and selected controls respond on one hardware and operating-system combination. Revision 16’s visual matrix adds evidence for four final rendered states, not for unseen devices or users. None of these layers automatically establishes the next one. A correct function may be absent from a package, an installable package may render poorly, and a coherent interface may still be misunderstood by users.

This layered reading prevents a common reporting error in prototype research: promoting evidence because several individually positive checks appear together. Ten successful checks are not equivalent to one check of a broader proposition. For example, a passing threshold test and a successful device launch do not jointly prove that the threshold represents unlawful processing. They establish two different facts, one logical and one operational. Privacy Lens therefore records evidence by claim type. The resulting report may look more cautious than a conventional demonstration, but it is more useful to a maintainer because it shows exactly which additional study could change a conclusion.

The phrase “initial store-submission engineering candidate” is deliberately narrow. It means that the repository can produce the expected release formats, that the product has a coherent user journey and disclosures, and that no known in-scope engineering defect makes controlled submission work irrational. It does not mean that the application is signed with the owner’s production key, that a store has reviewed its declarations, or that the legal mappings have been approved by counsel. Separating these meanings is particularly important for privacy software because a polished security or compliance interface can borrow authority from the domain even when its evidence is limited.

There is also a difference between absence of an observed failure and affirmative evidence of a desirable property. No crash during a short device session is evidence about that session; it is not positive evidence of long-term stability. A small memory range across a few restarts is a useful anomaly check; it is not a leak proof. No sensitive runtime permission in the merged manifest is a strong package fact, but it does not by itself prove that every future acquisition connector will remain privacy preserving. The evaluation uses explicit negative claims to preserve these distinctions.

The result should therefore be understood as a reproducible decision point. A future reviewer can rebuild the same revision, inspect the same rules and fixtures, and ask whether later evidence warrants promotion from one claim class to another. That property is more durable than a one-time declaration of readiness. It converts the thesis from a retrospective narrative into a maintainable assurance record.

6.2 Technical Validity Boundaries

The central technical limitations concern the authority of acquisition, the meaning of prototype thresholds, temporal aggregation, and the governance of rule data. These issues are grouped because each can make a deterministic output reproducible while still leaving its real-world

interpretation uncertain.

6.2.1 Acquisition Authority and Observation Quality

The largest external-validity constraint is not the visual interface or the threshold arithmetic; it is the authority and completeness of the observation source. Android deliberately limits one ordinary application’s access to sensitive information about other applications. Privacy Lens does not bypass that security model. The current native bridge and imported-record contracts can carry observations when an authorised deployment supplies them, while the simulator supports deterministic explanation and testing. The prototype must not imply that installing it from a public store grants unrestricted historical visibility.

This constraint affects every downstream conclusion. An apparently quiet history may mean that accesses did not occur, that the selected API does not expose them, that an OEM shortened retention, that the collection job was delayed, or that the current deployment lacks authority. These explanations are not interchangeable. The evidence schema therefore includes source and time-window fields, and an evidence-gap state is preferred when acquisition sufficiency is unknown. A “no technical concern” result refers only to accepted records, not to all activity on the device.

Observation quality has at least five dimensions. Coverage asks which applications, operations, and time intervals are visible. Resolution asks whether records are event-level or aggregated. Timeliness asks how long after an event a record becomes available. Integrity asks whether records can be altered, replayed, or reordered. Semantic fidelity asks whether the operation label actually corresponds to the personal-data behaviour under review. A future connector should publish a capability document addressing all five dimensions. Merely conforming to the TypeScript shape is insufficient.

Aggregate records are attractive because they reduce storage and may reduce the privacy footprint of the reviewing application. They also erase context. Ten accesses could represent ten user-requested map refreshes or repeated background tracking; a total cannot distinguish them. Conversely, event-level traces can support better interpretation but may themselves reveal sensitive routines, application use, or location-related behaviour. The acquisition design must therefore minimise collected fields, bound retention, protect local storage, and explain who is authorised to inspect the evidence. More data is not automatically a more privacy-preserving audit.

An authorised enterprise or research deployment could improve coverage through device-owner controls, instrumented test applications, or controlled firmware. Each option changes the population to which the findings apply. Device-owner evidence may be relevant to managed fleets but not to ordinary consumers. Instrumented application evidence can establish ground truth for that application but not for unrelated packages. Research firmware may provide high-resolution traces but changes the operating environment. Evaluation reports should name the authority model rather than treating all Android evidence as equivalent.

Future acquisition validation should pair the connector output with an independent event generator. The generator should produce timestamped, permission-specific actions under

known foreground and background conditions. The study should compare generated and acquired events, report missing and duplicate records, measure timestamp error, and repeat the process across Android versions and OEMs. This would create an empirical coverage statement. Until then, the architecture supports authorised evidence, but the public-store product remains primarily a review and demonstration layer.

6.2.2 Threshold Meaning, Calibration, and Decision Error

The implemented baselines and multipliers are transparent prototype parameters. Transparency makes the rule reproducible; it does not make the parameters representative. A location threshold of thirty-six accesses in a twenty-four-hour record can be inspected and tested exactly, but its normative meaning depends on application category, user action, service duration, and declared purpose. A navigation application and a static calculator cannot be judged by the same behavioural expectation merely because both run on Android.

False positives and false negatives also have asymmetric consequences. A false positive may alarm a user, damage confidence in a legitimate application, and contribute to warning fatigue. A false negative may reassure the user or reviewer despite activity that deserves investigation. The appropriate balance is not a purely mathematical optimisation because the costs depend on use context. The current design therefore avoids converting a below-threshold result into a compliance statement and avoids converting an exceedance into a violation verdict.

Empirical calibration would require a sampling frame, not simply a large collection of counts. Applications should be stratified by category, operating mode, permission purpose, foreground status, and user task. Collection periods should capture weekdays, weekends, intermittent use, and long sessions. The study should document missing observations and OEM-specific retention. Independent reviewers would need an annotation protocol for deciding which cases warrant investigation, and disagreement should be reported rather than removed through silent consensus.

The calibration target should also be explicit. Optimising agreement with human reviewers is different from predicting store-policy enforcement, detecting malware, or identifying GDPR accountability questions. Each target requires different labels and evidence. Privacy Lens targets review prioritisation: the desired output is a defensible prompt for further examination. That target supports conservative wording and structured missing-context questions rather than a binary compliance classifier.

Multi-scale features improve coverage but add parameters. The daily excess rule identifies sustained volume, the burst rule identifies concentrated activity, and cross-window logic identifies repeated near-threshold behaviour. These signals may correlate, so simply summing them could exaggerate risk. A future scoring design should explain whether each signal contributes independently, whether one supersedes another, and how uncertainty changes when source resolution is coarse. It should also preserve the raw contributing values so a reviewer can reconstruct the decision.

Parameter governance is as important as parameter selection. Every released pack should identify its effective version, parameter owner, change rationale, fixtures, and compatibility

requirements. A threshold change can alter stored findings without any change in observed activity. The application should therefore retain the pack version used for a historical decision and should not silently relabel old evidence when a new pack becomes active. Re-evaluation, if offered, should create a new result linked to the prior one.

A responsible calibration programme would begin with controlled application traces and explicit scenarios, progress to a consented multi-device observational study, and only then consider population thresholds. Performance should be reported by category and source capability, with confidence intervals and error analysis. Even strong empirical performance would support prioritisation accuracy, not legal adjudication. That claim boundary should remain invariant as the data set grows.

6.2.3 Temporal Evidence and Replay Semantics

Time-windowed evidence introduces failure modes that are easy to overlook in static test cases. Two records can describe the same underlying accesses, adjacent periods, partially overlapping periods, or unrelated periods with identical totals. If overlapping aggregates are added, one event may be counted twice. If late records move a retention watermark backwards, replayed evidence can alter later decisions. If future timestamps are accepted without tolerance, clock errors can create impossible histories.

The current engine addresses these risks with safe integer checks, a maximum window duration, bounded future tolerance, replacement of identical windows, a finite per-key history, and non-overlap requirements for cross-window evidence. These are meaningful controls because they specify how the repository behaves under replay. They do not create event-level identity where none exists. Partially overlapping aggregates cannot always be reconciled without access to their constituent events.

The conservative response is to decline cumulative interpretation when independence cannot be established. This may reduce sensitivity, but it avoids manufacturing precision. A future event-level connector could attach stable event identifiers or monotonic collection cursors. A future aggregate connector could publish disjoint collection intervals and a watermark signed or otherwise protected by the producer. Both designs would strengthen deduplication, but each must account for device clock changes and interrupted collection.

Retention limits deserve separate scrutiny. Keeping only the most recent records bounds memory and reduces local exposure, yet an arbitrary truncation point can remove context needed for a decision. The history limit of 4,096 entries is an engineering safeguard, not a legal retention determination. Production policy should relate retention to the stated review purpose, source frequency, device capacity, and user deletion expectations. The Settings clear action should remove review data predictably, while audit exports or externally managed evidence require their own lifecycle.

Temporal testing should cover more than threshold boundaries. Useful properties include idempotence under exact replay, invariance to arrival order for independent windows, rejection of unsafe durations, monotonic retention watermarks, bounded storage, and stable decisions after process restart. Tests should also inject clock rollback, daylight-saving changes, long

device sleep, and records delivered after the active pack changes. These cases would reveal whether the implementation depends on wall-clock assumptions that are not documented.

The broader lesson is that time is part of the evidence contract, not metadata attached after classification. A risk statement about repeated behaviour is only as defensible as the window relationships used to compute it. Making those relationships visible helps reviewers distinguish sustained behaviour from an artefact of aggregation.

6.2.4 Regulation-Pack Governance Beyond Software Modularity

The replaceable pack contract solves a software architecture problem: it prevents jurisdiction-specific references, prompts, and thresholds from being scattered across screens, storage, and native integration code. It does not solve the institutional problem of deciding what a pack should say. A pack can be syntactically valid and legally misleading. Consequently, extensibility must be paired with admission and maintenance governance.

Admission should begin with scope. The pack must name the jurisdiction, regulated context, legal instruments, effective date, intended users, and exclusions. It should distinguish binding provisions from regulator guidance and from research heuristics. Every rule should link a technical condition to a review question and explain the missing facts that prevent automated legal determination. A reviewer should be able to trace displayed wording to the pack source and from the source to the governing material.

Revision 15 operationalises part of this distinction by recording each source as binding law, final guidance, or consultation material. This prevents the 2026 EDPB DPIA template and Article 6(1)(f) consultation text from silently acquiring the same status as the Regulation or final guidance [14, 16]. The status is still maintainer-authored metadata. Without signed provenance, source retrieval, content hashing, and qualified review, the validator can enforce consistency but cannot establish authenticity or legal correctness.

Legal review should be independent of the developer who encoded the rule. The reviewer should assess citation accuracy, applicability, exceptions, terminology, translations, and the risk that interface wording implies a verdict. Approval should apply to a specific pack version and date. Where the law or guidance changes, the registry should deprecate the old version, explain the transition, and preserve historical identity for prior findings.

Technical admission should then validate identifiers, semantic versions, rule uniqueness, supported evidence fields, bounded numerical parameters, localisation completeness, and fixture coverage. The application should reject an unknown or malformed pack rather than fall back silently to a default. If a selected pack is unavailable after an update, the interface should present an explicit blocked state and retain the user's selection for recovery.

Regional variation is not merely a change of article numbers. Concepts such as controller, covered business, sensitive information, consent, sale, sharing, child, or data localisation may not map one-to-one. Enforcement actors and private rights may differ. A common interface model should therefore contain generic evidence and review primitives while allowing packs to declare jurisdiction-specific questions. Forcing every regime into GDPR terminology would create superficial flexibility and substantive error.

Localisation also requires more than translating strings. Legal terms may have authoritative language versions, while plain-language prompts need comprehension testing in the user’s locale. Text length affects layout, and right-to-left scripts affect navigation and icon direction. A pack release should therefore combine legal-language review, plain-language review, and rendered interface testing. The interface should identify both the selected region and pack version so the user understands which framework generated the prompt.

Distribution is another trust boundary. If packs are ever downloaded, the system will need authenticity, integrity, rollback, and availability controls. Signed manifests, pinned publisher identities, checksum verification, and an offline last-known-good version are plausible mechanisms. Remote updates should not introduce executable code; the contract should remain declarative and narrowly validated. The present bundled registry avoids network update risk, which is appropriate for the current candidate.

The Research Baseline pack illustrates an important disclosure pattern. It proves that the product can switch rule sets and preserve identity, but it intentionally does not masquerade as a validated regional law. Its name and explanatory text should remain visibly non-legal. A future second jurisdiction should only be advertised after completing the same source, review, fixture, localisation, and update-governance process as the GDPR pack.

6.3 Human, Ethical, and Operational Trust

The product’s credibility depends on more than a correct rule. Users must understand the status, the reviewing application must minimise its own privacy and security risks, research evidence must be collected responsibly, and release claims must match the package that is actually distributed.

6.3.1 Trust, Comprehension, and Willingness to Use

Visual coherence is necessary because users make rapid judgements about whether a privacy tool is maintained and safe. Consistent typography, restrained colour, stable navigation, and plain-language hierarchy reduce avoidable friction. They are not sufficient for trustworthiness. A visually polished product can be more dangerous if polish causes users to accept unsupported conclusions. Privacy Lens therefore treats calibrated trust, rather than maximum trust, as the design objective.

Calibrated trust means that confidence changes with evidence quality. A native or authorised imported observation can be described according to its documented capability. A simulator record must remain labelled synthetic. Missing purpose or lawful-basis context must remain visible even when the numerical signal is normal. The interface should make the most important limitation available at the decision point, not bury it exclusively in a policy page.

Willingness to use has several components. Initial adoption depends on a credible value proposition and low setup burden. Continued use depends on whether findings arrive at useful times, whether repeated alerts are suppressed, whether explanations support action, and whether local data controls behave predictably. Recommendation to others may depend on institutional endorsement, independent review, and evidence that the tool does not create

new privacy risks. A short expert inspection can assess obvious barriers but cannot establish these behavioural outcomes.

A participant study should therefore use realistic tasks rather than preference ratings alone. Participants could identify the source of a finding, compare two risk states, explain why a low-count record is not a compliance certificate, switch regulation packs, clear local evidence, and decide what follow-up information is needed. Measures should include completion, error, time, comprehension, perceived workload, appropriate reliance, and confidence. Interviews can reveal whether terms such as evidence gap or technical concern match users' mental models.

The sample should include ordinary Android users, privacy-aware users, accessibility participants, and professional reviewers. Their goals differ. Ordinary users may need a concise decision path; professionals may need exportable provenance and exact rule details. Accessibility participants may expose focus order, verbosity, contrast, magnification, or motor-control barriers that visual review misses. The study should avoid treating one group's preference as universal.

Longitudinal testing is necessary for notification design. A single session cannot reproduce fatigue or habituation. A multi-week study could compare immediate alerts, daily summaries, and state-change-only notifications. It should record dismissals and follow-up actions while avoiding collection of unnecessary behavioural telemetry. The hypothesis is not that fewer notifications are always better, but that each interruption should correspond to a meaningful change.

Trust measures also need a correct-answer component. Asking whether participants trust the application can reward unjustified confidence. Better evaluation presents cases with strong evidence, synthetic evidence, missing context, and malformed input, then tests whether confidence tracks those differences. An effective design should sometimes cause the user to withhold judgement. That is a successful outcome for an evidence-bounded review tool.

6.3.2 Privacy and Security of the Reviewing Application

A privacy-review application can itself become a source of sensitive data. Permission histories may reveal health routines, mobility, communication, or working patterns even when they contain no content. Application identifiers can expose relationships or interests. Findings and notification text may be visible to other people with physical access to the device. The product's local-first design reduces network disclosure, but local processing is not automatically risk free.

The current package requests no sensitive runtime permission and disables Android backup. These are valuable minimisation facts for the evaluated build. The application also provides a clear-data control and does not require an account. Future connectors could change this profile. Every new acquisition or synchronisation feature should therefore trigger a new data-flow inventory, manifest reconciliation, threat review, privacy notice update, and store declaration review.

At-rest protection should be proportionate to the threat model. Platform sandboxing protects data from ordinary applications, while device compromise, debug builds, rooted

environments, and unlocked backups present different risks. If future deployments store event-level histories or legal annotations, encryption and key lifecycle may become appropriate. Encryption does not eliminate exposure through rendered screens, notifications, logs, or exported files, so those channels also require controls.

Logging is a frequent source of accidental disclosure. Production builds should avoid writing raw evidence, application inventories, or legal-context notes to Logcat. Crash reports, if introduced, should be opt-in where required, minimise payloads, document processors and retention, and avoid attaching the evidence database. Analytics should not be added merely to measure engagement; a privacy tool carries a higher burden to justify behavioural tracking.

Pack supply-chain risk must also be considered. A malicious or compromised pack could present alarming text, suppress review, or direct users to unsafe resources even without executable code. Declarative schema restrictions, signed distribution, trusted publisher identities, URL allow-lists where appropriate, and human review reduce this risk. Bundled packs should be covered by the same source-control and release-review process as application code.

Abuse scenarios include importing fabricated evidence to defame an application, replaying old records as current, selecting a misleading jurisdiction, and showing a screenshot without its source label. Privacy Lens addresses some of these risks through provenance, timestamps, pack identity, and replay controls. It cannot authenticate every external producer in the current prototype. Export formats should preserve those fields and should state that the document is a review artefact rather than a regulator finding.

Security maintenance is part of store readiness. Revision 14 removed inherited Internet and network-state permissions and reduced npm audit results from 30 findings including two critical items to 20 findings with no critical item. The remainder is concentrated in Expo/Metro build-time image parsing and Apple configuration dependency paths, for which npm proposes a breaking framework transition. The absence of application networking limits on-device reachability but does not erase supply-chain exposure during development. Dependencies, Android target requirements, signing keys, build provenance, and vulnerability disclosures therefore require continuing ownership and a separately validated Expo upgrade.

6.3.3 Research Ethics and Responsible Evaluation

The safest current evaluation uses synthetic or controlled evidence because real permission histories can expose third-party and participant behaviour. Synthetic data provides precise ground truth and repeatability, but it must be labelled so screenshots and metrics are not mistaken for field evidence. The simulator label is therefore an ethical as well as methodological control.

A future participant or device study would require informed consent describing what is collected, why it is needed, how long it is retained, who can access it, and how withdrawal operates. Researchers should avoid collecting unrelated application inventories or fine-grained histories when controlled test applications can answer the research question. Test accounts and scripted actions should be preferred for validation of acquisition correctness.

The legal subject matter creates an additional risk of participant misunderstanding. Study

materials should explain that findings are prompts and that the research team is not providing individual legal advice. Participants should not be encouraged to confront application developers or change essential device settings solely because of a synthetic warning. Debriefing should clarify which cases were fabricated and how the evidence boundaries work.

Reporting should include failed runs, missing observations, exclusions, and changes made after pilot testing. Selectively presenting only coherent screenshots would exaggerate maturity. The repository’s separation of logical, package, physical, and visual evidence provides a structure for ethical reporting because each claim can be tied to its acquisition method.

If external legal experts annotate cases, their independence, jurisdiction, instructions, and disagreement process should be recorded. A single expert opinion is valuable but should not be transformed into universal ground truth. Where reasonable interpretations differ, the pack can present the ambiguity as a question rather than forcing a binary label.

Finally, public release creates responsibilities to people who were not research participants. Store text and onboarding must not imply regulatory endorsement. Support channels should handle correction and deletion questions. Pack updates should communicate material interpretation changes. Responsible deployment is therefore a continuing governance activity, not the final administrative step after technical acceptance.

6.3.4 Store and Operational Readiness

The repository contains target-SDK configuration, AAB output, privacy-policy content, a stable package identifier, branded assets, and a store checklist. Google Play nevertheless requires owner-controlled signing and console information, including a Data safety declaration and current policy compliance [28,29]. The present AAB uses a QA signing fallback when no upload-key environment is provided. Public submission therefore remains blocked on an upload key, HTTPS privacy-policy publication, support contact, final listing content, screenshots selected for the listing, console declarations, internal-track testing, and store review.

6.4 Cross-Disciplinary Assurance Case

A high-quality privacy-engineering thesis must remain coherent when read by different expert communities. Table 10 summarises the strongest claim each discipline can accept from the current evidence, the residual weakness, and the evidence needed to advance it. This prevents one discipline’s success from compensating rhetorically for another discipline’s missing validation.

Table 10: Cross-disciplinary assurance case and next evidence

Perspective	Supported claim	Residual weakness	Evidence required next
Algorithms	The engine deterministically implements the stated admission, threshold, burst, replay, and pack rules	Oracle and implementation share the same threshold specification; parameters are not field-calibrated	Independent oracle, mutation/property tests, and a labelled external dataset
Software engineering	Source, release artefacts, manifest, package identity, hashes, and one device path are traceable	One OEM and short duration do not establish operational reliability	Multi-device matrix, 24-hour scheduling, performance, and upgrade/migration tests
Human-computer interaction	The interface implements progressive disclosure, non-colour cues, provenance, caveats, and direct actions	Expert inspection cannot establish comprehension, efficiency, satisfaction, or calibrated reliance	Task-based participant study, baseline comparison, and assistive-technology campaign
Law and governance	The EU pack exposes official anchors and missing legal facts while prohibiting automated infringement claims	Numerical prompts are engineering heuristics and the pack has no independent legal approval record	Qualified review, provision-level decision records, reviewer identity, effective date, and controlled update process
Research integrity	Evidence classes and negative claims are separated; historical versions and failures are retained	Evaluation was led within the project and does not provide full evaluator independence	External replication, preregistered protocol, and archived immutable evidence bundle

The table also clarifies what a high thesis score cannot legitimately mean. A polished architecture and a clean release can justify strong software-engineering marks, but they cannot manufacture participant data or legal authority. Conversely, explicit limitations are not a weakness when they are connected to a concrete validation programme; they become a weakness only when the central contribution depends on evidence that is entirely absent.

6.5 Prioritised Next Work

The next programme should proceed in this order:

1. obtain independent GDPR-pack review and define pack-update governance;
2. provision production signing and complete Play Console privacy and Data safety declarations;
3. run accessibility and localisation matrices, followed by task-based participant evaluation;
4. execute multi-OEM, multi-version, 24-hour scheduling, battery, memory, and recovery tests;
5. replace aggregate-only acquisition where authorised event-level evidence is available;
6. design and independently validate one additional regional pack before claiming multi-jurisdiction support.

This order prioritises validity and operational trust over adding more visible features.

7 Conclusion

This thesis set out to determine whether bounded Android permission evidence could support a useful privacy-review product without being misrepresented as automated legal judgement. The conclusion answers the four research questions, integrates the architectural and methodological contributions, states the final product position, and identifies the next evidence required. Operational deployment details already analysed in Chapter 6 are not repeated here.

7.1 Answers to the Research Questions

RQ1: deterministic and safe transformation. Permission summaries are accepted through a conservative evidence contract, rejected when identifiers, numbers, sources, windows, timestamps, or contexts are unsafe, and evaluated through a deterministic rule-pack engine. Temporal isolation addresses replay and overlap within the limits of aggregate windows.

The answer is supported by more than nominal examples. Runtime validation treats TypeScript declarations as developer guidance rather than as a security boundary. Permission identifiers must be owned values from the admitted set; numeric fields must be safe integers; windows must be ordered, bounded, and plausibly timed; and source values must match the closed provenance vocabulary. A rejected observation does not become a normal finding. This fail-closed behaviour is important because silent normalisation would create reassurance from evidence whose meaning is unknown.

Determinism also includes state. An identical record does not accumulate repeatedly, a late record cannot move the repository’s retention watermark backwards, and cross-window evaluation uses completed non-overlapping history. These controls make repeated test runs explainable and prevent several aggregate replay errors. The qualification remains that aggregates do not contain stable identities for their constituent events. The implementation can refuse unsafe combination; it cannot reconstruct information the producer did not supply.

RQ2: calibrated communication. Each finding preserves provenance, temporal scope, rule-pack identity, rationale, and missing-context questions. The interface distinguishes needs-review, evidence-gap, and no-technical-concern states, and repeatedly states that technical output is not a legal verdict. Synthetic demonstrations remain labelled throughout the data path.

Calibrated communication is implemented as an information architecture, not as a disclaimer appended to an otherwise absolute score. Overview explains the product’s current evidence and offers a clear next action. Findings organises review states and exposes detail progressively. Settings identifies the selected region and provides data controls. Within a finding, status, source, time, numerical reason, pack identity, and follow-up questions appear in an order intended to support both rapid triage and deeper inspection.

The communication result remains an expert assessment of the implemented design. The device render and interaction checks show that the path exists and is visually coherent on the tested handset. They do not prove that a population will understand the distinction between technical concern and legal compliance. A task-based participant study is required to establish

that claim. The thesis therefore answers RQ2 at the level of traceable interface design and bounded device evidence, not validated human comprehension.

RQ3: regional extensibility. Jurisdiction-dependent material is loaded from registered, versioned packs behind a stable contract. The EU GDPR pack is the evaluated legal objective; the Research Baseline demonstrates software switching without claiming another jurisdiction. New regional packs require their own legal and empirical validation.

This result separates two forms of extensibility that are often conflated. Software extensibility means that another admitted pack can supply identifiers, labels, thresholds, legal references, rationales, and questions without rewriting navigation, persistence, or Android integration. Substantive legal extensibility means that those contents are correct for a jurisdiction and use context. The project demonstrates the first and defines governance requirements for the second. It does not use a placeholder pack as evidence of comparative-law validity.

The contract also improves maintenance of the evaluated GDPR objective. A finding stores the pack and rule version that produced it, so a later rule update need not erase historical meaning. Registry loading validates mandatory authored, reviewed, effective-date, review-authority, review-scope, release-scope, locale, trigger, lifecycle, and source-status fields; inconsistent approval claims, impossible dates, guidance-only legal packs, and unknown identifiers fail closed. Basis-specific evidence questions prevent a lawful-basis label alone from producing an evidence-complete result. Explicit selection avoids hidden regional inference. This design provides a foundation for future jurisdictions, but every production pack still needs scoped sources, independent review, versioned fixtures, localisation, and an update policy.

RQ4: supported claims. Compilation, lint, deterministic and adversarial suites, release APK/AAB assembly, merged-manifest inspection, physical installation, cold launch, and rendered workflows passed for the recorded candidate. These results support an initial store-submission engineering candidate. They do not prove legal compliance, unrestricted Android visibility, long-run reliability, broad usability, accessibility conformance, production signing, or store acceptance.

The answer to RQ4 is therefore a claim map rather than a maturity slogan. Logical suites support engine conformance to specified rules. Package inspection supports statements about identifiers, SDK targets, permissions, backup configuration, and incorporated resources. Physical checks support only the recorded handset, build, and interaction interval. Visual review supports an expert judgement that the candidate is coherent enough for controlled submission work. External owner actions and broader empirical studies remain explicit gates.

This structure makes negative results useful. If a future OEM test misses observations, the acquisition claim can be narrowed without invalidating the regulation-pack architecture. If participants misunderstand a label, the presentation layer can change without rewriting historical evidence. If legal review changes a mapping, a new pack version can be admitted without silently rewriting old findings. The research questions are therefore connected by contracts while retaining separable evidence.

7.2 Contributions and Research Significance

The first contribution, the evidence-bounded architecture, establishes the governing principle: every conclusion must retain the source, temporal scope, rule identity, and known missing context needed to interpret it. This principle influences data types, validation, storage, interface language, testing, and release documentation. It is not a wrapper around the implementation. Without it, individual technical features could still work while the product as a whole overstated its authority.

The second contribution, the fail-closed decision pipeline, operationalises that principle at the point where data enters the system. A permissive pipeline would transform malformed or inherited values into ordinary results, creating an especially harmful false sense of safety. Closed identifiers, safe numbers, bounded windows, explicit sources, replay semantics, and structured evidence gaps ensure that uncertainty remains visible. The pipeline is intentionally deterministic so an auditor can reconstruct why a record was accepted and why a finding changed.

The third contribution, the regulation-pack contract, locates jurisdictional interpretation in a reviewable component. This reduces coupling and clarifies the boundary between generic signal processing and legally motivated questions. The contract includes identity and versioning because a rule without provenance cannot support historical accountability. The contribution is the architecture and governance model, not a claim that arbitrary legal systems can be translated automatically.

The fourth contribution, the local-first interface, turns the internal evidence model into a product experience. The three-destination structure removes experimental navigation and emphasises current state, actionable review, and controls. Status is not encoded by colour alone. Synthetic provenance and claim limits are presented near relevant decisions. Local storage, no account requirement, no sensitive runtime permission, disabled backup, and a clear-data action align the product behaviour with its privacy message.

The fifth contribution, the verification record, connects artefacts to claims. Deterministic fixtures, adversarial boundary cases, build outputs, manifest facts, device screenshots, and store-readiness gaps are not collapsed into one test total. This makes the evaluation auditable and limits accidental overstatement. It also defines a repeatable path for future releases: every material change should identify which evidence layers must be renewed.

Methodologically, these contributions form a design-science result: the artefact embodies a solution principle, demonstration shows that the principle can be realised, and evaluation tests the artefact against explicit objectives [9, 10]. The resulting knowledge claim is conditional rather than universal. Within the recorded evidence sources and Android release candidate, preserving provenance and missing context makes the warning path more auditable. Whether that architecture improves comprehension, decisions, or organisational accountability remains an empirical question for the next study.

Together, the contributions support a form of trustworthy incompleteness. Privacy Lens does not solve every acquisition, legal, or usability problem, but it makes those unsolved parts structurally difficult to hide. In an accountability tool, that property is itself useful. A system

that clearly refuses an unsupported verdict can better support responsible human review than a more ambitious classifier whose authority is unclear.

7.3 Final Product Position

Privacy Lens 1.2.0 previously received an internal 95.1/100 expert-proxy score under a release-engineering rubric. The stricter publication-target rubric resets the App-thesis predecessor to 70/100, assesses Revision 15 at 79/100, and assesses Revision 16 at 84/100. The difference is intentional: the strict rubric reserves substantial credit for independent legal review, participant and assistive-technology evidence, and publication-quality external replication. Version 1.4.0 uses version code 5, retains strengthened legal-evidence governance, exposes source authority in the interface, uses non-colour status structure, builds as ARM64 APK/AAB, passes package inspection, and completes five final cold-start rounds plus a four-screen visual matrix on the recorded OPPO device. None of these scores is a university mark or publication decision.

The candidate is not yet a publicly publishable production release. Production signing, stable public policy and support endpoints, console declarations, internal-track evaluation, and store review remain external owner actions. This distinction is part of the result rather than an administrative footnote: trustworthy privacy software must disclose where its evidence ends.

7.4 Implications for Privacy-Engineering Practice

The first lesson is that provenance must survive every transformation. Capturing a source label at ingestion is insufficient if storage, notifications, exports, or screenshots omit it. Provenance should be treated as part of the value being computed. Tests should therefore follow it from observation through decision and presentation.

The second lesson is that missing context requires a first-class state. Many privacy questions cannot be answered from technical frequency. Returning a neutral or positive result when purpose, authority, or acquisition coverage is unknown creates hidden assumptions. An evidence-gap state makes the limitation actionable: it can ask for a privacy notice, controller record, consent basis, or authorised trace rather than fabricating certainty.

The third lesson is that extensibility needs governance. An interface or schema that can load new legal text is technically flexible but may accelerate error if sources, versions, review, and updates are uncontrolled. The pack boundary is valuable precisely because it creates a unit that can be admitted, tested, rejected, and deprecated.

The fourth lesson is that polished design should reduce ambiguity rather than conceal it. Visual hierarchy can foreground evidence strength, state, and next action. Calm colour and concise language can lower cognitive burden. Yet critical caveats must remain visible. The ethical objective is not to maximise engagement or confidence; it is to help the user reach an appropriately bounded understanding.

The fifth lesson is that store readiness is an evidence package. An AAB is necessary but not sufficient. Signing ownership, manifest reconciliation, data declarations, policy endpoints, screenshots, accessibility, update testing, and support processes all affect whether

the distributed product matches the evaluated one. Treating these as afterthoughts would break the thesis’s claim chain at the point of deployment.

The final lesson is that disciplined limitation can increase, rather than reduce, practical value. A reviewer can act on a clearly sourced question, ask for missing evidence, and revisit the decision after a pack update. An unsupported compliance badge offers none of those affordances. Privacy Lens is most credible when it remains specific about what it observed, how it reasoned, and what a human must still decide.

7.5 Future Work

Future work should validate the GDPR pack independently, add a genuinely reviewed regional pack, improve authorised event-level acquisition, test multiple Android implementations over long durations, and conduct accessibility and participant studies. These activities should retain the Revision 16 claim-to-evidence discipline and its separation between implemented affordances and observed human outcomes. Additional features are valuable only when their authority, provenance, and validation remain visible.

The programme should also investigate formal and property-based techniques for evidence admission and replay. Useful properties include idempotence, bounded storage, monotonic watermarks, source isolation, and stability under reordered independent windows. Mutation testing could assess whether the suite detects weakened validation. A separately implemented oracle would reduce the correlated-defect risk of deriving implementation and expected values from one specification.

Interface research should compare alternative presentations of uncertainty. For example, a three-state label, an evidence-strength scale, and a question-led presentation may produce different reliance patterns. The comparison should measure comprehension and decisions, not only preference. Accessibility evaluation should be incorporated from recruitment through reporting rather than added as a final conformance scan.

The pack framework could support signed declarative updates after its threat model is tested. This would require publisher identity, schema versioning, offline recovery, rollback protection, and clear user communication. Remote executable rules would introduce unnecessary risk; the safest extension preserves a narrow validated data contract.

Finally, future scholarship should study the relationship between accountability prompts and organisational action. A technically clear question has value only if users or reviewers can obtain the missing policy, consent, retention, or controller evidence. Research with data-protection practitioners could examine which prompts accelerate review, which create noise, and how findings should integrate with established assessment records without becoming an unreviewed legal decision.

7.6 Closing Statement

The central finding is that privacy accountability on Android is strengthened not by presenting the strongest possible conclusion, but by presenting the strongest conclusion the available evidence actually supports. Privacy Lens operationalises that principle in code, interface,

evaluation, and thesis structure.

References

- [1] European Parliament and Council of the European Union, “Regulation (eu) 2016/679 of 27 april 2016,” 2016.
- [2] European Data Protection Board, “Guidelines 4/2019 on article 25 data protection by design and by default,” 2020.
- [3] H. Nissenbaum, “A contextual approach to privacy online,” *Daedalus*, vol. 140, no. 4, pp. 32–48, 2011.
- [4] P. Wijesekera, A. Baokar, A. Hosseini, S. Egelman, D. Wagner, and K. Beznosov, “Android permissions remystified: A field study on contextual integrity,” in *24th USENIX Security Symposium*, pp. 499–514, USENIX Association, 2015.
- [5] W. Cao, C. Xia, S. T. Peddinti, D. Lie, N. Taft, and L. M. Austin, “A large scale study of user behavior, expectations and engagement with android permissions,” in *30th USENIX Security Symposium*, USENIX Association, 2021.
- [6] S. Prange, P. Knierim, G. Knoll, F. Dietz, A. D. Luca, and F. Alt, “I do (not) need that feature: Understanding users’ awareness and control of privacy permissions on android smartphones,” in *Twentieth Symposium on Usable Privacy and Security*, pp. 453–472, USENIX Association, 2024.
- [7] A. M. McDonald and L. F. Cranor, “The cost of reading privacy policies,” *I/S: A Journal of Law and Policy for the Information Society*, vol. 4, p. 543, 2008.
- [8] J. A. Obar and A. Oeldorf-Hirsch, “The biggest lie on the internet: Ignoring privacy policies and terms of service anymore?,” *Information, Communication & Society*, vol. 23, no. 1, pp. 128–147, 2018.
- [9] A. R. Hevner, S. T. March, J. Park, and S. Ram, “Design science in information systems research,” *MIS Quarterly*, vol. 28, no. 1, pp. 75–105, 2004.
- [10] K. Peffers, T. Tuunanen, M. A. Rothenberger, and S. Chatterjee, “A design science research methodology for information systems research,” *Journal of Management Information Systems*, vol. 24, no. 3, pp. 45–77, 2007.
- [11] S. Brooks, M. Garcia, N. Lefkowitz, S. Lightman, and E. Nadeau, “An introduction to privacy engineering and risk management in federal systems,” Tech. Rep. NISTIR 8062, National Institute of Standards and Technology, 2017.
- [12] K. Boeckl and N. Lefkowitz, “Nist privacy framework: A tool for improving privacy through enterprise risk management, version 1.0,” Tech. Rep. NIST CSWP 01162020, National Institute of Standards and Technology, 2020.
- [13] European Data Protection Board, “Guidelines 05/2020 on consent under regulation 2016/679,” 2020. Final guidance; accessed 10 August 2026.
- [14] European Data Protection Board, “Guidelines 1/2024 on processing of personal data based on article 6(1)(f) gdpr,” 2024. Public-consultation version, not cited as final guidance; accessed 10 August 2026.
- [15] European Data Protection Board, “Endorsed wp29 guidelines, including transparency and data protection impact assessment,” 2018. Accessed 10 August 2026.

- [16] European Data Protection Board, “Template for data protection impact assessment,” 2026. Consultation material pending finalisation; accessed 10 August 2026.
- [17] European Data Protection Board, “Guidelines 03/2022 on deceptive design patterns in social media platform interfaces,” 2023.
- [18] Android Developers, “Permissions on android,” 2026. Accessed 9 August 2026.
- [19] M. Egele, C. Kruegel, E. Kirda, and G. Vigna, “Pios: Detecting privacy leaks in ios applications,” in *Proceedings of the 18th Network and Distributed System Security Symposium (NDSS)*, 2011.
- [20] Android Developers, “Appopsmanager api reference,” 2026. Accessed 9 August 2026.
- [21] International Organization for Standardization, *ISO 9241-11:2018 Ergonomics of Human-System Interaction—Part 11: Usability: Definitions and Concepts*. Geneva: ISO, 2018.
- [22] World Wide Web Consortium, “Web content accessibility guidelines (wcag) 2.2,” 2024.
- [23] Android Developers, “Principles for improving app accessibility,” 2026.
- [24] J. Sweller, P. Ayres, and S. Kalyuga, *Cognitive Load Theory*. Springer Science & Business Media, 2011.
- [25] D. Kahneman, *Thinking, Fast and Slow*. Farrar, Straus and Giroux, 2011.
- [26] J. D. Lee and K. A. See, “Trust in automation: Designing for appropriate reliance,” *Human Factors*, vol. 46, no. 1, pp. 50–80, 2004.
- [27] Android Developers, “About android app bundles,” 2026. Accessed 9 August 2026.
- [28] Google Play Console Help, “Meet google play’s target api level requirement,” 2026. Accessed 9 August 2026.
- [29] Google Play Console Help, “Provide information for google play’s data safety section,” 2026. Accessed 9 August 2026.