

Bridging the Semantic and Control Gaps: A Human-Centric Privacy Architecture for Cross-Border mHealth Regulations.

Zhiheng Zhang

Student ID: 54560484

Supervisor: Richard Sinnott

Word Count: 5931

April 29, 2026

Contents

1	Abstract	1
2	Introduction and Background	2
2.1	Evolution of the mHealth Ecosystem	2
2.2	The Transparency Paradox and the Illusion of Control	2
2.3	Architectural Lineage: From Spatial De-identification to Semantic Control	3
2.4	Regulatory Fragmentation and the 2026 Landscape	4
2.5	Proposed Solution: The Privacy-Enhancing Architecture (PEA)	4
2.6	Socio-Technical Roots and Dark Patterns	4
3	Global Privacy Regulation: A Three-Party Technical and Philosophical Comparison	5
3.1	Regulatory Fundamentals and Legal Philosophy	5
3.2	Mechanism of Consent: From Informing to Verifying	5
3.3	Data Subject Rights and the “Verification Gap”	6
3.4	Data Portability and Interoperability (2026 Mandates)	6
3.5	Legal Penalties and Architectural Accountability	6
4	Literature Review	8
4.1	The Socio-Technical Roots of Privacy Resignation	8
4.2	The Evolution of Informed Consent	8
4.3	Cognitive Load Theory and Decision-Making	8
4.4	Information Visualization and Pre-attentive Processing	9
4.5	Spatial Privacy and Attribute Inference	9
4.6	The Control Paradox and Native-Layer Enforcement	9
4.7	Trust Calibration and the Digital Divide	9
5	Research Gaps and Problem Statement	11
5.1	Gap 1: UI/UX Semantic and Visual Overload	11
5.2	Gap 2: The Control Paradox and Verification Gap	11
5.3	Overarching Problem Statement	11
5.4	Theoretical Framework and Conceptual Model	12
6	Research Questions	13
6.1	RQ1: Psychological Load and Cognitive Mitigation	13
6.2	RQ2: Dual-Layer Trust and Perceived Control	13
6.3	RQ3: Interaction Efficiency and Decision-Making Paradigms	13

7	Significance and Implications of the Study	14
7.1	Theoretical Contribution: Modeling Cognition and Control	14
7.2	Methodological Significance: Standardizing Evaluation	14
7.3	Practical and Industry Implications	15
8	Research Methodology and Technical Architecture	16
8.1	Phase 1: Privacy-Enhancing Architecture (PEA) Implementation	16
8.1.1	Technical Justification and Comparison of Alternatives	17
8.1.2	Native Interceptor Middleware and OS-Verification Bridge	17
8.1.3	LDP Noise-to-Visual Transformation Pipeline	19
8.2	Phase 2 & 3: HCI Evaluation and Statistical Modeling	20
9	Research Schedule and Milestones	21
10	Evaluation and Data Analysis	22
10.1	User-Centric Evaluation: Experimental Design	23
10.2	Technical Evaluation: Strong Consistency and Performance	23
10.3	Psychometric Quantification via NASA-TLX	25
10.4	Statistical Robustness and PLS-SEM Modeling	25
10.5	Hypotheses and Boundary Conditions	25
11	Research Limitations	26
11.1	Sample Representativeness and the Digital Divide	26
11.2	Cognitive Interpretation and Pre-attentive Efficiency	26
11.3	The Verification-Friction Trade-off	26
11.4	Subjectivity in Psychometric Measurement	26
11.5	Ecological Validity and Experimental Context	26
11.6	Platform Specificity and Architectural Stability	27
11.7	Regulatory Volatility and Compliance Gaps	27
11.8	Architectural Compatibility with HarmonyOS NEXT	27
12	Conclusion	27
13	Proposed Thesis Structure and Outline	29
13.1	Chapter Content Summary and Word Distribution	29

1 Abstract

The global mobile health (mHealth) sector currently operates under a “Fractured Compliance Model” where legal requirements remain decoupled from technical execution [1], [2].

This market is projected to reach \$403.02 billion by 2031 [3]. However, the proliferation of advanced sensors has created a “Transparency Paradox”: excessive visual noise triggers privacy fatigue, leading users — cognitively overwhelmed — to resort to “blind agreement” [4], [5]. Consequently, users receive notifications they lack the cognitive utility to process.

To address this gap, this thesis presents the Privacy-Enhancing Architecture (PEA) — a middleware system that synchronizes human cognition with native-layer enforcement mechanisms. Developed using the React Native framework, the system features a Skia-driven dynamic dashboard that translates Local Differential Privacy (LDP) parameters into “Blurry Path” visualizations to address the “Semantic Gap.”

By incorporating a native-layer Privacy Interceptor Middleware and an OS-Verification Bridge, the architecture implements a dual-layer design that addresses the Control Paradox through kernel-level API pipeline management. This approach leverages smartphone OS settings as an external “Ground Truth” to restore user trust.

A Randomized Controlled Trial ($n \approx 100$) will be conducted to evaluate the proposed framework. This study utilizes the NASA-TLX scale [6] to quantify mental demand and frustration levels. Furthermore, Partial Least Squares Structural Equation Modeling (PLS-SEM; [7]) will test the moderating effect of cognitive load on Perceived Behavioral Control (PBC), grounded in Privacy Calculus Theory [8].

By translating abstract legal mandates into quantifiable technical controls, this thesis bridges the divide between privacy policy and user experience — delivering a verifiable Privacy-by-Design framework for complex medical data ecosystems.

2 Introduction and Background

2.1 Evolution of the mHealth Ecosystem

The mobile health (mHealth) sector has undergone a significant structural evolution driven by the ubiquitous adoption of smartphone-integrated biosensors. Empirical data indicates that innovation in this field peaked between 2016 and 2021. By 2024, the sector reached a *technical saturation point* [9], shifting the focus from general health tracking to specialized clinical interventions.

The market is projected to reach \$403.02 billion by 2031, representing a compound annual growth rate (CAGR) of 25.4% [3], [9]. Put differently, mHealth is transitioning from a consumer utility to a foundational component of digital healthcare delivery. This expansion necessitates a more robust examination of how sensitive biological metrics are handled within these high-frequency data environments.

2.2 The Transparency Paradox and the Illusion of Control

Despite rapid market expansion, existing mHealth privacy safeguards remain inadequate [5], [10]. Users are routinely required to share “special category data”, such as hormone fluctuations and genetic sequences, which are protected under frameworks like the GDPR [1]. However, existing authorization models often provide only an “illusion of control”.

In standard application-layer architectures, a user’s decision to revoke consent frequently fails to synchronize with native-layer sensor execution. Research reveals that background data harvesting often continues even after UI-level revocation [11]. This discrepancy results in a “Transparency Paradox”: empirical studies indicate that 63–97% of users bypass lengthy text agreements depending on interface complexity, default settings, and device type [12], [13]. In mHealth contexts specifically, the granularity of sensitive data requests amplifies this heuristic behavior[14].

This fractured compliance model and the resulting cycle of privacy fatigue, which necessitates the proposed PEA solution, is visually summarized in Figure 1.




Figure 1: The Fractured Compliance Model and the Proposed PEA Solution Path.

2.3 Architectural Lineage: From Spatial De-identification to Semantic Control

The technical challenges identified in mHealth build upon a decade of research into spatial and attribute-based privacy. For example, earlier exploration of spatial privacy within the Melbourne region demonstrated that sensitive location identities could be identified using machine learning on aggregated geo-tagged data [15]. Sun (2015) demonstrated that user behavioral patterns could be extracted from seemingly benign spatial metadata.

Building upon these spatial risks, research by Altaf (2017) showed that location-specific attributes could be inferred with 100% accuracy from discretized social media timelines [16]. In effect, an adversary can re-identify individual status changes if the privacy budget is not strictly enforced via robust noise models. This highlights a fundamental vulnerability in systems that rely on manually-calibrated de-identification.

This research lineage reached a significant milestone with the introduction of a three-layer semantically-driven access control framework for health record linkage [17]. Lu established the foundational logic for

semantic reasoning in secure data exchange. However, a critical discrepancy remains in the psychological accessibility of these semantic controls for non-technical users. This project seeks to bridge this “Semantic-Interaction Gap” through real-time native-layer enforcement.

2.4 Regulatory Fragmentation and the 2026 Landscape

By 2026, the global mHealth landscape is governed by diverse legal frameworks, including the GDPR, PIPL, and the Australian Privacy Principles. While these regulations mandate “ease of withdrawal,” the industry faces a significant “Compliance Gap”. Research as of 2026 indicates that 42.3% of health applications remain non-compliant due to the structural disconnect between mandates and execution [18].

This fragmentation is exacerbated by the 2026 APP 1.7 amendments in Australia and the European Health Data Space (EHDS) framework [2], [19]. These mandates impose mandatory transparency for Automated Decision-Making (ADM) and granular interoperability, compelling mHealth providers to adopt a unified architecture that can handle divergent regional requirements without increasing cognitive friction.

2.5 Proposed Solution: The Privacy-Enhancing Architecture (PEA)

This thesis presents the Privacy-Enhancing Architecture (PEA) — an end-to-end framework that unifies cognitive clarity with physical-layer assurance. At its core, the PEA integrates two components: a Skia-driven dynamic dashboard and a native-layer *OS-Verification Bridge*. The dashboard provides immediate, verifiable translation of user intent into native-layer API execution.

By establishing an external “Ground Truth” that resolves the *Control Paradox*, the PEA achieves data sovereignty that is both psychologically accessible and technically absolute. Rather than optimizing the user interface at a superficial level, this architecture implements a robust, verifiable audit mechanism at the native layer.

2.6 Socio-Technical Roots and Dark Patterns

The mismatch between user intent and action is frequently categorized as the “Privacy Paradox.” In the mHealth context, this is often exacerbated by “Dark Patterns”—manipulative UI designs that nudge users toward over-sharing. As biosensor data like continuous PPG tracking becomes more complex, traditional “Notice and Choice” models become cognitively hazardous.

The proposed PEA framework seeks to dismantle these patterns by replacing manipulative designs with objective architectural guarantees. It seeks to restore the user’s sense of agency by providing high-fidelity visual feedback. This transforms the relationship between the user and the system from a black-box honor system into a verifiable, resilient digital health ecosystem.

3 Global Privacy Regulation: A Three-Party Technical and Philosophical Comparison

In the contemporary 2026 landscape, the regulatory environment for mHealth technologies is defined by three primary frameworks: the GDPR, PIPL, and the Australian Privacy Act. These frameworks present divergent focus areas that stem from unique regional political, cultural, and historical contexts [1], [20], [21]. For developers and service providers, compliance has transitioned from a peripheral legal requirement to a core architectural constraint. Rather than treating legal mandates as external patches, this thesis integrates them into the fundamental system design.

This section analyzes these frameworks through the lens of the “Verification Gap” and the architectural requirements they impose on the proposed PEA. By evaluating the discrepancy between legal claims and technical execution, the analysis identifies necessary system-level interventions. The goal is to ensure that privacy remains a functional reality rather than a legal promise.

3.1 Regulatory Fundamentals and Legal Philosophy

The philosophical foundations of these frameworks shape the “Mental Model” a system must support to achieve perceived legitimacy. The EU GDPR is rooted in the philosophy of “Rights and Dignity,” viewing data protection as a fundamental human right [1]. From an HCI perspective, this demands a “Rights-Orientation” where users maintain active sovereignty over their entire data life-cycle, governing how their health data is processed and stored at every stage.

In contrast, China’s PIPL prioritizes “Security and Order,” integrating individual protection with national data sovereignty [20]. Its “Separate Consent” mandate creates a high-density interaction requirement that may lead to “Visual Overload” if the system fails to utilize dynamic visualization. To prevent “Privacy Fatigue,” the interface must streamline the authorization process for sensitive data items.

Historically “Principle-Based,” the Australian framework balanced utility with individual privacy [21]. However, the 2024–2026 reforms have shifted the Australian Privacy Principles (APP) toward a “Rights-Based” model [2]. This shift introduces mandatory transparency for Automated Decision-Making (ADM), requiring granular explanation of algorithmic logic. Such transparency necessitates that systems offer users clear insight into how biometric or health data influences automated health outcomes.

3.2 Mechanism of Consent: From Informing to Verifying

When collecting sensitive health data, such as reproductive metrics or biometrics, the three frameworks impose specific interaction granularities. Under GDPR Article 7.3, withdrawing consent must be as straightforward as granting it [1]. Current “deep-nested” settings frequently violate this principle by complicating the opt-out process. This functional imbalance undermines the spirit of “Ease of Withdrawal.”

To address this requirement, the PEA architecture provides a “One-Click Physical Severing” mechanism located directly on the primary dashboard. User intent to withdraw is physically enforced at the native system layer, ensuring data access terminates at the API level immediately upon request.

The PIPL mandates individual, itemized authorizations for sensitive processing and cross-border transfers [20]. Skia-driven visualizations can bridge the “Semantic Gap” by grouping these requirements into intuitive paths, translating complex legal authorizations into context-aware visual flows. This design satisfies the legal requirement while mitigating cognitive burden.

3.3 Data Subject Rights and the “Verification Gap”

All examined frameworks grant rights to erasure and portability, yet they suffer from a fundamental technical “Verification Gap.” Under GDPR Article 17 and PIPL Article 47, users possess the legal right to request total data destruction [1], [20]. In current application-layer models, users cannot confirm whether the “Delete” button triggers kernel-level severance. Authorization revocation may merely change a UI flag while data remains stored in the background.

The **OS-Verification Bridge** addresses this gap by enabling users to cross-reference an application’s claims with the smartphone’s native permission status. This bridge provides a “Ground Truth” that verifies whether revoked consent retains physical effect, resolving the trust deficit between users and service providers through cross-layer verification.

3.4 Data Portability and Interoperability (2026 Mandates)

The 2026 expansion of data portability rights transforms interoperability from a functional goal into an architectural constraint. The EU’s EHDS mandates the European Electronic Health Record Exchange Format (EEHRxF), while the Australian CDR shifts toward real-time FHIR standards [2], [19]. Adherence necessitates system designs capable of processing highly structured medical data formats as a baseline compliance requirement.

For portability rights to be meaningful, the dashboard must translate complex exchange formats into intuitive “Export Flows,” ensuring “Data Sovereignty” is a functional reality rather than a legal claim. Users must be able to move data between providers without technical friction. The PEA abstracts underlying technical complexity into a user-centric visual interface.

3.5 Legal Penalties and Architectural Accountability

Non-compliance carries severe financial risks. GDPR violations incur penalties up to 4% of global annual turnover; PIPL fines reach 5% of the previous year’s revenue; Australia imposes thresholds of \$50 million AUD or 30% of turnover for serious breaches [1], [2], [20].

In the 2026 landscape, application-level “honor systems” fail to meet institutional risk management

demands. These penalties redefine technical overhead as a necessary investment. The PEA middleware provides verifiable proof of compliance, serving as a pragmatic safeguard against legal and financial exposure.

Table 1: Regulatory Frameworks as HCI and Architectural Constraints

Dimension	China: PIPL	EU: GDPR	Australia: APP	Technical & HCI Impact
Consent Logic	“Separate Consent” for sensitive items and cross-border transfers.	“Explicit & Specific” action required for processing.	“Express Consent” mandated for ADM (APP 1.7, 2026).	Requires LDP Visualization to reduce interaction density.
Withdrawal	Right to post-revocation deletion; state-verified.	Withdrawal must be as easy as granting.	Mandatory easy-opt-out mechanism.	Requires Native Interceptor for physical API severance.
Verification	State-monitored auditability with periodic reporting.	Individual accountability with documented consent trails.	Mandatory Privacy Impact Assessment (PIA).	Uses OS-Verification Bridge for cross-layer consistency.
Interoperability	Refined interface standards for domestic data flows.	Mandatory EEHRxF exchange format for cross-border sharing.	Real-time FHIR-based CDR integration.	Visual “Export Flows” for user-facing data portability.

4 Literature Review

This review examines the intersection of human-computer interaction (HCI), regulatory compliance, and system architecture, establishing the theoretical foundations for the proposed Privacy-Enhancing Architecture (PEA)[5].

4.1 The Socio-Technical Roots of Privacy Resignation

The mismatch between user intent and action is frequently categorized as the “Privacy Paradox” [8]. However, mHealth users often suffer from “Privacy Resignation,” perceiving data management as beyond their control[14].

Resignation emerges as a rational response to the power asymmetry between users and service providers [10]. Dual-process theory suggests that current consent models predominantly engage “System 1” (automatic, heuristic processing) rather than “System 2” (deliberate, effortful reasoning), causing users to bypass rational decision-making when accessing services [22].

4.2 The Evolution of Informed Consent

Early works focused on spatial de-identification; however, the field has shifted toward semantically-driven access control [17]. Lu demonstrated that seemingly innocuous health data could erode privacy in decentralized environments, establishing the foundational logic for the “Semantic Gap” addressed in this study.

The concept of Dynamic Informed Consent (DIC) offers a patient-centered alternative to static notices [23]. Platforms such as ADDN and ATHENA enable granular data control through dynamic consent models [24], [25]. Despite these advances, existing implementations rely predominantly on static list-based interfaces that fail to provide sufficient cognitive clarity, producing “Granularity Fatigue” as itemized choices accumulate [14], [26].

In contrast to serial-processing models, this thesis shifts toward pre-attentive visual metaphors. Traditional DIC requires high-effort “System 2” processing to evaluate each permission item [22]. The Skia-driven “Blurry Path” leverages pre-attentive variables (blur, transparency) to communicate risk levels before conscious focus is applied, mitigating cognitive burden in mHealth ecosystems.

4.3 Cognitive Load Theory and Decision-Making

The efficacy of a privacy dashboard is constrained by the user’s cognitive bandwidth. Cognitive Load Theory (CLT) posits that excessive visual noise induces cognitive shortcuts, undermining informed consent[27]. Information visualization principles further suggest that well-designed displays can restore the user’s “Rational Privacy Calculus” without exceeding memory limits [28], [29].

The PEA addresses this challenge by mapping data flows to intuitive visual metaphors. Rather than presenting users with exhaustive textual disclosures, the system communicates privacy risk through visual

artifacts, reducing extraneous cognitive load while preserving decision-relevant information density.

4.4 Information Visualization and Pre-attentive Processing

The “Blurry Path” visualization draws on pre-attentive processing theory. Visual variables such as blur and transparency are detected by the perceptual system within 200–250 milliseconds, prior to focal attention [30], [31]. This pre-conscious detection enables rapid risk communication without requiring technical literacy.

Local Differential Privacy (LDP) provides mathematical guarantees for data obfuscation, yet remains cognitively opaque to end-users unfamiliar with epsilon (ϵ) parameters [32]. By visualizing LDP noise as visual blur, the system translates abstract privacy budgets into intuitive risk indicators, bridging the semantic gap between formal privacy mechanisms and user comprehension.

4.5 Spatial Privacy and Attribute Inference

Prior research demonstrated that user behavioral patterns can be identified from aggregated geo-tagged data [15]. Machine learning techniques can automatically re-identify sensitive location identities from seemingly benign spatial metadata. Altaf extended this work by showing that attributes such as house occupancy could be inferred with 100% accuracy from discretized social media timelines [16].

These findings highlight a fundamental vulnerability in systems relying on manually-calibrated de-identification: adversaries can exploit temporal correlations to infer individual status changes unless privacy budgets are strictly enforced via robust noise models. The pre-attentive metaphors proposed in this study aim to communicate these risks transparently, enabling users to make informed decisions before data is conceded.

4.6 The Control Paradox and Native-Layer Enforcement

Approximately 88% of mHealth applications share data in ways inconsistent with their stated policies. This discrepancy creates a “Control Paradox”: UI-level consent revocation fails to synchronize with kernel-level sensor execution [11].

Most existing technologies operate exclusively at the application layer, relying on developer self-regulation — effectively an “honor system” that provides no verifiable guarantees [18]. The PEA middleware addresses this vulnerability by shifting the enforcement point to a native-layer interceptor, ensuring that user intent translates into physical API-level blockade rather than superficial UI toggles.

4.7 Trust Calibration and the Digital Divide

Trust in digital systems comprises both affective and cognitive dimensions [33]. Current mHealth applications frequently foster “Over-trust” through aesthetically polished interfaces while obscuring back-end data practices [32].

The OS-Verification Bridge is designed to achieve “Calibrated Trust” — a balance where user confidence aligns with the system’s actual reliability — through independent verification mechanisms [34]. Additionally, this study examines whether Skia-driven diagrams provide a universal mental model that reduces the Digital Divide for elderly populations with limited technical literacy [35], [36].

5 Research Gaps and Problem Statement

Dynamic Informed Consent (DIC) has yet to resolve the persistent challenges of the “Transparency Paradox.” A critical disconnection remains between abstract legal requirements and user-centric technical execution within the mHealth ecosystem [37].

5.1 Gap 1: UI/UX Semantic and Visual Overload

Current graphical interfaces frequently fail to bridge the cross-jurisdictional semantic gap arising from divergent legal definitions [38]. Most privacy dashboards aggregate legal clauses without contextual explanation, producing “meaningless disclosure” where users are technically informed but cognitively unable to process technical nuances [12], [14].

When information density exceeds cognitive thresholds, users resort to heuristic processing to manage visual complexity [28], [29]. Empirical studies indicate that users bypass rational evaluation in favor of shortcuts such as scrolling directly to the consent button [12]. Current dashboards employ a static, post-hoc disclosure model that presents privacy decisions only after data has already been accessed [26]. However, effective privacy decisions require real-time, context-aware feedback before data is conceded. A UI/UX framework characterized by low interaction density and pre-decisional disclosure remains absent from existing mHealth ecosystems.

5.2 Gap 2: The Control Paradox and Verification Gap

Existing mHealth models exhibit a fundamental disconnect between user intent and technical execution [11]. Aggregate data interfaces remain vulnerable to re-identification via neighboring dataset attacks when de-identification parameters are not strictly enforced [16].

The proposed PEA middleware addresses this vulnerability by shifting privacy enforcement from the application layer to a native-layer interceptor. This structural transition moves the architecture from manually-calibrated de-identification models toward a “strong-consistency” framework, ensuring that user-initiated consent revocation translates directly into kernel-level API blockade. The result is the elimination of the discrepancy between UI-level intent and hardware-level execution.

In current implementations, UI revocations frequently function as “soft-deletions” that leave native APIs accessible to background components [11]. Users possess no mechanism to verify whether an in-app toggle severs the data pipeline. This absence of a verifiable “Ground Truth” undermines the GDPR requirement for effortless consent withdrawal and creates a structural “Verification Gap” that erodes user trust.

5.3 Overarching Problem Statement

Current mHealth applications operate under a “Fractured Compliance Model”: they provide neither cognitive clarity at the front-end nor technical guarantees at the back-end. This thesis bridges this gap by translating

informed decisions directly into native-layer API execution, aligning psychological accessibility with physical-layer autonomy.

5.4 Theoretical Framework and Conceptual Model

This research integrates the Theory of Planned Behavior (TPB) and Privacy Calculus Theory (PCT) to frame the relationship between cognitive burden, perceived control, and behavioral intention [8], [39]. Skia-driven visualizations are hypothesized to reduce extraneous cognitive load, thereby restoring the user’s rational privacy decision-making capacity [27], [40].

Within the structural model, the inverse cognitive load ($1/CL$) serves as a gatekeeper variable for testing the moderation hypothesis. The theoretical logic posits that as NASA-TLX mental demand increases, cognitive load (CL) rises proportionally, decreasing the value of $1/CL$. Mathematically, elevated cognitive load acts as a nullifying factor: when CL exceeds threshold levels, the positive relationship between physical API blockade (operationalized as PBC) and behavioral intention is neutralized. Conversely, when the Skia-driven interface optimizes comprehensibility, the gatekeeper variable $1/CL$ maximizes the path coefficient between perceived behavioral control and intention. This integration ensures that technical robustness is functionally dependent on psychological accessibility within the mHealth context.

The model enhances Perceived Behavioral Control (PBC) through a native-layer interceptor that provides verifiable API-level blockade [39]. The research further hypothesizes that reducing cognitive load moderates the positive effect of enhanced PBC on behavioral intention. This moderation logic is expressed as:

$$\text{Behavioral Intention} \approx \beta_1(\text{Attitude}) + \beta_2(\text{Subjective Norms}) + \beta_3\left(\text{PBC} \times \frac{1}{CL}\right) \quad (1)$$

Where CL represents measured cognitive load (via NASA-TLX), and the interaction term $\text{PBC} \times \frac{1}{CL}$ captures the moderating effect of cognitive accessibility on the perceived control-behavior relationship. The theoretical relationship between perceived control, cognitive gate moderation, and behavioral intention is illustrated in the conceptual model in Figure 2.




Figure 2: Conceptual model integrating Theory of Planned Behavior and Cognitive Load moderation.

6 Research Questions

This thesis evaluates the effectiveness of the Skia-based visualization and the Privacy-Enhancing Architecture (PEA) across three research questions that address the intersection of human cognition and technical enforcement.

6.1 RQ1: Psychological Load and Cognitive Mitigation

To what extent can a Skia-based real-time feedback mechanism significantly reduce user “Mental Demand” and “Frustration” compared to traditional text-aggregated dashboards?

This question quantifies the reduction in extraneous cognitive load achieved by replacing abstract legal terminology with intuitive visual blur [29]. The hypothesis is that pre-attentive visual metaphors reduce the cognitive effort required to interpret privacy risk, thereby lowering mental demand and frustration during consent interactions. Measurement employs NASA-TLX sub-scales for mental demand and frustration [6], supplemented by post-task subjective ratings.

6.2 RQ2: Dual-Layer Trust and Perceived Control

How does the integration of an in-app “OS-Verification Bridge” with a native-layer Privacy Interceptor enhance user system trust and Perceived Behavioral Control (PBC)?

This question evaluates whether physical autonomy and independent verification restore user trust in mHealth applications [34], [39]. The OS-Verification Bridge provides a “Ground Truth” mechanism that cross-references application-layer consent status with native-layer API state. Perceived Behavioral Control (PBC) is hypothesized to increase when users can verify that UI actions produce kernel-level effects. Metrics include PBC scores and Trust scales assessed via Structural Equation Modeling (SEM) [41].

6.3 RQ3: Interaction Efficiency and Decision-Making Paradigms

Does the integration of context-aware UI elements optimize the time cost and accuracy of consent withdrawal tasks?

This question determines whether visual feedback dismantles visual noise and facilitates a transition from heuristic “blind agreement” to rational privacy calculus under high-frequency data transmission scenarios [12], [26]. The shift from System 1 (automatic) to System 2 (deliberate) processing is measured via task completion time, withdrawal accuracy, and post-hoc comprehension scores [22]. The hypothesis is that context-aware visual flows reduce decision latency while preserving informed consent quality.

7 Significance and Implications of the Study

This research bridges abstract legal mandates with quantifiable technical benchmarks, correlating the Privacy-Enhancing Architecture (PEA) with the user-centric objectives of data sovereignty. The significance of this work is organized across three dimensions: theoretical, methodological, and practical.

7.1 Theoretical Contribution: Modeling Cognition and Control

This study advances the theoretical understanding of privacy decision-making by integrating Cognitive Load Theory (CLT) into the Privacy Calculus model [8], [27], validating how visual metaphors dismantle visual noise and facilitate rational processing. The result is a theoretical foundation for the transition from heuristic agreement to rational privacy calculus.

This integration builds upon the evolution of semantic control frameworks established in clinical contexts [17]. Prior work demonstrated the logic for health record linkage but frequently neglected the interaction layer. This thesis addresses the cognitive-semantic disconnect by moving beyond policy-level flags toward real-time enforcement, refining the foundational logic of semantic reasoning in secure data exchange.

Additionally, the study introduces the concept of “Calibrated Trust” within the mHealth ecosystem [34]. Rather than relying on aesthetic-mediated affective trust, this concept positions cognitive trust grounded in verifiable reliability as the primary driver of user confidence. Conceptually, user intention is modeled as a function of verifiable sovereignty rather than interface polish. This addresses the re-identification risks identified in spatial research, where sensitive identities can be extracted from aggregated location metadata [15].

7.2 Methodological Significance: Standardizing Evaluation

The research proposes NASA-TLX scores and “Verification Frequency” as key performance indicators for privacy-compliant software [6], [34]. By operationalizing temporal demand and effort as evaluative dimensions, the methodology provides a replicable framework for comparing traditional text-based consent protocols against dynamic visualizations, standardizing the measurement of privacy-by-design effectiveness.

This methodological shift addresses audit vulnerabilities exposed in neighboring dataset attacks [16]. When records are dynamically added or deleted, adversaries may re-identify individual status changes unless privacy budgets are strictly enforced. The study implements a native-layer interceptor to ensure consistency between user intent and physical API-level blockade, offering a defensible method for technical audit.

Structural Equation Modeling (SEM) demonstrates that Privacy-by-Design is contingent upon psychological evaluation [41]. The framework posits that back-end robustness is effectively neutralized if it remains cognitively opaque to the end-user. Technical security and psychological accessibility are therefore not independent variables but interdependent design constraints that must be jointly optimized. This paradigm establishes an evaluative benchmark that quantifies software efficacy through the prism of the user’s cognitive

bandwidth.

7.3 Practical and Industry Implications

The findings offer a technical blueprint for the mHealth industry to resolve the Control Paradox. Empirical data demonstrate that Local Differential Privacy (LDP) visualizations manage perceived risk effectively, reducing cognitive friction without expanding the information leakage surface. These results provide evidence-based design patterns for complex medical data environments.

The study provides a functional pathway to satisfy mandatory transparency obligations for Automated Decision-Making (ADM) under the 2026 APP 1.7 amendments in Australia and the EU’s EHDS framework [2], [19]. By synthesizing multi-jurisdictional requirements into a unified verification architecture, this work formalizes a new standard for verifiable, psychologically accessible data sovereignty in the global mHealth landscape [32].

8 Research Methodology and Technical Architecture

This study employs a two-pronged “strong-consistency” approach to bridge the semantic and control gaps in mHealth ecosystems. The methodology ensures that findings are both technically defensible and psychologically valid.

8.1 Phase 1: Privacy-Enhancing Architecture (PEA) Implementation

The PEA moves the enforcement point from the application layer to the native OS layer, extending the lineage of three-layer semantically-driven access control frameworks established for health record linkage [17]. Rather than relying on centralized security policies, the architecture utilizes a native-layer interceptor proxy(see [42]), ensuring that cross-border compliance constitutes a physical blockade executed at the OS level rather than a UI-level flag [1], [20].

Figure 3 provides a high-level overview of the PEA workflow, demonstrating the synchronization between user intent and physical layer enforcement.




Figure 3: High-level workflow of the Privacy-Enhancing Architecture (PEA).

8.1.1 Technical Justification and Comparison of Alternatives

Three alternative privacy enforcement pathways were evaluated and subsequently rejected to ensure the technical appropriateness of the PEA:

- **Application-Layer Logic (The “Honor System”):** Most mHealth frameworks rely on internal application toggles. This approach was rejected because it functions as a developer self-regulation mechanism with no verifiable guarantees. Empirical data indicates that approximately 88% of mHealth applications share data in ways inconsistent with their stated policies, rendering this layer insufficient for “Strong Consistency” [11], [43].
- **Centrally Managed OS Policies:** While robust, centralized security policies (e.g., SELinux) are often too rigid for the granular, itemized consent requirements mandated by the GDPR and PIPL [1], [20]. Additionally, modifying the OS kernel introduces generalizability constraints for diverse participant populations, particularly elderly users with limited technical literacy [36].
- **Pure UI/UX Optimization:** Dashboard enhancements may improve “Perceived Control” but fail to bridge the technical “Verification Gap” [14]. Without a native-layer interceptor, the interface remains a “black-box” in which users cannot independently verify whether a toggle produces actual API-level severance [11].

The native-layer interception approach achieves physical autonomy while retaining the granular flexibility required for modern health data ecosystems.

8.1.2 Native Interceptor Middleware and OS-Verification Bridge

The middleware utilizes system-level API hooking to bypass the limitations of application-layer toggles. On iOS, the implementation employs Method Swizzling on sensitive `HKHealthStore` query methods. When a revoke command is active, the interceptor forces the query handler to return a null descriptor, causing the application to perceive data unavailability without triggering a system-level failure.

On Android, the system functions as a proxy for the AIDL interface to block sensor callbacks. This enables real-time injection of Local Differential Privacy (LDP) noise directly into the Binder IPC pipeline [44], [45]. The OS-Verification Bridge employs a deep-link observer to monitor OS-level permission statuses, maintaining the UI in a pending state until cross-layer consistency is confirmed. The specific technical pathways for achieving kernel-level API blockade across different operating systems are detailed in Figure 4.




Figure 4: Technical implementation of kernel-level API blockade for iOS and Android.

8.1.3 LDP Noise-to-Visual Transformation Pipeline




Figure 5: LDP Noise-to-Visual transformation pipeline using Skia Graphics Engine.

The system maps LDP parameters to visual blur radii via the Skia graphics engine. The Laplace Mechanism obfuscates continuous mHealth metrics (e.g., heart rate variability) to prevent temporal correlation attacks [46]. A cumulative privacy budget monitor tracks correlation risks across sequential data points. When the budget exceeds thresholds defined by regulatory frameworks such as the EHDS, the system dynamically increases visual opacity to communicate elevated risk levels to users, bridging the gap between abstract epsilon parameters and intuitive visual cues [19], [29]. The transformation pipeline that bridges the semantic gap by visualizing LDP noise as intuitive blurry paths is depicted in Figure 5

8.2 Phase 2 & 3: HCI Evaluation and Statistical Modeling

A randomized controlled trial ($n \approx 100$) validates the moderating effects of cognitive load on privacy decision-making. Participants are stratified by age and technical literacy to address digital divide concerns [35], [36].

Statistical analysis employs Partial Least Squares Structural Equation Modeling (PLS-SEM) to evaluate complex predictive relationships and interaction effects [7], [47]. PLS-SEM is selected for its high statistical power with small-to-medium sample sizes and its capacity to model both measurement and structural relationships simultaneously [48].

9 Research Schedule and Milestones

The research is structured into four distinct phases over two consecutive semesters. Specifically, Semester 1 focuses on foundation and the research proposal, while Semester 2 covers implementation and synthesis.

- **Phase 1: Foundation (Semester 1):** Focuses on problem identification and the comprehensive review of relevant literature.
- **Phase 2: Proposal and Ethics (Semester 1):** Involves the formal submission of the Research Proposal and the initiation of the ethics application process.
- **Phase 3: Implementation and Evaluation (Semester 2):** Covers technical prototype development and user-centric data collection.
- **Phase 4: Synthesis and Thesis (Semester 2):** Focuses on result interpretation, critical analysis, and the final thesis write-up.

Table 2: Research Project Timeline and Milestones

Timeline	Task / Milestone	Assessment/Deliverable
Semester 1		
Weeks 1–4	Research problem identification and research question formulation.	Decide and finalize Research Question
Weeks 5–6	Comprehensive literature review of relevant prior work and law regulations.	Lit Review Draft
Weeks 6–7	Draft RP to supervisor for review; leave time for receiving feedback	N/A
Week 8	Milestone: Research Proposal Submission.	30% (Hurdle, May 3 2026)
Weeks 9–12	Ethics application (HREC/IRB) and PEA development.	IRB Application
Weeks 13–14	Oral presentation preparation for a general CS audience.	Presentation Slides
Week 15	Milestone: Oral Presentation.	5% Assessment
Semester 2		
Weeks 1–2	Finalize PEA middleware	Technical Validation
Weeks 2–4	Pilot study ($n = 10$) and PEA stress testing	Pilot Validation.
Weeks 5–8	User-centric evaluation and NASA-TLX data collection.	Experimental Data
Weeks 9–10	Data analysis and interpretation of research results.	Results Analysis
Weeks 10–11	Thesis draft submission to supervisor for feedback.	N/A
Week 12	Final thesis synthesis and formatting (25,000–30,000 words).	Completed Thesis
Exam Week 1	Milestone: Final Thesis Submission.	65% Assessment

10 Evaluation and Data Analysis

The evaluation phase quantifies the psychological impact of the PEA-driven dashboard and validates the moderating effect of Cognitive Load. The design premise holds that technical enforcement and cognitive clarity must operate in tandem to restore user autonomy; back-end robustness is insufficient if the front-end remains cognitively opaque to end-users.

10.1 User-Centric Evaluation: Experimental Design




Figure 6: The five-stage experimental workflow: from recruitment to PLS-SEM analysis.

A randomized controlled trial (RCT) with a target sample of $n \approx 100$ ensures generalizable results. Participants represent a diverse range of technical literacy levels, addressing digital divide concerns and ensuring that the proposed architecture remains accessible to non-technical and elderly populations [35], [36].

Participants are randomly assigned to baseline, intervention, or visual-only control groups to isolate the effects of each architectural component. Group A utilizes standard list-based toggles; Group B engages with the full Skia-based dashboard and OS-Verification Bridge. This comparative structure isolates the specific contribution of each architectural layer to user trust and decision quality.

The five-stage experimental workflow, spanning from recruitment and stratification to PLS-SEM statistical synthesis, is mapped out in Figure 6.

10.2 Technical Evaluation: Strong Consistency and Performance

Technical evaluation focuses on the operational efficacy of the Privacy Interceptor Middleware and its synchronization with the user interface. Synchronization latency is measured against a target threshold of less than 50ms, ensuring that users perceive enforcement as instantaneous and reinforcing the mental model of physical autonomy.

Bridge consistency is quantified to verify that the OS-Verification Bridge accurately reflects the ground truth of native settings. Specifically, zero discrepancy should exist between UI-level permission status and OS-level API state. Stress testing with concurrent simulated users verifies that PEA overhead remains below 10% of total system resources.

Figure 7 clarifies the operational logic of the OS-Verification Bridge in identifying potential verification gaps between the application UI and kernel state.




Figure 7: Operational logic of the OS-Verification Bridge for ground truth checks.

10.3 Psychometric Quantification via NASA-TLX

The NASA-TLX scale captures multiple dimensions of user experience to evaluate the mitigation of privacy fatigue [6]. Analysis focuses on mental demand, temporal demand, and frustration sub-scales, quantifying the reduction in extraneous load achieved by replacing text-based notices with pre-attentive visual metaphors.

Frustration measurement is particularly significant for assessing the impact of the OS-Verification Bridge. Independent verification is expected to offset the anxiety associated with black-box data flows in standard mHealth applications. Providing a verifiable audit mechanism transforms the interaction from a developer self-regulation model into a transparent, user-verifiable system.

10.4 Statistical Robustness and PLS-SEM Modeling

Partial Least Squares Structural Equation Modeling (PLS-SEM) analyzes complex interaction effects within the data [7], [47]. A sample size of $n \approx 100$ provides sufficient statistical power to detect relevant path coefficients. The model employs 5,000 bootstrapping iterations to ensure robustness and generate precise confidence intervals.

Within the structural model, the term $1/CL$ operationalizes the moderation hypothesis: elevated NASA-TLX mental demand scores reduce the magnitude of this term, attenuating the positive influence of physical API blockade on user behavioral intention. When Skia-driven visualizations optimize comprehensibility, the gatekeeper variable $1/CL$ maximizes the path coefficient between perceived behavioral control and intent. This statistical relationship confirms that technical robustness requires cognitive accessibility to translate into functional user empowerment.

10.5 Hypotheses and Boundary Conditions

H1: Visual metaphors significantly lower perceived mental demand during privacy tasks. Pre-attentive risk cues such as the “Blurry Path” enable users to perceive privacy risk without engaging in deliberate cognitive effort, facilitating the transition from heuristic to rational privacy calculus [30].

H2: The OS-Verification Bridge produces more accurate trust calibration compared to isolated indicators. Ground truth linkage aligns user perception with the actual technical state of the API blockade, preventing both over-trust and under-trust in system privacy claims [34].

H3: Cognitive load significantly moderates the relationship between perceived control and usage intention. High cognitive load is expected to neutralize the benefits of technical control mechanisms.

H4: External verification reduces user frustration by addressing the verification gap. Transparency in API-level execution transforms the interaction from opaque processing to auditable consent management.

These hypotheses are subject to boundary conditions including contextual urgency, prior technical self-efficacy, and device heterogeneity across the participant sample.

11 Research Limitations

This study acknowledges several boundary conditions that may constrain the generalizability and technical stability of the Privacy-Enhancing Architecture (PEA).

11.1 Sample Representativeness and the Digital Divide

The online randomized sampling strategy ($n \approx 100$) may skew toward younger, technically literate cohorts, underrepresenting elderly users and their specific aging barriers [36]. This demographic limitation restricts the external validity of the proposed mental models to populations with lower digital literacy levels [35].

11.2 Cognitive Interpretation and Pre-attentive Efficiency

The effectiveness of visual metaphors is contingent upon the user’s baseline technical self-efficacy. Prior exposure to privacy-preserving technologies may modulate the efficiency of pre-attentive processing [29]; users with higher technical knowledge may interpret the “Blurry Path” more readily than novices. This differential baseline could attenuate the cognitive mitigation effects observed in the study sample.

11.3 The Verification-Friction Trade-off

The OS-Verification Bridge introduces a verification-friction trade-off. Navigating external settings to confirm privacy state increases operational burden and temporal demand, presenting a risk of “Verification Fatigue” wherein users abandon external validation in favor of heuristic trust judgments [40]. The threshold at which visual metaphors sufficiently substitute for manual verification remains an open empirical question.

11.4 Subjectivity in Psychometric Measurement

Self-reported instruments (NASA-TLX and SUS) introduce measurement subjectivity through social desirability bias and individual interpretation variance [6]. Since cognitive load is operationalized through perceived rather than neurological indicators, the findings reflect user perception of mental demand rather than objective cognitive transitions during privacy decision-making.

11.5 Ecological Validity and Experimental Context

The controlled laboratory setting may not replicate the ecological conditions of real-world clinical environments. In life-critical scenarios, privacy considerations frequently yield to immediate health imperatives; recorded performance metrics thus represent an idealized baseline rather than a behavioral proxy under genuine urgency.

11.6 Platform Specificity and Architectural Stability

The technical architecture faces constraints arising from mobile OS sandboxing and proprietary API evolution. Interceptor middleware may encounter permission penetration risks when interacting with proprietary native modules, and protocol changes may destabilize the OS-Verification Bridge as platform vendors update UI and kernel architectures [11], [49].

11.7 Regulatory Volatility and Compliance Gaps

The evolving 2026 regulatory landscape—including the EHDS and Australia’s APP 1.7 amendments—demands perpetual adaptability rather than fixed architectural stability [2], [19]. Mandated transparency for Automated Decision-Making (ADM) prescribes an inherently iterative development path for PEA logic. The architecture functions not as a static solution but as an evolving system capable of navigating cross-border compliance shifts.

11.8 Architectural Compatibility with HarmonyOS NEXT

A boundary condition of this research is its current focus on AOSP-based Android and iOS ecosystems. HarmonyOS NEXT employs a microkernel architecture and distributed security bus, presenting a distinct technical environment for privacy enforcement requirements under the PIPL [38]. While the PEA’s core philosophy remains applicable, the physical blockade mechanism would require migration from Binder-based proxies to HarmonyOS-specific “AbilityStage” interceptors for kernel-level sensor management. This platform variation limits immediate cross-platform generalizability but identifies a robust pathway for future inquiry into distributed operating systems privacy architectures.

12 Conclusion

This thesis bridges abstract legal mandates and technical implementation, formalizing a scalable framework that reconciles data utility with user sovereignty. The operationalization of GDPR’s “ease of withdrawal” requirement demonstrates that the Control Paradox yields to strategic synchronization between UI-level intent and kernel-level execution [1]. Collectively, these findings reposition Privacy-by-Design as a foundational structural requirement rather than a cosmetic UI enhancement.

The OS-Verification Bridge resolves the Verification Gap by providing an independent ground truth that addresses the historical absence of trust in application-layer toggles. Technical robustness is functionally dependent on psychological accessibility; physical autonomy must be matched by visual transparency to achieve Calibrated Trust.

The moderation finding—that cognitive load attenuates the relationship between control and intention—provides a design blueprint for future mHealth applications. This blueprint prioritizes user cognitive

capacity alongside data security measures. By ensuring that privacy intent is both intuitively formed and verifiably enforced, this work establishes a foundation for a more resilient digital health ecosystem capable of restoring data sovereignty in the 2026 regulatory landscape.

13 Proposed Thesis Structure and Outline

The thesis comprises approximately 25,000 to 30,000 words across seven chapters, documenting the end-to-end development and validation of the PEA framework.

13.1 Chapter Content Summary and Word Distribution

- **Chapter 1: Introduction (approx. 2,000 words)** defines the research scope, centering on the “Compliance Gap” and the “Verification Gap” in mHealth applications. It establishes the rationale for transitioning beyond honor-based privacy enforcement toward physically verifiable consent mechanisms.
- **Chapter 2: Literature Review (approx. 5,000 words)** provides a comparative legal analysis across three jurisdictions (EU, China, and Australia), critically examining the Semantic-Interaction Gap and employing Cognitive Load Theory to explain the phenomenon of Privacy Resignation [14], [40].
- **Chapter 3: Technical Architecture and Selection Logic (approx. 4,500 words)** justifies the native-layer interception approach over application-layer toggles through the following subsections:
 - *Section 3.1*: Justification of native-layer interception versus application-layer toggles (the honor system model).
 - *Section 3.2*: Analysis of Method Swizzling (iOS) and AIDL (Android Interface Definition Language) proxy (Android) as mechanisms for achieving strong consistency.
 - *Section 3.3*: Design of the OS-Verification Bridge as an independent ground truth provider.
- **Chapter 4: Implementation and Shader Logic (approx. 4,500 words)** documents the realization of the Skia-driven dashboard, providing a technical exposition of the GLSL shader logic that maps Laplace Mechanism (LDP) parameters into the “Blurry Path” visualization.
- **Chapter 5: Evaluation and Results (approx. 6,000 words)** presents the assessment of psychological load via NASA-TLX [6] and PLS-SEM modeling, validating how cognitive load ($1/CL$) moderates the path between technical control and behavioral intention [7], [47].
- **Chapter 6: Discussion and Cross-Platform Scalability (approx. 4,500 words)** addresses three key dimensions:
 - *Section 6.1*: The verification-friction trade-off in real-world clinical urgency contexts.
 - *Section 6.2*: Cross-ecosystem portability—evaluating PEA logic migration from monolithic AOSP kernels to HarmonyOS NEXT microkernel environments [38].
 - *Section 6.3*: Mitigating programmer-induced user suffering through HCI-centric design principles.

- **Chapter 7: Conclusion and Future Work (approx. 1,500 words)** synthesizes how the study bridges abstract legal mandates with physical enforcement, outlining implications for Privacy-by-Design in the 2026 regulatory landscape and identifying directions for future mHealth privacy research.

References

- [1] General Data Protection Regulation (GDPR), *Regulation (EU) 2016/679 of the european parliament and of the council*, 2016.
- [2] Parliament of Australia, *Privacy and other legislation amendment act 2024 (cth)*, Provisions for Automated Decision-Making (APP 1.7) mandatory in late 2026, 2024. [Online]. Available: <https://www.aph.gov.au/>.
- [3] Mordor Intelligence, *mHealth market size & share analysis - growth trends & forecasts (2026 - 2031)*, 2026.
- [4] Deloitte, “2017 global mobile consumer survey: US edition,” Deloitte, Tech. Rep., 2017.
- [5] H. Nissenbaum, “A contextual approach to privacy online,” *Daedalus*, vol. 140, no. 4, pp. 32–48, 2011.
- [6] S. G. Hart and L. E. Staveland, “Development of NASA-TLX (Task Load Index): Results of empirical and theoretical research,” *Advances in Psychology*, vol. 52, pp. 139–183, 1988.
- [7] J. F. Hair, G. T. M. Hult, C. M. Ringle, and M. Sarstedt, *A primer on partial least squares structural equation modeling (PLS-SEM)*, 3rd. Sage Publications, 2022.
- [8] T. Dinev and P. Hart, “An extended privacy calculus model for e-commerce transactions,” *Information Systems Research*, vol. 17, no. 1, pp. 61–80, 2006.
- [9] IQVIA Institute for Human Data Science, “Digital health trends 2021: Innovation, evidence, regulation, and adoption,” IQVIA Institute for Human Data Science, Tech. Rep., Jul. 2021, Accessed: 2026-03-24. [Online]. Available: <https://www.iqvia.com/insights/the-iqvia-institute/reports/digital-health-trends-2021>.
- [10] C. P. Hoffman et al., “The mechanics of privacy resignation: Why users quit managing their data,” *New Media & Society*, 2024, Forthcoming/Online First.
- [11] L. Parker et al., “88% of android mHealth apps share data: A systematic assessment of privacy and security of mobile health apps,” *BMJ*, vol. 373, no. n1248, 2021.
- [12] A. M. McDonald and L. F. Cranor, “The cost of reading privacy policies,” *I/S: A Journal of Law and Policy for the Information Society*, vol. 4, p. 543, 2008.
- [13] J. A. Obar and A. Oeldorf-Hirsch, “The biggest lie on the internet: Ignoring privacy policies and terms of service anymore?” *Information, Communication & Society*, vol. 23, no. 1, pp. 128–147, 2018.
- [14] S. Liao et al., “Beyond notice and choice: Investigating the impact of visual feedback latency on privacy resignation in mHealth ecosystems,” *Computers in Human Behavior*, vol. 152, pp. 108–120, 2024.
- [15] S. Sun, “An exploration of spatial privacy,” The University of Melbourne, Capstone Project Report, 2015, Supervised by Prof. Richard O. Sinnott. [Online]. Available: <https://github.com/Siqis/spatialprivacyexploration>.

- [16] M. U. Altaf, “Privacy oriented data analysis, with a focus on location-specific privacy issues,” Supervised by Prof. Richard O. Sinnott. 19,384 words., M.S. thesis, The University of Melbourne, 2017. [Online]. Available: <https://github.com/mumeraltaf/Thesis>.
- [17] Y. Lu, “Development and evaluation of a semantic-based access framework for health record linkage,” Supervised by Prof. Richard O. Sinnott. Focuses on semantic security and decentralized record linkage., Ph.D. dissertation, The University of Melbourne, 2018.
- [18] R. Guo and F. Chen, “Privacy policy compliance of mobile sports and health apps in china: Scale development, data analysis, and prospects for regulatory reform,” *JMIR mHealth and uHealth*, vol. 14, 2026, ISSN: 2291-5222. DOI: <https://doi.org/10.2196/73651>. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S2291522226000288>.
- [19] European Commission, “Proposal for a regulation of the european parliament and of the council on the european health data space,” European Commission, COM(2022) 197 final, 2022. [Online]. Available: https://health.ec.europa.eu/publications/proposal-regulation-european-health-data-space_en.
- [20] National People’s Congress, *Personal information protection law of the people’s republic of china*, 2021. [Online]. Available: <http://www.npc.gov.cn/>.
- [21] Parliament of Australia, *Privacy act 1988 (cth)*, 1988.
- [22] D. Kahneman, *Thinking, Fast and Slow*. Farrar, Straus and Giroux, 2011.
- [23] J. Kaye et al., “Dynamic consent: A patient-centred approach to e-health,” *European Journal of Human Genetics*, vol. 23, no. 2, pp. 141–146, 2015.
- [24] Z. Wang, A. Stell, and R. O. Sinnott, “The australasian diabetes data network (addn) dynamic consent model,” *PMC9957235*, 2023.
- [25] E. Xiong et al., “The australians together health initiative (athena): Dynamic consent platform,” *JMIR Formative Research*, 2024.
- [26] J. Muchagata, P. Vieira-Marques, and A. Ferreira, “Mhealth applications: User-adaptive visualization,” in *Proceedings of the 21st ICEIS*, 2019.
- [27] J. Sweller, P. Ayres, and S. Kalyuga, *Cognitive Load Theory*. Springer Science & Business Media, 2011.
- [28] E. R. Tufte, *The visual display of quantitative information*. Graphics Press, 2001.
- [29] M. Agrawala, W. Li, and F. Berthouzoz, “Design principles for visual communication,” *Communications of the ACM*, vol. 54, no. 4, pp. 60–69, 2011.
- [30] P. Cavanagh, “Visual attention,” *Vision Research*, vol. 51, no. 13, pp. 1538–1551, 2011.
- [31] C. G. Healey and J. T. Enns, “Attention and visual memory in visualization,” *IEEE Transactions on Visualization and Computer Graphics*, vol. 18, no. 7, pp. 1170–1188, 2012.

- [32] A. Corman, R. Canaway, C. Culnane, and V. Teague, “Technical guarantees of privacy-preserving technologies and public understanding,” *PubMed*, 2022.
- [33] D. H. McKnight, V. Choudhury, and C. Kacmar, “Developing and validating trust measures for e-commerce: An integrative typology,” *Information Systems Research*, vol. 13, no. 3, pp. 334–359, 2002.
- [34] J. D. Lee and K. A. See, “Trust in automation: Designing for appropriate reliance,” *Human Factors*, vol. 46, no. 1, pp. 50–80, 2004.
- [35] J. van Dijk, *The Digital Divide*. Cambridge, UK: Polity Press, 2020, ISBN: 9781509534449.
- [36] G. A. Wildenbos, L. Peute, and M. Jaspers, “Aging barriers influencing mobile health usability for older adults: A literature segmentation strategy (lss),” *International Journal of Medical Informatics*, vol. 114, pp. 66–75, 2018. DOI: 10.1016/j.ijmedinf.2018.03.011.
- [37] S. Turner and L. M. Tanczer, “In principle vs in practice: User, expert and policymaker attitudes towards the right to data portability in the internet of things,” *Computer Law & Security Review*, vol. 52, p. 105912, 2024, ISSN: 2212-473X. DOI: <https://doi.org/10.1016/j.clsr.2023.105912>. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S026736492300122X>.
- [38] H. Fan and Y. Zhao, “A comparative study of personal information protection in china and the EU: From the perspective of mHealth data sovereignty,” *International Data Privacy Law*, vol. 13, no. 2, pp. 112–128, 2023.
- [39] I. Ajzen, “The theory of planned behavior,” *Organizational Behavior and Human Decision Processes*, vol. 50, no. 2, pp. 179–211, 1991.
- [40] S. H. Ali and H. Thu, “When more becomes less: The rising toll of app fatigue,” *PMC12637926*, 2025.
- [41] R. B. Kline, *Principles and Practice of Structural Equation Modeling*, 4th. Guilford Publications, 2016.
- [42] R. Xu, H. Saïdi, and R. Anderson, “Aurasium: Practical enforcement of user privacy policies via app repackaging,” in *Proceedings of the 21st USENIX Security Symposium*, 2012, pp. 27–42.
- [43] N. Artenstein and I. Revivo, *Man in the binder: He who controls the binder, controls the android*, Black Hat USA, 2014.
- [44] X. Wang et al., “Deepdroid: Dynamically enforcing enterprise-level weapons-grade security policy on android,” in *Proceedings of the 2015 Network and Distributed System Security Symposium (NDSS)*, 2015. [Online]. Available: <https://www.ndss-symposium.org/ndss2015/deepdroid/>.
- [45] D. Wu and S. Bratus, “A context-aware kernel ipc firewall for android,” Dartmouth College, Computer Science, Tech. Rep., 2016. [Online]. Available: https://www.cs.dartmouth.edu/~sergey/binderfilter/binderfilter_tr.pdf.
- [46] C. Dwork, “Differential privacy,” *International Colloquium on Automata, Languages, and Programming*, pp. 1–12, 2006.

- [47] N. Kock and P. Hadaya, “Minimum sample size estimation in PLS-SEM: The inverse square root and gamma-exponential methods,” *Information Systems Journal*, vol. 28, no. 1, pp. 227–261, 2018.
- [48] D. Barclay, C. Higgins, and R. Thompson, “The partial least squares (PLS) approach to causal modeling: Personal computer adoption and use as an illustration,” *Technology Studies*, vol. 2, no. 2, pp. 285–309, 1995.
- [49] M. Egele, C. Kruegel, E. Kirda, and G. Vigna, “Pios: Detecting privacy leaks in ios applications,” in *Proceedings of the 18th Network and Distributed System Security Symposium (NDSS)*, 2011.